Frozen versions: full-18 validation and test trends

ALL18 = equal weight per task; frozen evaluations only; versions are not independent repetitions.

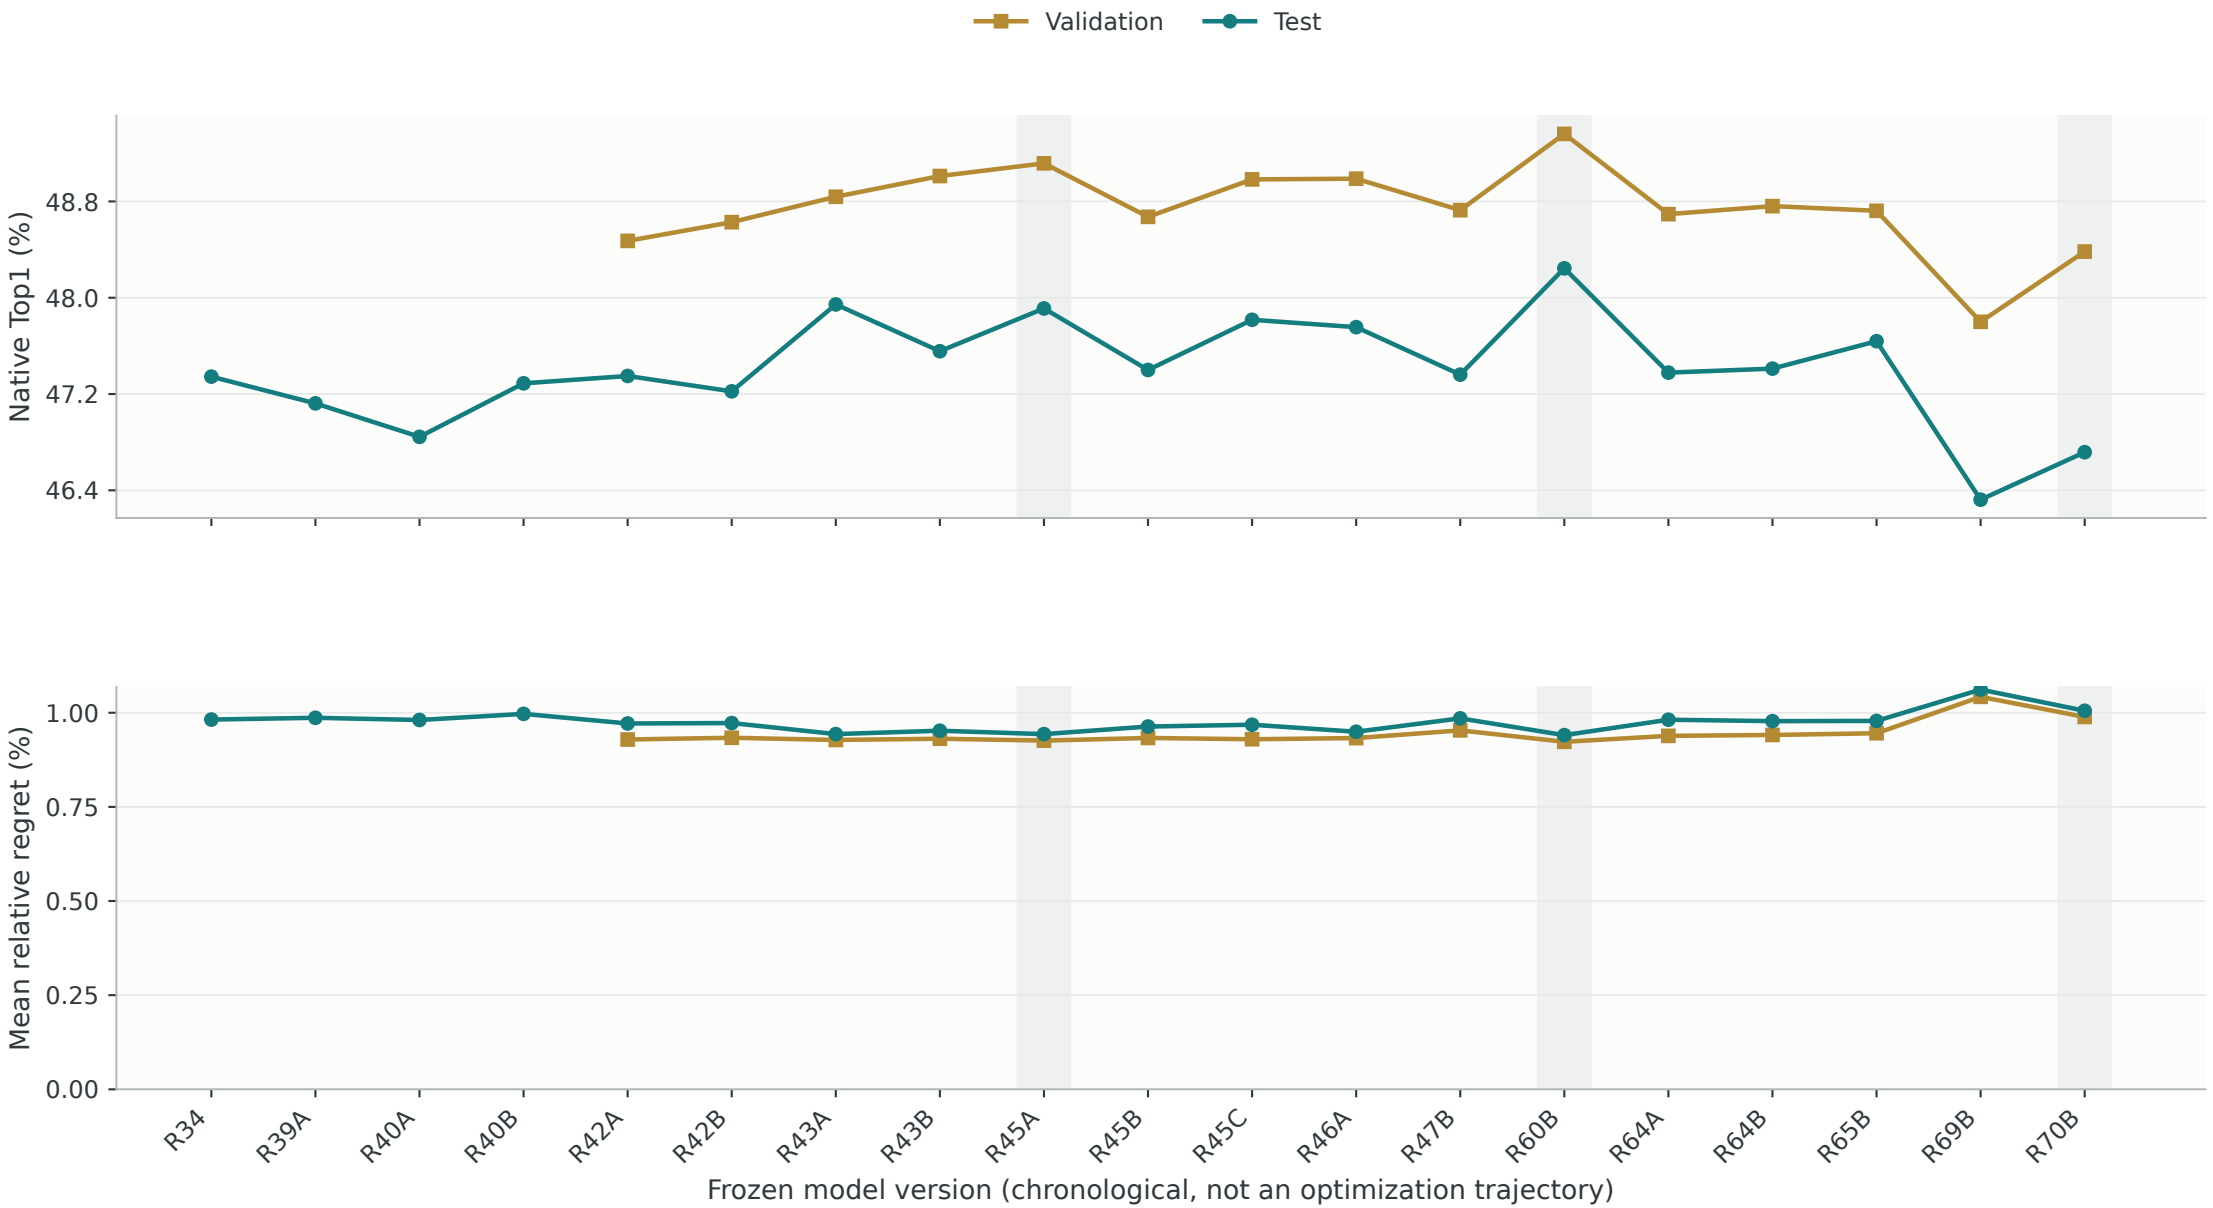

Regret = mean[100 x (selected cost / instance-best cost - 1)]. Missing validation points remain gaps. No four-dataset aggregate and no fitted performance ceiling are shown.