Frozen models agree on predictions and share error sets

Test | Every off-diagonal entry is an equal-task mean over all 18 problems; NSS is not a full-18 model.

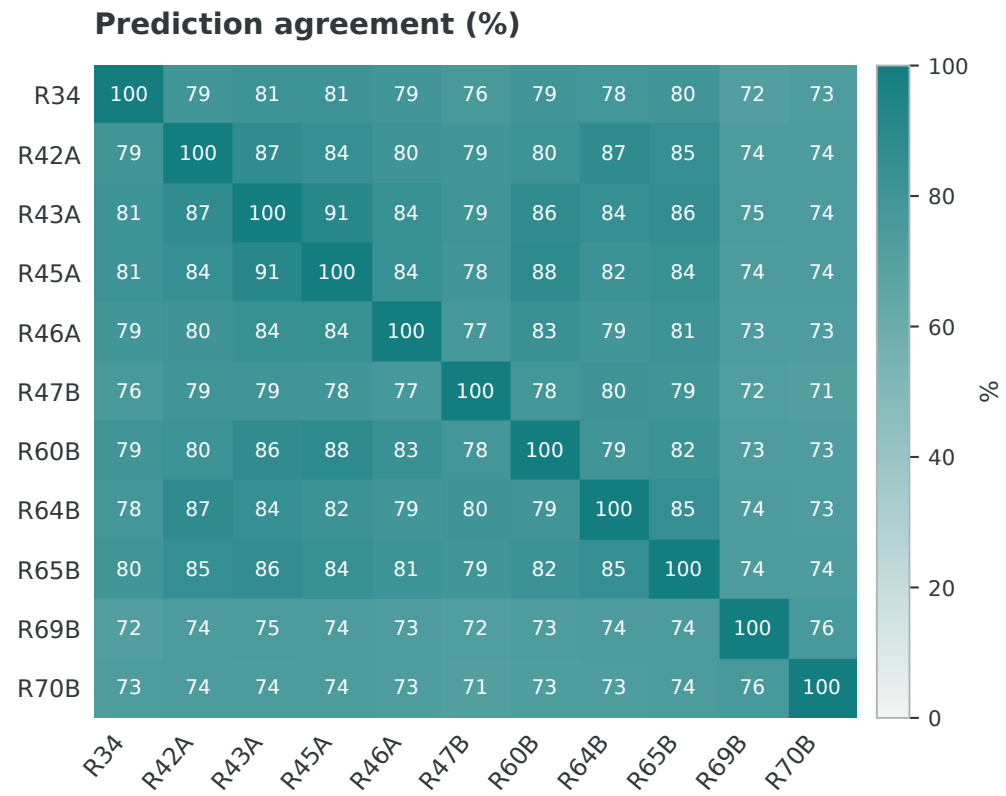

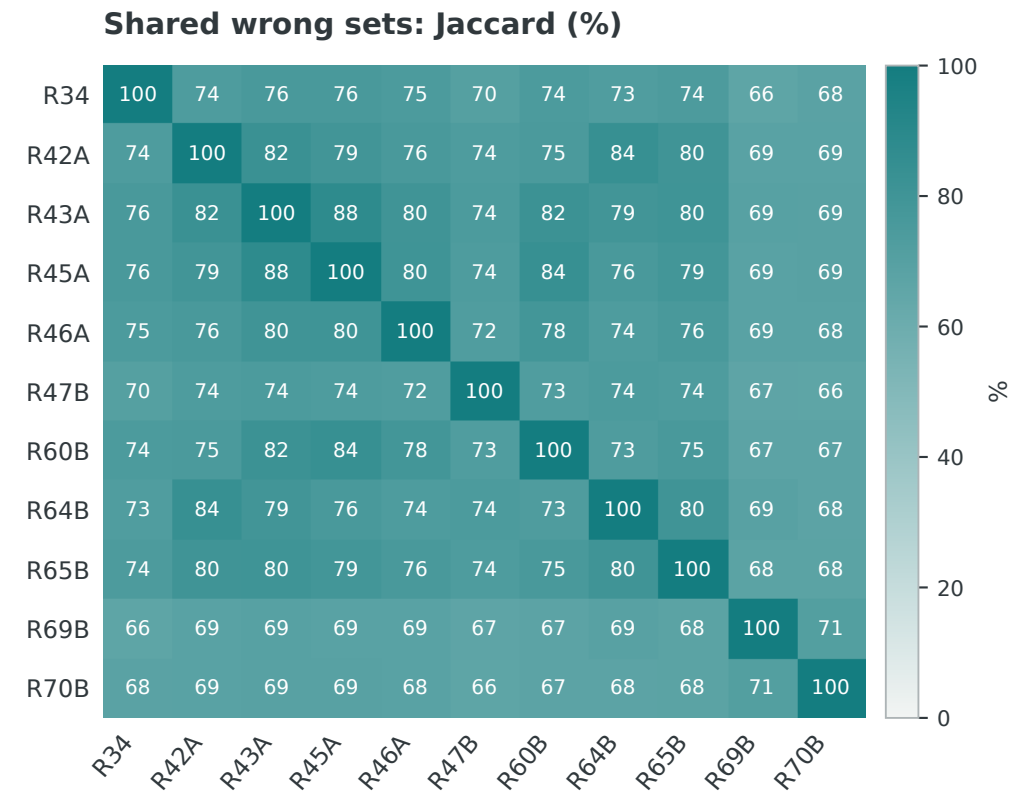

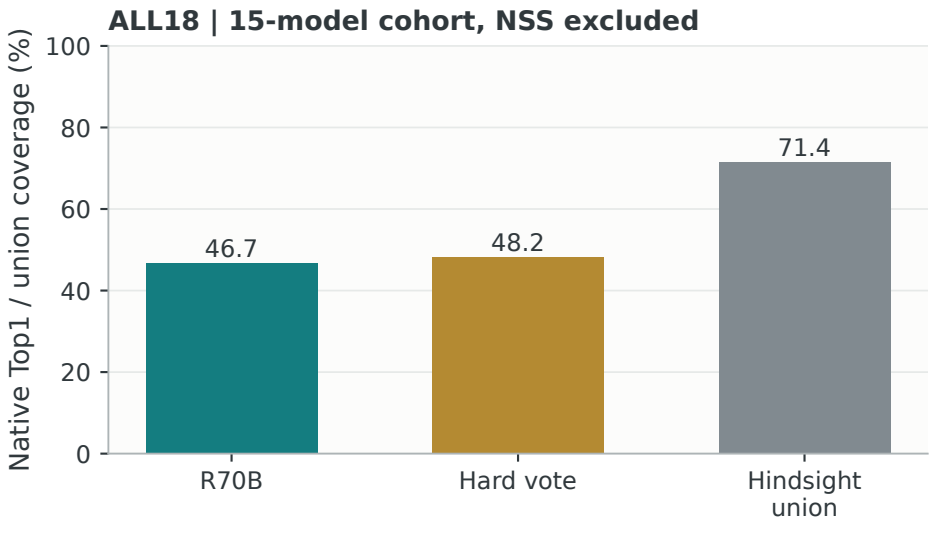

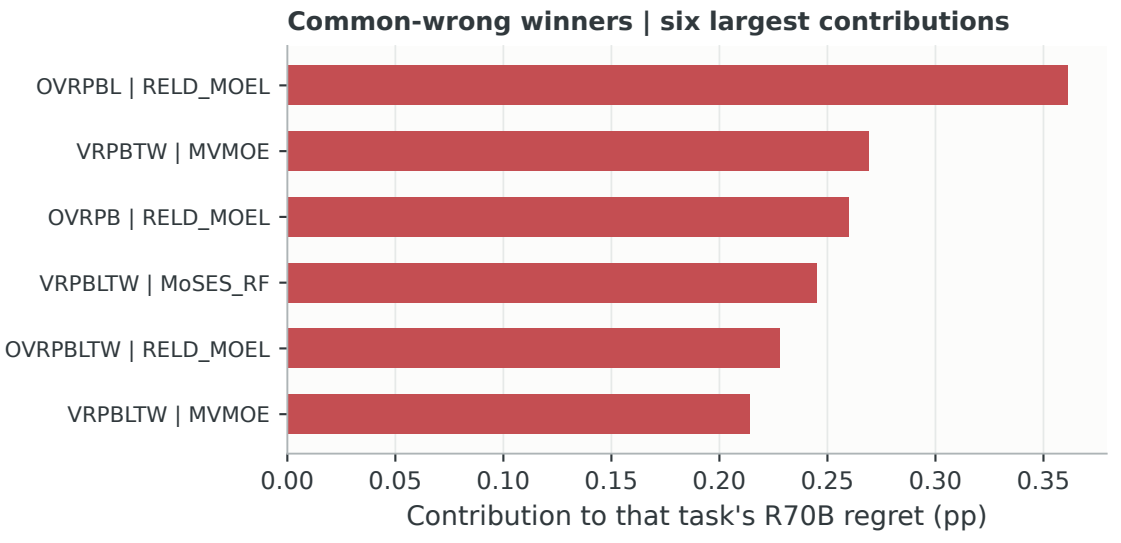

Error Jaccard = wrong for both / wrong for either, within each task before averaging; not joint error rate. Blank matrix cells denote incomplete task pairs. Bottom cohort excludes NSS and can exceed the matrix subset. Union is hindsight coverage, not a deployable selector or a ceiling. Hard vote has unmeasured multi-model overhead; right rows are test-sorted diagnostic task/solver pairs, not a pooled macro.