Score confidence, empirical correctness, and cost of wrong picks

Test | Max-softmax scores are descriptive rankings, not validated or calibrated probabilities. R25 is excluded.

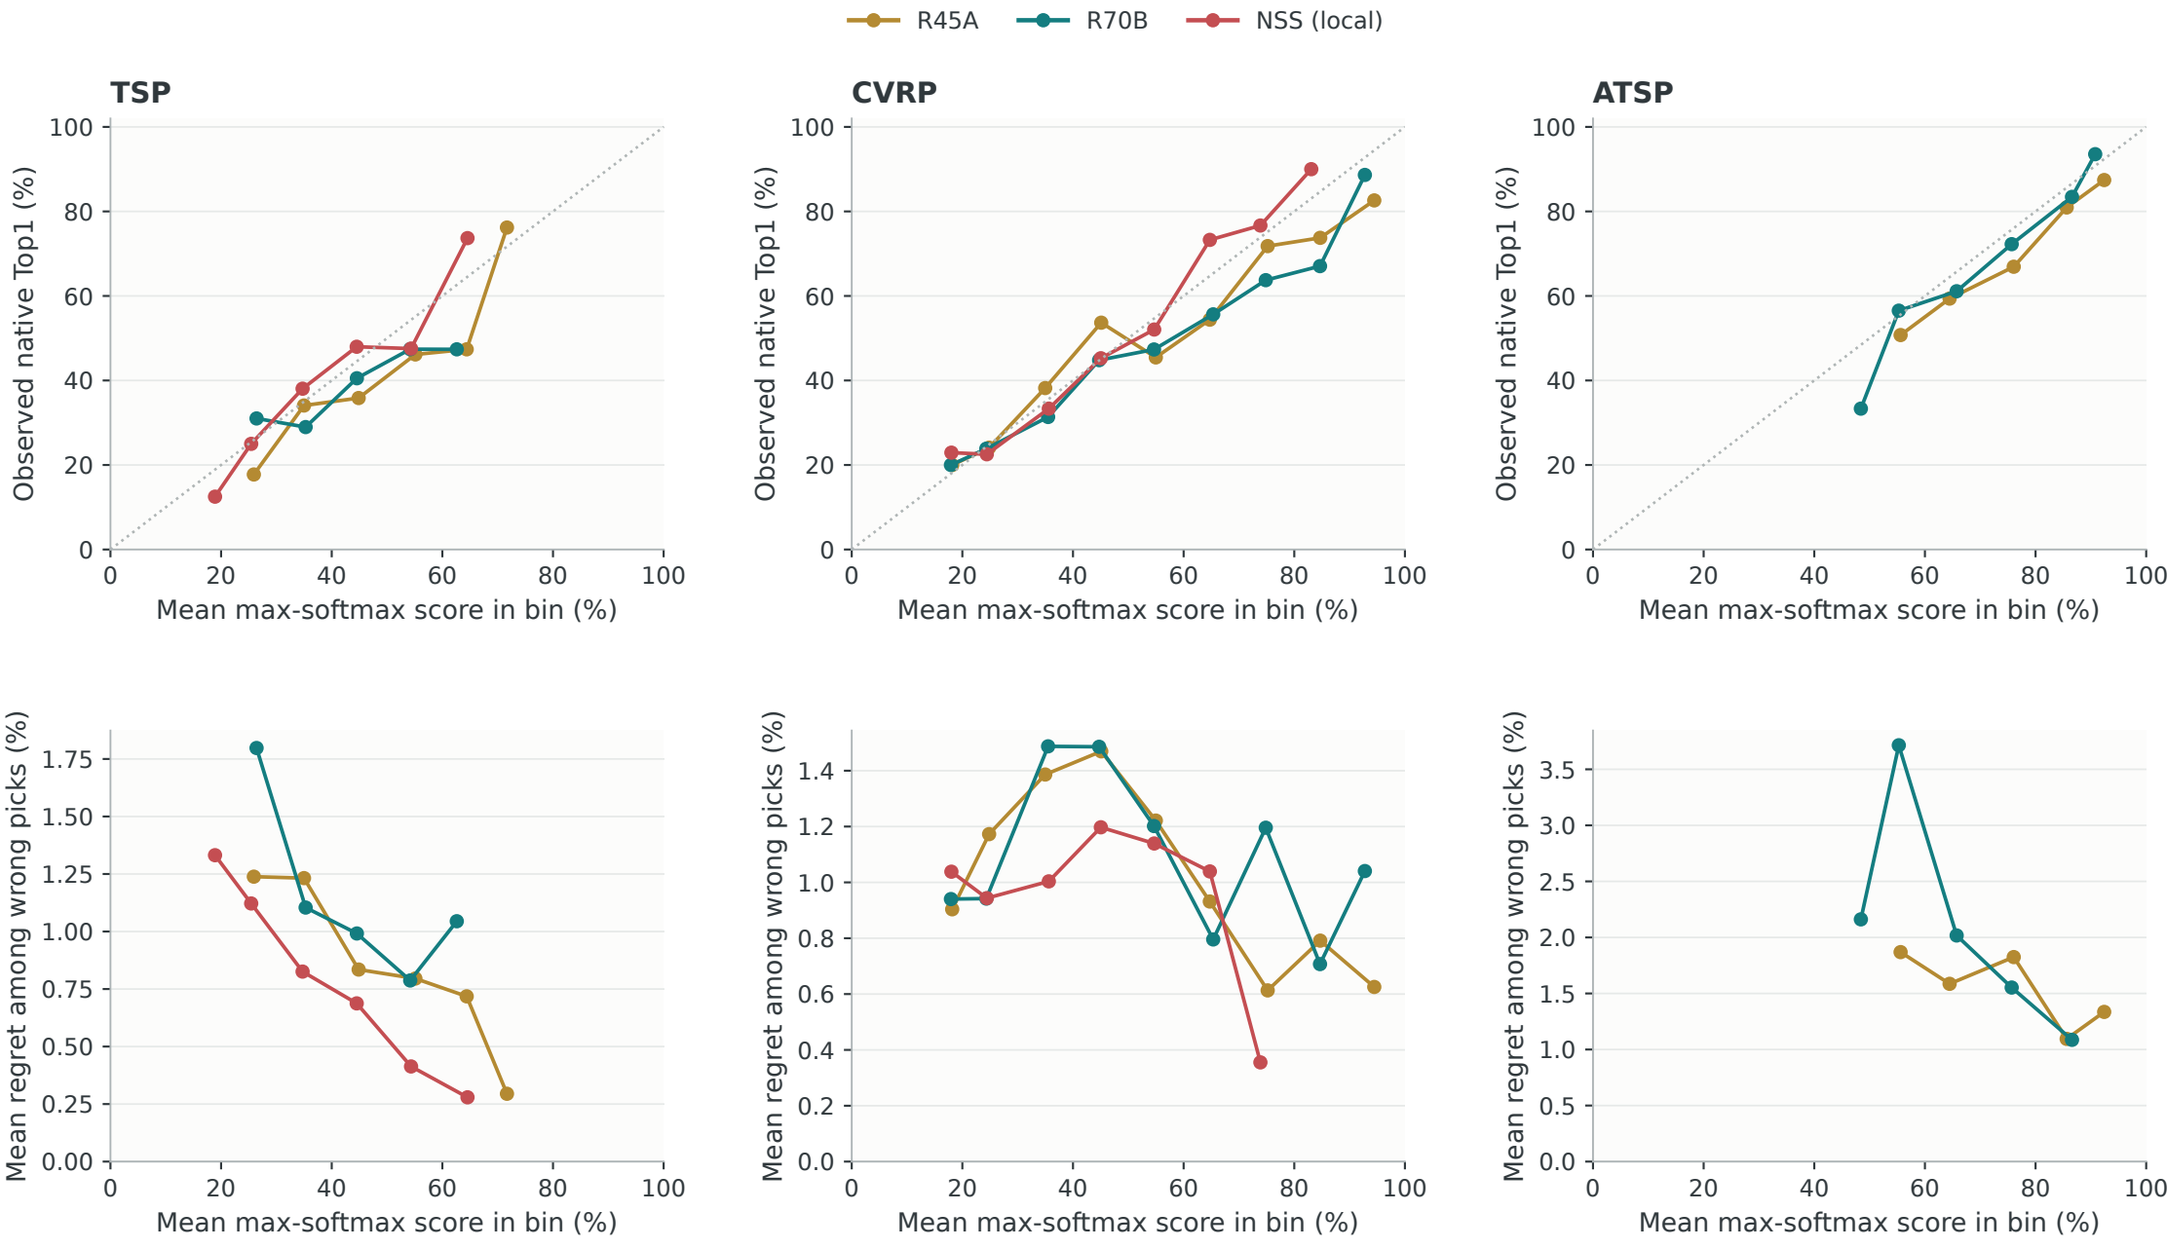

Ten fixed score bins; top row requires >=10 instances, bottom also >=5 wrong picks. Wrong-only costs are calculated directly from existing prediction NPZs, not inferred by dividing bin-average regret by error rate. Dotted diagonal is a visual reference, not a calibration claim.