R85C | Locked result and missed-method recovery

Epoch 25 / update 36000; historical R84A update 37440. Same NEW90k release; no new inference.

Split	Model	Top1 (%)	mean_cost	Actual regret (%)
VAL	R84A	71.0852	11.416611477	0.403494
TEST	R84A	71.2358	11.373843044	0.394847
VAL	R85C	71.0086	11.416206518	0.397969
TEST	R85C	71.3111	11.373420994	0.389713

Joint VAL gate: FAIL. Lower cost/regret do not compensate for lower VAL Top1.

Published 36 zero/near-zero targets: 4,625 VAL winners. No target redefinition using C or TEST.

Split/model	Winners	Correct	Recall %	Picks	Precision %	Rescue/lost	New false
VAL/R84A	4625	11	0.238	17	64.71	0/0	0
VAL/R85C	4625	22	0.476	35	62.86	11/0	10
TEST/R84A	4598	13	0.283	25	52.00	0/0	0
TEST/R85C	4598	24	0.522	46	52.17	14/3	13

Split	Decision change	Count	Regret delta (pp)	Mean cost delta
VAL	corrected	1,154	-0.020942	-0.001841182
VAL	harmed	1,216	0.018585	0.001708139
VAL	both_wrong_cheaper	655	-0.008042	-0.000760773
VAL	both_wrong_dearer	526	0.004874	0.000488856
TEST	corrected	1,193	-0.020206	-0.001910734
TEST	harmed	1,132	0.017398	0.001715761
TEST	both_wrong_cheaper	633	-0.007864	-0.000766954
TEST	both_wrong_dearer	528	0.005538	0.000539877

Negative deltas favor C. Category contributions use all instances per task, then equal task weights.

ALL mean_cost is descriptive only; cross-task raw-cost sums are not pooled. TEST cannot override VAL.

R85C | Full18: VAL and TEST kept separate

Each task: 4,500 instances per split; native/FP64 winner. Deltas = C minus historical R84A.

Task	VAL Top1 pp	VAL regret pp	VAL cost	VAL fix/harm	TEST Top1 pp	TEST regret pp	TEST cost	TEST fix/harm
TSP	-0.2444	0.004100	0.0004019	68/79	0.5111	0.001192	0.0001779	92/69
CVRP	-1.2222	-0.040449	-0.0026379	161/216	0.0667	-0.050510	-0.0049659	193/190
ATSP	0.0222	-0.005905	-0.0000943	41/40	0.0667	-0.003746	-0.0000673	32/29
OVRP	0.7556	-0.001673	0.0000134	287/253	-0.0889	-0.015831	-0.0009793	274/278
VRPB	-0.3111	0.013279	0.0009708	7/21	0.0444	-0.000869	0.0000187	11/9
VRPL	-0.0889	-0.003431	-0.0005242	216/220	0.5556	-0.011091	-0.0012222	209/184
VRPTW	0.0667	-0.003046	-0.0004596	6/3	0.0667	-0.001103	-0.0002482	5/2
OVRPTW	0.0222	-0.003301	-0.0002652	2/1	-0.0667	0.001577	0.0001320	1/4
OVRPB	0.0000	0.000000	0.0000000	0/0	0.0444	-0.002314	-0.0001414	3/1
OVRPL	-0.5778	-0.016185	-0.0012477	319/345	0.5111	-0.022502	-0.0015745	333/310
VRPBL	-0.1556	0.001547	0.0001425	7/14	-0.0222	0.006705	0.0006222	12/13
VRPBTW	0.0444	-0.000844	0.0001169	9/7	-0.0444	0.000412	-0.0001375	6/8
VRPLTW	0.0889	-0.005272	-0.0008798	8/4	-0.2222	0.006759	0.0011377	2/12
OVRPBL	0.0222	-0.008169	-0.0004051	3/2	0.0667	0.003418	0.0001718	5/2
OVRPBTW	0.1333	-0.025145	-0.0020401	6/0	-0.0667	0.002479	0.0002784	1/4
OVRPLTW	0.0444	-0.000603	-0.0002113	5/3	0.0000	-0.009728	-0.0008998	2/2
VRPBLTW	0.0000	0.003480	0.0006091	5/5	-0.0222	0.001481	0.0000182	10/11
OVRPBLTW	0.0222	-0.007833	-0.0007788	4/3	-0.0444	0.001256	0.0000824	2/4
ALL	-0.0765	-0.005525	-0.0004050	1154/1216	0.0753	-0.005134	-0.0004221	1193/1132

Complete absolute metrics, TRAIN-fitted SBS, Oracle, vs_SBS and vs_Oracle are in metrics_by_problem.csv.

Top1 improvements and regret reductions need not occur on the same tasks. No per-task model assembly.

Only the prelocked VAL screen can qualify C; this report does not select another checkpoint or inference rule.

R85C | TRAIN suppression and existing LIB evidence

TRAIN predictions are already saved for each selected checkpoint; this analysis does not occupy a GPU.

Model	TRAIN Top1 %	Zero groups	Zero winner mass	Target correct/support	Recall %
R84A	71.2439	26	25,627	297/83279	0.357
R85C	71.1817	27	39,906	419/83279	0.503

TRAIN task/method	Winners	R84A picks/correct	C picks/correct	C recall %	C precision %
TSP/LEHD	14,279	15/7	0/0	0.000	undefined
TSP/DIFUSCO	6,806	0/0	0/0	0.000	undefined
ATSP/MATPOENET	6,476	0/0	0/0	0.000	undefined
CVRP/RouteFinder	1,354	0/0	0/0	0.000	undefined

Decision suppression persists on TRAIN; unseen-instance overfitting alone is not an adequate explanation.

Same locked C model	n	Top1 (%)	mean_cost	Actual regret (%)
TSPLIB	49	18.3673	8.215885069	0.938997
CVRPLIB	100	63.0000	67.274507074	0.591596

LIB rows quote the existing completed report; no new inference or LIB-based model selection.

Missing evidence: none required here. TRAIN comparison reuses selected-state counts, not newly computed AUC.

One seed, historical R84A reference; R85A is still needed as the newly matched control.

No causal proof, no accuracy ceiling claim, no acceptance from fewer zero groups or a better TEST Top1.

Next: wait for other R85 arms. The research goal remains active and unresolved.