R85 | Locked full18 results

Split	Model	Top1 (%)	mean_cost	Actual regret (%)
VAL	R84A	71.0852	11.416611477	0.403494
TEST	R84A	71.2358	11.373843044	0.394847
VAL	R85A	70.9679	11.417059356	0.406901
TEST	R85A	71.2321	11.374097328	0.397280
VAL	R85B	70.9358	11.416844932	0.405866
TEST	R85B	71.2111	11.373750026	0.394872
VAL	R85C	71.0086	11.416206518	0.397969
TEST	R85C	71.3111	11.373420994	0.389713
VAL	R85D	71.1272	11.416646199	0.403383
TEST	R85D	71.4370	11.373489772	0.390487

Same NEW90k release and native/FP64 winner labels. SBS fitted on TRAIN; ALL = equal18-task mean.

Each model has one previously locked VAL checkpoint. TEST never changes a checkpoint or rule.

The long-term research goal remains active even if an empirical VAL screen passes.

R85 | Validation screens and matched EQL control

Model	Baseline	VAL dTop1 pp	VAL dCost	VAL dRegret pp	VAL screen	TEST dTop1 pp
R85A	R84A	-0.117284	0.000447879	0.003406	False	-0.003704
R85B	R84A	-0.149383	0.000233455	0.002371	False	-0.024691
R85C	R84A	-0.076543	-0.000404959	-0.005525	False	0.075309
R85D	R84A	0.041975	0.000034722	-0.000111	False	0.201235
R85B	R85A	-0.032099	-0.000214424	-0.001035	False	-0.020988

Screen: unrounded VAL Top1 strictly higher; mean_cost and actual regret not higher.

B versus R85A separates EQL from the common OVR classification change.

A small screen pass is evidence for review, not completion or deployment selection.

R85 | Published R84A zero / near-zero targets

Split	Model	Winners	Recall %	Rescued / lost	New false picks	New-FP regret pp
VAL	R84A	4625	0.238	0 / 0	0	0.000000
TEST	R84A	4598	0.283	0 / 0	0	0.000000
VAL	R85A	4625	0.108	0 / 6	2	0.000075
TEST	R85A	4598	0.065	1 / 11	4	0.000239
VAL	R85B	4625	0.324	7 / 3	8	0.000368
TEST	R85B	4598	0.196	3 / 7	8	0.000111
VAL	R85C	4625	0.476	11 / 0	10	0.000585
TEST	R85C	4598	0.522	14 / 3	13	0.000982
VAL	R85D	4625	0.368	8 / 2	13	0.001320
TEST	R85D	4598	0.435	10 / 3	23	0.000682

Frozen from R84A locked VAL: 36 groups / 4,625 winners, choice share <=0.1%.

Posthoc mechanism targets, not a training rule. Strict-zero results and precision are in target_summary.csv.

One false choice can shrink the zero list; true rescue and new false picks are reported together.

R85 | Same-instance corrections and harms

Split	Model / baseline	Corrected	Harmed	Both wrong cheaper / dearer	Cost delta	Regret delta pp
VAL	R85A / R84A	1023	1118	534 / 555	0.000447879	0.003406
TEST	R85A / R84A	1069	1072	548 / 513	0.000254284	0.002434
VAL	R85B / R84A	919	1040	527 / 538	0.000233455	0.002371
TEST	R85B / R84A	936	956	573 / 490	-0.000093018	0.000025
VAL	R85C / R84A	1154	1216	655 / 526	-0.000404959	-0.005525
TEST	R85C / R84A	1193	1132	633 / 528	-0.000422050	-0.005134
VAL	R85D / R84A	1003	969	479 / 449	0.000034722	-0.000111
TEST	R85D / R84A	1072	909	500 / 440	-0.000353272	-0.004360
VAL	R85B / R85A	780	806	421 / 421	-0.000214424	-0.001035
TEST	R85B / R85A	731	748	462 / 430	-0.000347303	-0.002409

Every pair uses identical saved instances, candidates, FP64 costs and native winners.

Native-winner errors can have zero extra cost in exact ties. Negative cost/regret delta is better.

Per-task results, all-method recall/precision, target effects and full checkpoint hashes accompany this PDF.