R89/A_adamw | Actual optimizer and fixed target recovery

One fixed-role original selector; optimizer is the only controlled training change.

Actual class: torch.optim.adamw.AdamW

Group	Tensors	Parameters	Initial LR	Weight decay
group0	148	1,842,953	0.0001	0.0001

Split	Winners	Correct	Rescued	Lost	New false picks	Recall (%)
TRAIN	83279	331	80	46	76	0.3975
VAL	4625	15	5	1	6	0.3243
TEST	4598	16	3	0	6	0.3480

Muon: Moonshot .2*sqrt(max dimensions), same engine LR/WD; no5x/aspect multiplier.

Actual per-step Muon application and direction RMS are saved in optimizer_runtime.csv.

TRAIN/VAL/TEST targets are frozen from R84A. Zero groups alone are not a success metric.

R89/A_adamw | Optimizer execution evidence
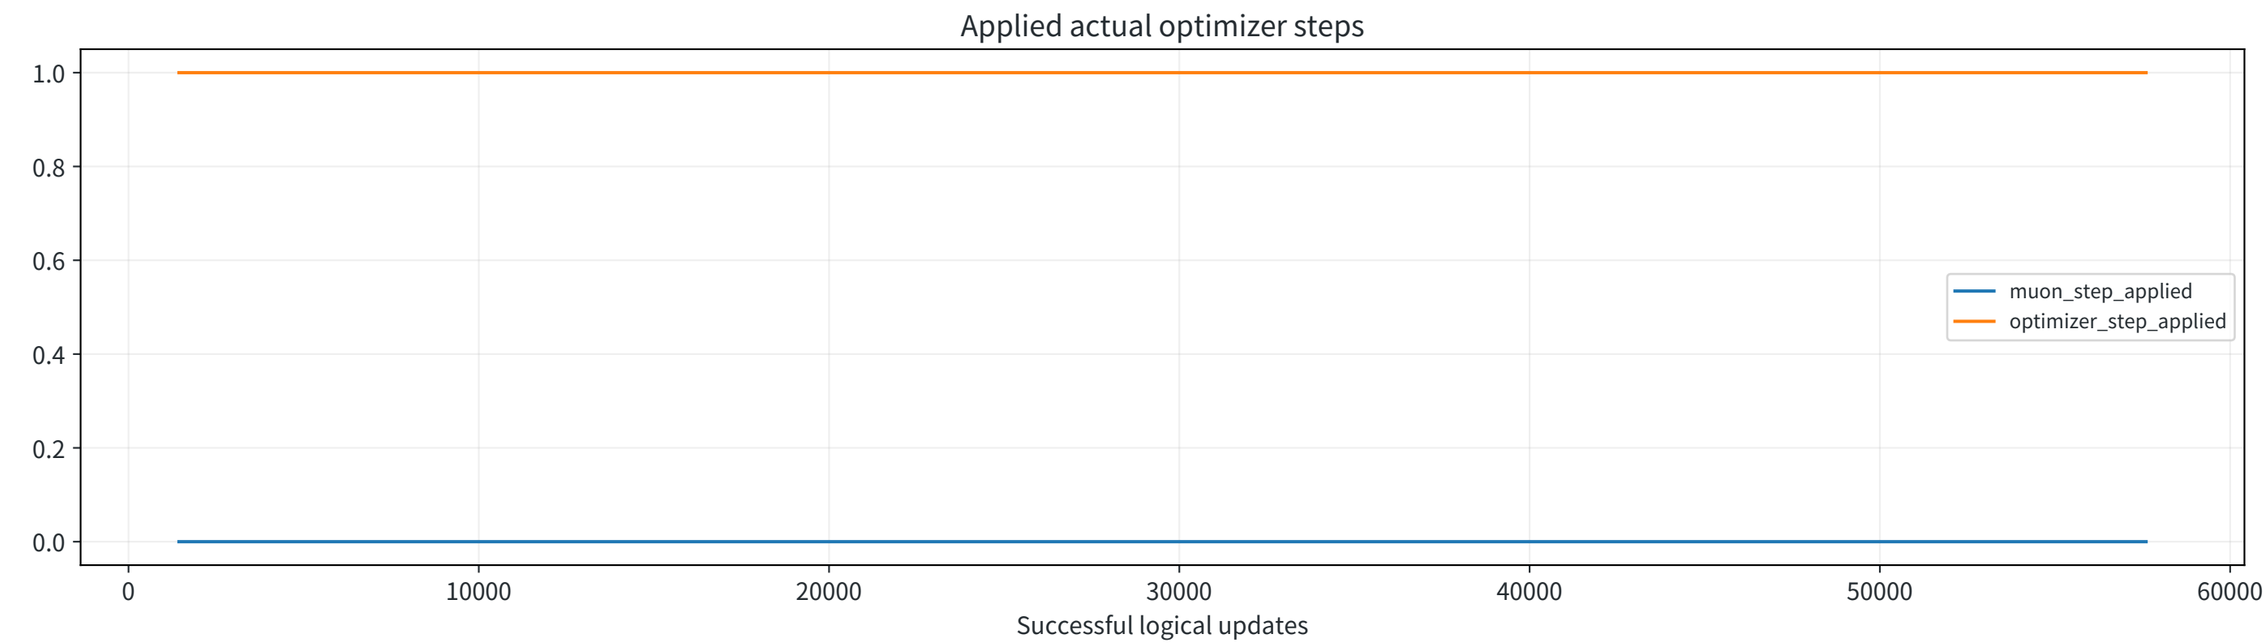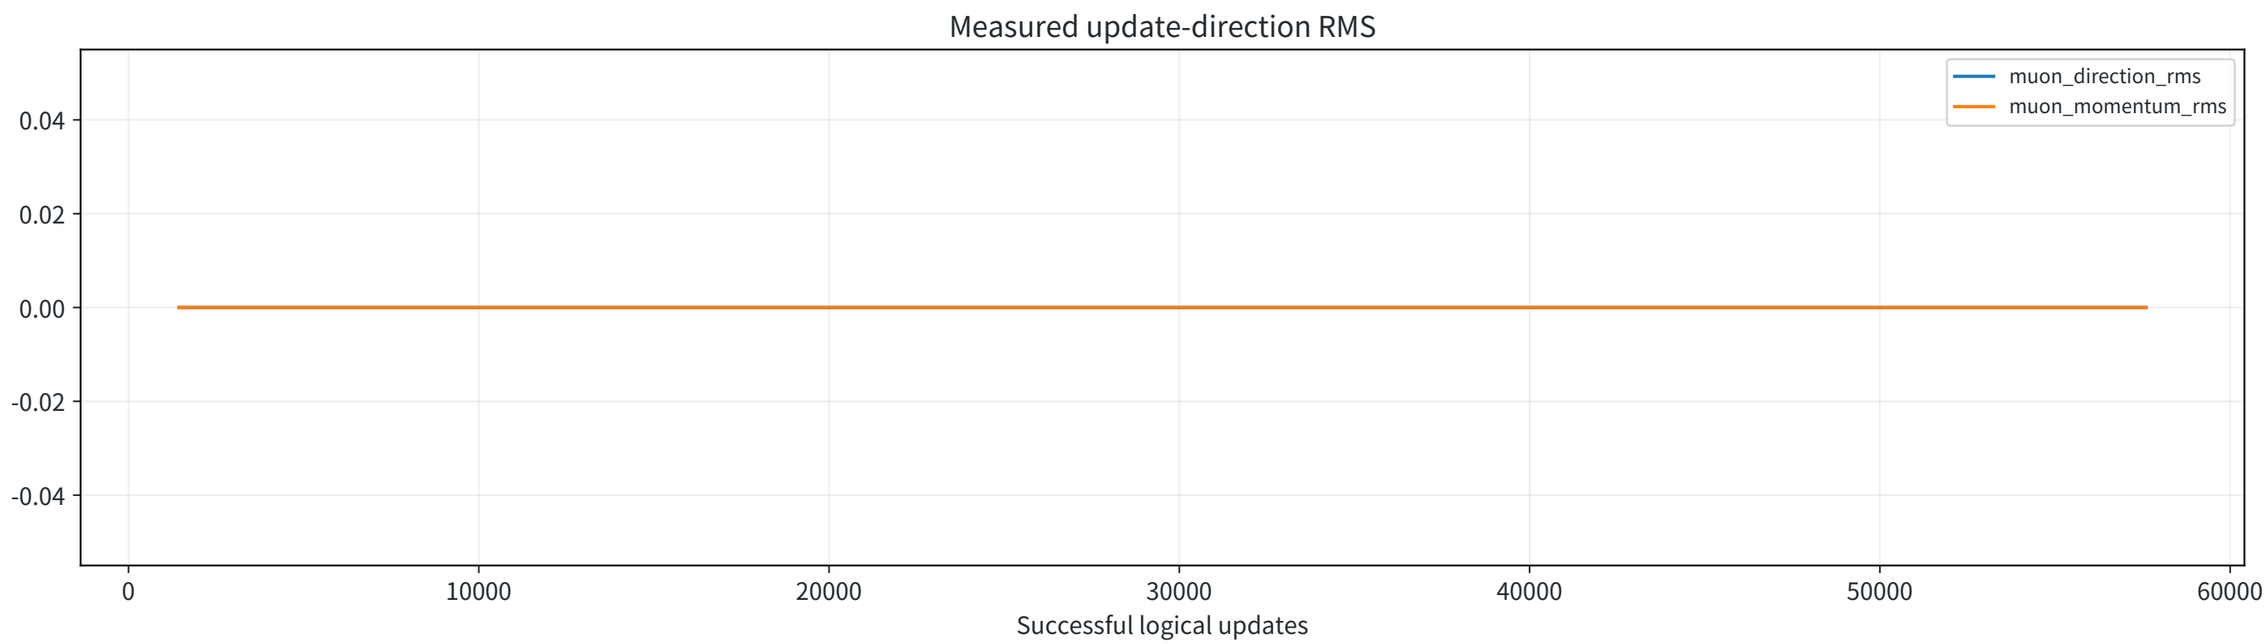
Interval means of actual whole-logical-batch steps; no per-micro optimizer statistics.