R89/B_muon_hybrid | Actual optimizer and fixed target recovery

One fixed-role original selector; optimizer is the only controlled training change.

Actual class: code.V4.r89_optim.MuonHybrid

Group	Tensors	Parameters	Initial LR	Weight decay
muon	55	1,785,856	0.0001	0.0001
adamw	93	57,097	0.0001	0.0001

Split	Winners	Correct	Rescued	Lost	New false picks	Recall (%)
TRAIN	83279	1024	765	38	603	1.2296
VAL	4625	52	42	1	35	1.1243
TEST	4598	45	34	2	42	0.9787

Muon: Moonshot .2*sqrt(max dimensions), same engine LR/WD; no5x/aspect multiplier.

Actual per-step Muon application and direction RMS are saved in optimizer_runtime.csv.

TRAIN/VAL/TEST targets are frozen from R84A. Zero groups alone are not a success metric.

R89/B_muon_hybrid | Optimizer execution evidence
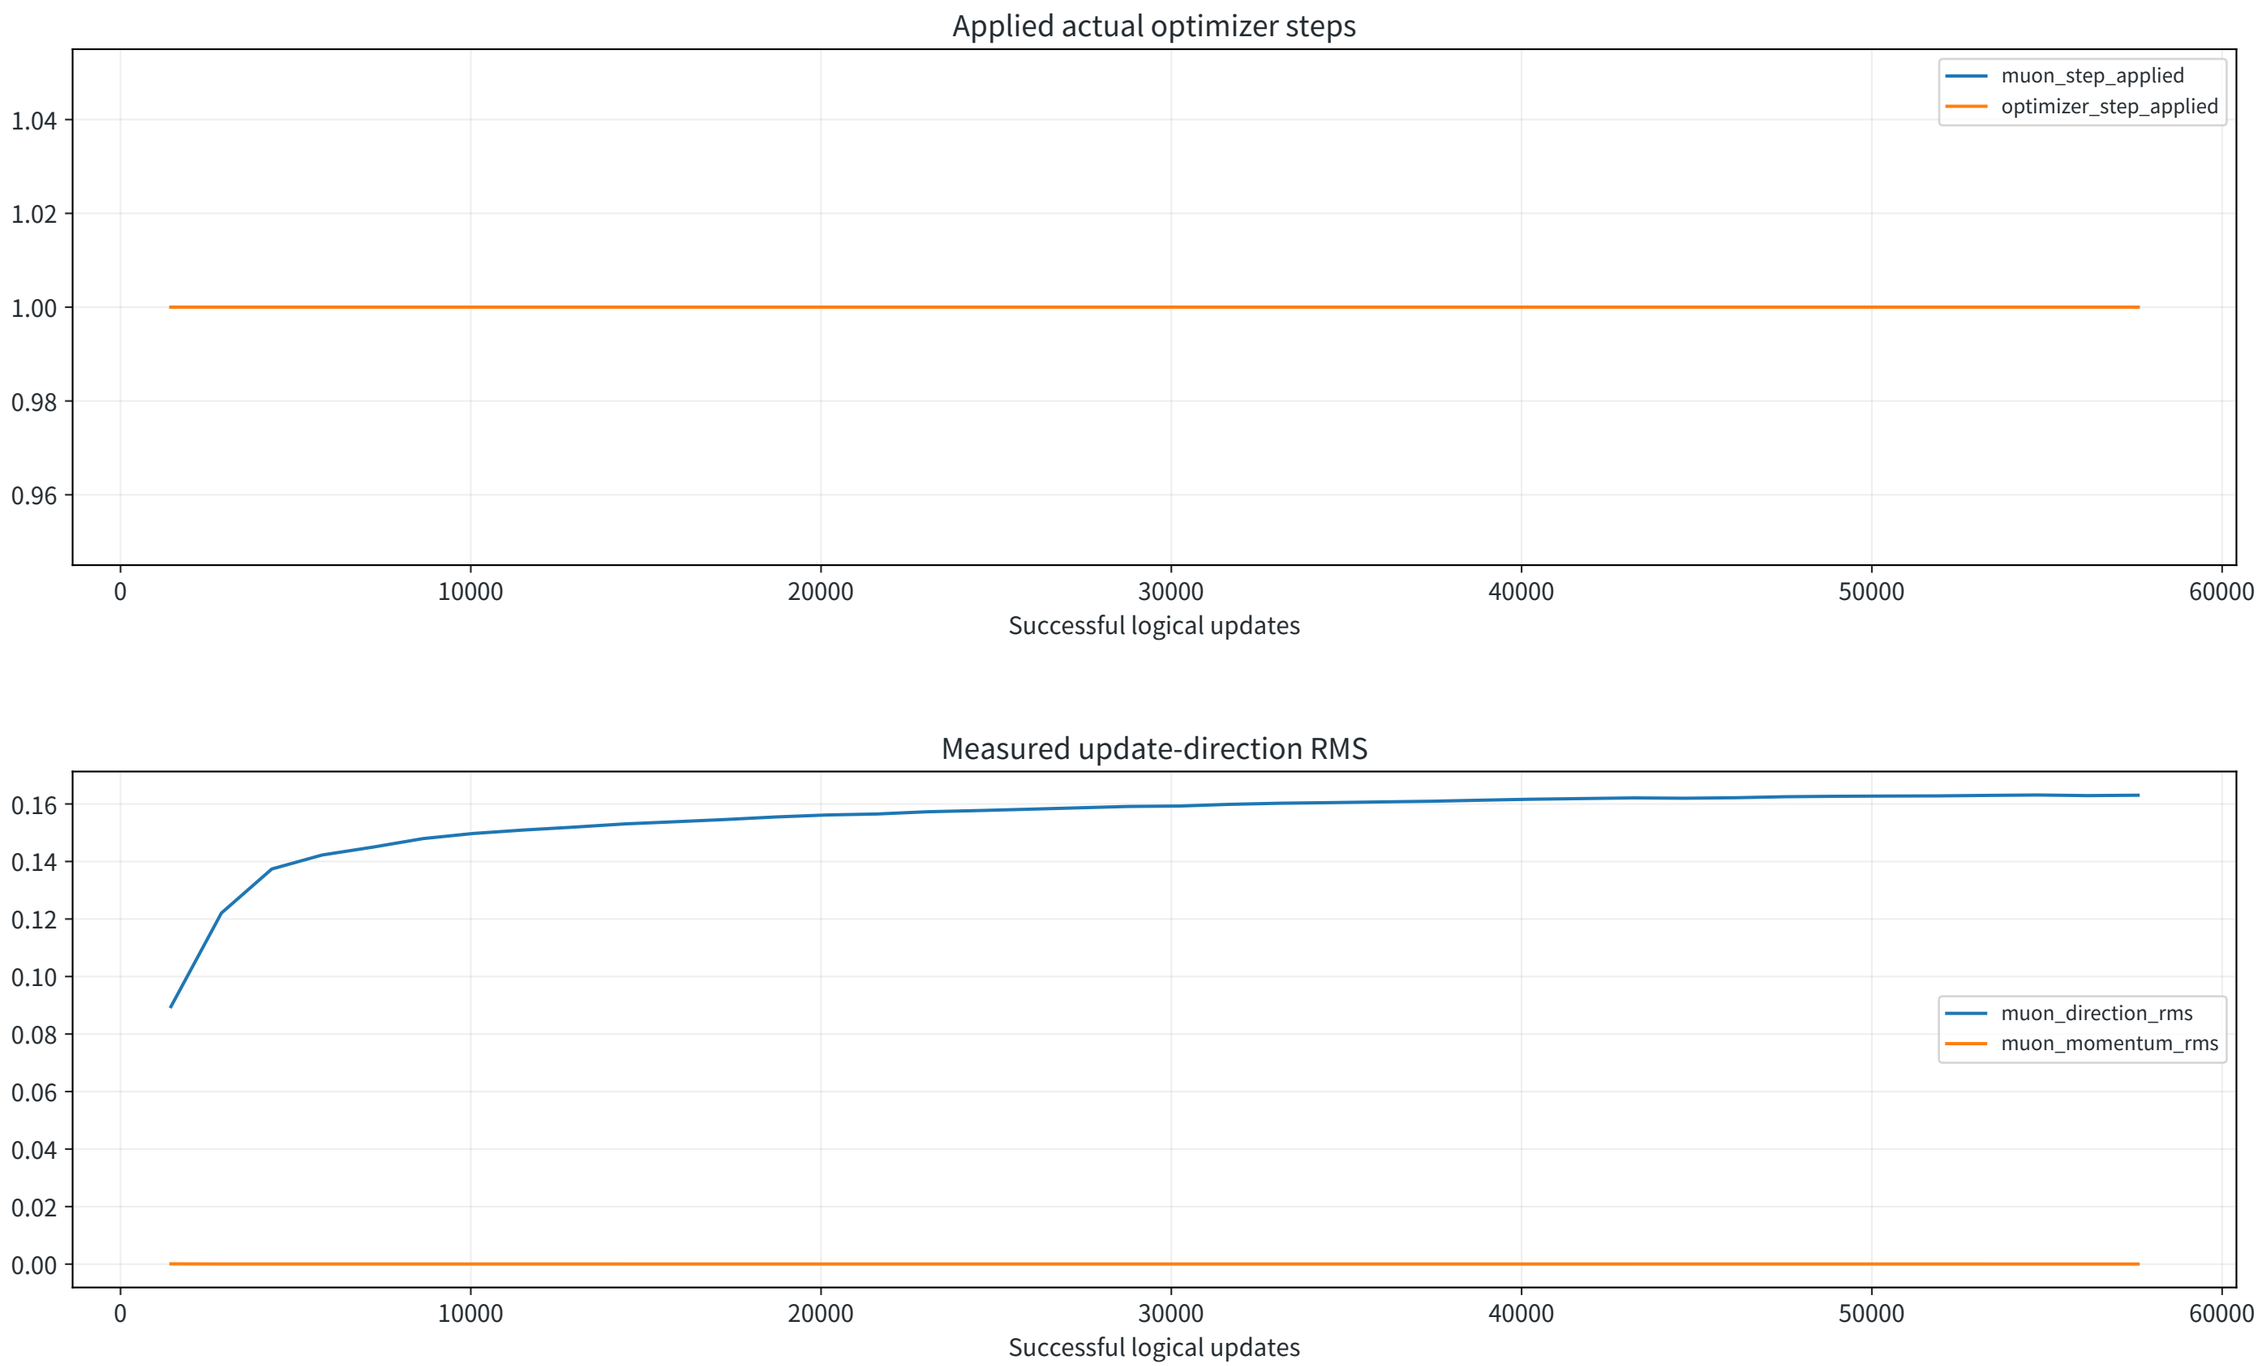
Interval means of actual whole-logical-batch steps; no per-micro optimizer statistics.