R89 | Muon hybrid minus AdamW

Identical model/input/objective; unique VAL checkpoint per arm. Saved predictions only.

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	0.446639	28988	22476	-0.001036326	-0.011067
VAL	0.197531	1492	1332	-0.000537980	-0.005005
TEST	0.187654	1505	1353	-0.000315034	-0.004590

Split	Target correct	Rescued	Lost	New false picks	Recall (%)
TRAIN	1024	745	52	589	1.2296
VAL	52	38	1	29	1.1243
TEST	45	31	2	39	0.9787

All18 tasks have equal macro weight; candidate pools and FP64 labels are the same.

Optimizer diagnostics are execution evidence, not acceptance criteria. TEST does not select a new rule.