R93/A | Fixed target36 vs R89B

Same unique VAL checkpoint; native/FP64 winner; no reselection.

Split	Winners	Correct	Rescued	Lost	New false picks	Recall (%)
TRAIN	83279	928	406	502	580	1.1143
VAL	4625	54	29	27	36	1.1676
TEST	4598	57	29	17	53	1.2397

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	-0.231207	21421	24792	0.001609953	0.016307
VAL	-0.122222	1202	1301	0.000768160	0.007198
TEST	-0.043210	1268	1303	0.001194993	0.012568

Historical R89B main / R84A secondary; not newly matched original-loss controls.

TRAIN CE = raw utility only. Eval macro_ce = original maximin diagnostic, not the optimized loss.

No comparison/risk/R-Drop; WD=1e-4 retained. Zero-count movement alone is not success.

A: released3090_1; B:4090_1. Not bitwise hardware-matched; tiny differences are not dropout causal proof.

R93/A | Fixed target36 vs R84A

Same unique VAL checkpoint; native/FP64 winner; no reselection.

Split	Winners	Correct	Rescued	Lost	New false picks	Recall (%)
TRAIN	83279	928	677	46	710	1.1143
VAL	4625	54	45	2	48	1.1676
TEST	4598	57	46	2	63	1.2397

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	0.216049	24557	21407	0.000525970	0.004743
VAL	0.040741	1271	1238	0.000450966	0.004380
TEST	0.140741	1345	1231	0.000822401	0.007367

Historical R89B main / R84A secondary; not newly matched original-loss controls.

TRAIN CE = raw utility only. Eval macro_ce = original maximin diagnostic, not the optimized loss.

No comparison/risk/R-Drop; WD=1e-4 retained. Zero-count movement alone is not success.

A: released3090_1; B:4090_1. Not bitwise hardware-matched; tiny differences are not dropout causal proof.