R93/B | Fixed target36 vs R89B

Same unique VAL checkpoint; native/FP64 winner; no reselection.

Split	Winners	Correct	Rescued	Lost	New false picks	Recall (%)
TRAIN	83279	1234	763	553	942	1.4818
VAL	4625	64	39	27	61	1.3838
TEST	4598	57	37	25	55	1.2397

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	-0.257064	22896	26644	0.001464474	0.014714
VAL	-0.138272	1303	1415	0.000758777	0.007243
TEST	-0.113580	1305	1397	0.001164677	0.012318

Historical R89B main / R84A secondary; not newly matched original-loss controls.

TRAIN CE = raw utility only. Eval macro_ce = original maximin diagnostic, not the optimized loss.

No comparison/risk/R-Drop; WD=1e-4 retained. Zero-count movement alone is not success.

A: released3090_1; B:4090_1. Not bitwise hardware-matched; tiny differences are not dropout causal proof.

R93/B | Fixed target36 vs R84A

Same unique VAL checkpoint; native/FP64 winner; no reselection.

Split	Winners	Correct	Rescued	Lost	New false picks	Recall (%)
TRAIN	83279	1234	1002	65	1034	1.4818
VAL	4625	64	55	2	68	1.3838
TEST	4598	57	48	4	65	1.2397

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	0.190192	25039	22266	0.000380491	0.003149
VAL	0.024691	1327	1307	0.000441584	0.004425
TEST	0.070370	1340	1283	0.000792085	0.007118

Historical R89B main / R84A secondary; not newly matched original-loss controls.

TRAIN CE = raw utility only. Eval macro_ce = original maximin diagnostic, not the optimized loss.

No comparison/risk/R-Drop; WD=1e-4 retained. Zero-count movement alone is not success.

A: released3090_1; B:4090_1. Not bitwise hardware-matched; tiny differences are not dropout causal proof.