R94/B | Fixed target36 vs R89B

Native FP64 winners; same unique full18 VAL checkpoint on TRAIN/VAL/TEST.

Split	Winners	Correct	Rescued	Lost	New false	Recall (%)
TRAIN	83279	2730	2031	325	1294	3.2781
VAL	4625	126	92	18	97	2.7243
TEST	4598	117	86	14	94	2.5446

Split	Top1 delta pp	Corrected	Harmed	Net cost delta	Regret delta pp
TRAIN	0.256722	32713	28970	-0.001786039	-0.015835
VAL	0.180247	1805	1659	-0.000015320	-0.001597
TEST	0.158025	1756	1628	-0.001236847	-0.010531

R89B main / R84A secondary historical; A is the matched random-donor control.

Original R77 CE .35 / comparison .35 / risk .02 / R-Drop .35; dropout .1.

Early trainable inter-node fusion, not R50 pool-first readout. Improvement remains unproven.

16-task donor input coverage, not16-task pretraining; TSP/ATSP raw paths unchanged.

External donor pretraining budget unmatched. No decoder/routes/new cost labels/MTL.

R94/B | Fixed target36 vs R84A

Native FP64 winners; same unique full18 VAL checkpoint on TRAIN/VAL/TEST.

Split	Winners	Correct	Rescued	Lost	New false	Recall (%)
TRAIN	83279	2730	2495	62	1438	3.2781
VAL	4625	126	119	4	110	2.7243
TEST	4598	117	106	2	109	2.5446

Split	Top1 delta pp	Corrected	Harmed	Net cost delta	Regret delta pp
TRAIN	0.703978	37277	27013	-0.002870023	-0.027399
VAL	0.343210	1944	1666	-0.000332514	-0.004415
TEST	0.341975	1913	1636	-0.001609439	-0.015731

R89B main / R84A secondary historical; A is the matched random-donor control.

Original R77 CE .35 / comparison .35 / risk .02 / R-Drop .35; dropout .1.

Early trainable inter-node fusion, not R50 pool-first readout. Improvement remains unproven.

16-task donor input coverage, not16-task pretraining; TSP/ATSP raw paths unchanged.

External donor pretraining budget unmatched. No decoder/routes/new cost labels/MTL.

R94/B | B Inference Timing

First128 locked TEST inputs/task; synchronized GPU wall time, three repetitions; saved evidence only.

Problem	Batch	Donor + selector ms / instance	Cache-hit selector ms / instance
TSP	128	0.262168	0.261500
CVRP	128	0.317999	0.317428
ATSP	128	0.064552	0.064129
OVRP	128	0.068539	0.068381
VRPB	128	0.068998	0.068472
VRPL	128	0.068735	0.068517
VRPTW	128	0.068802	0.068693
OVRPTW	128	0.068672	0.068616
OVRPB	128	0.068922	0.068748
OVRPL	128	0.068784	0.068531
VRPBL	128	0.069443	0.068753
VRPBTW	128	0.068777	0.068460
VRPLTW	128	0.068952	0.068537
OVRPBL	128	0.068727	0.067991
OVRPBTW	128	0.068671	0.068324
OVRPLTW	128	0.068635	0.068589
VRPBLTW	128	0.068858	0.068624
OVRPBLTW	128	0.068739	0.068639

Inputs already on GPU. Excludes dataset I/O and one-time feature-cache preparation.

Cached-feature selector latency is NOT full-system latency; donor+selector is not I/O-inclusive E2E.

TSP/ATSP have unchanged raw paths and no donor; external pretraining budget is unmatched.

Device: cuda:0