R95 | Adaptive corrected sampling B versus uniform A

Original R88B/R77 full18 model and CE/comparison/risk/two-pass R-Drop retained.

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	0.080864	16508	15329	0.000150971	0.000978
VAL	0.098765	914	834	0.000175161	0.001494
TEST	0.004938	882	878	0.000346255	0.002279

Arm / VAL ref	Top1 delta pp	Cost delta	Regret delta pp	Pass
R95/A / R89B	-0.248148	0.000366210	0.003559	False
R95/A / R84A	-0.085185	0.000049017	0.000741	False
R95/B / R89B	-0.149383	0.000541371	0.005054	False
R95/B / R84A	0.013580	0.000224178	0.002236	False

Original R77 CE .35 / comparison .35 / risk .02 / R-Drop ramp .35; dropout .1 in both arms.

A:3090_0; B:4090_0. Not bitwise hardware-matched. No causal claim or guaranteed variance reduction.

Same natural first7200; then A uniform IID / B own-TRAIN adaptive q, .5 uniform mix, whole-row 1/(N*q).

R89B main / R84A secondary historical; no TEST/LIB checkpoint or inference-rule selection.

R95 | Fixed target36 genuine recovery and false picks

B versus A on the unchanged historical36 targets; native FP64 winners.

Split	Winners	Correct	Rescued	Lost	New false picks	Recall (%)
TRAIN	83279	420	157	48	123	0.5043
VAL	4625	19	9	5	8	0.4108
TEST	4598	19	8	4	7	0.4132

Detailed per-task recovery, false-pick costs and regret remain in the paired/target CSVs.

Changing zero-selection counts alone is not success. Full18 cost, regret and native Top1 all matter.

No causal claim; the independent arms run on different GPU types.

R95 | Own-TRAIN sampling proposal diagnostics

Summaries only; never full-network/stochastic gradient variance or unbiased optimizer updates.

Arm	Update	Used for draws	Task-equal proposal ESS	Max inverse weight	Proxy second-moment ratio
A	7200	False	67136.65	1.811189	0.620407
A	14400	False	64828.11	1.915344	0.583729
A	21600	False	64402.05	1.946027	0.575919
A	28800	False	64604.01	1.952767	0.577758
A	36000	False	65297.91	1.960115	0.588203
A	43200	False	65831.04	1.970443	0.597519
A	50400	False	65322.28	1.970190	0.590160
A	57600	False	65352.96	1.968946	0.590439
B	7200	True	65527.87	1.827204	0.592939
B	14400	True	65842.35	1.901765	0.601639
B	21600	True	63418.56	1.938534	0.559996
B	28800	True	64979.52	1.959383	0.583235
B	36000	True	65012.93	1.970140	0.585322
B	43200	True	65292.37	1.970784	0.588036
B	50400	True	65224.79	1.972470	0.588043
B	57600	False	65375.26	1.972289	0.591274

A proposal summaries are diagnostic only; A draws stay uniform. Final57600 proposal has no next interval.

Final-affine CE+comparison+risk eval proxy; KL=0 in eval, not full-parameter norm.

Whole-row correction, arithmetic mean, no self-normalization. No guaranteed variance reduction or causal claim.