R96 | Where do completed models still fail?

21 fixed models | 18 tasks x 4,500 VAL rows | Reference R89B Model Top1=71.25%

Model	Top1 (%)	Regret (%)	Rescued	Harmed
R84A	71.085	0.40349	1,332	1,464
R84B	71.000	0.40500	1,300	1,501
R84C	71.023	0.40830	1,262	1,444
R84D	70.954	0.40810	1,384	1,622
R85A	70.968	0.40690	1,523	1,750
R85B	70.936	0.40587	1,496	1,749
R85C	71.009	0.39797	1,366	1,560
R85D	71.127	0.40338	1,198	1,296
R87A	71.148	0.40635	1,224	1,305
R87B	71.079	0.40206	1,160	1,297
R88A	71.058	0.40544	1,301	1,455
R88B	71.058	0.40513	1,223	1,377
R89A	71.051	0.40568	1,332	1,492
R89B	71.248	0.40068	0	0
R91A	71.073	0.40598	1,324	1,466
R91B	71.091	0.40837	1,184	1,311
R92B	69.558	0.43758	2,904	4,273
R92C	71.286	0.40324	1,616	1,585
R92D	69.425	0.43066	2,747	4,224
R93A	71.126	0.40787	1,202	1,301
R93B	71.110	0.40792	1,303	1,415

Common misses across all21
17,332

Share of reference errors
74.42%

Share of reference regret
75.90%

Fixed36 winner cases
4,625
Reference correct: 52
Any model correct: 146 (3.16%)
All models miss: 4,479 (96.84%)

VAL only / native FP64 winners / frozen NEW90k / no new fit or inference rule.
Coverage uses known answers across saved models: NOT deployable performance or an accuracy ceiling.

Common failure is measured per task, not just ALL Top1

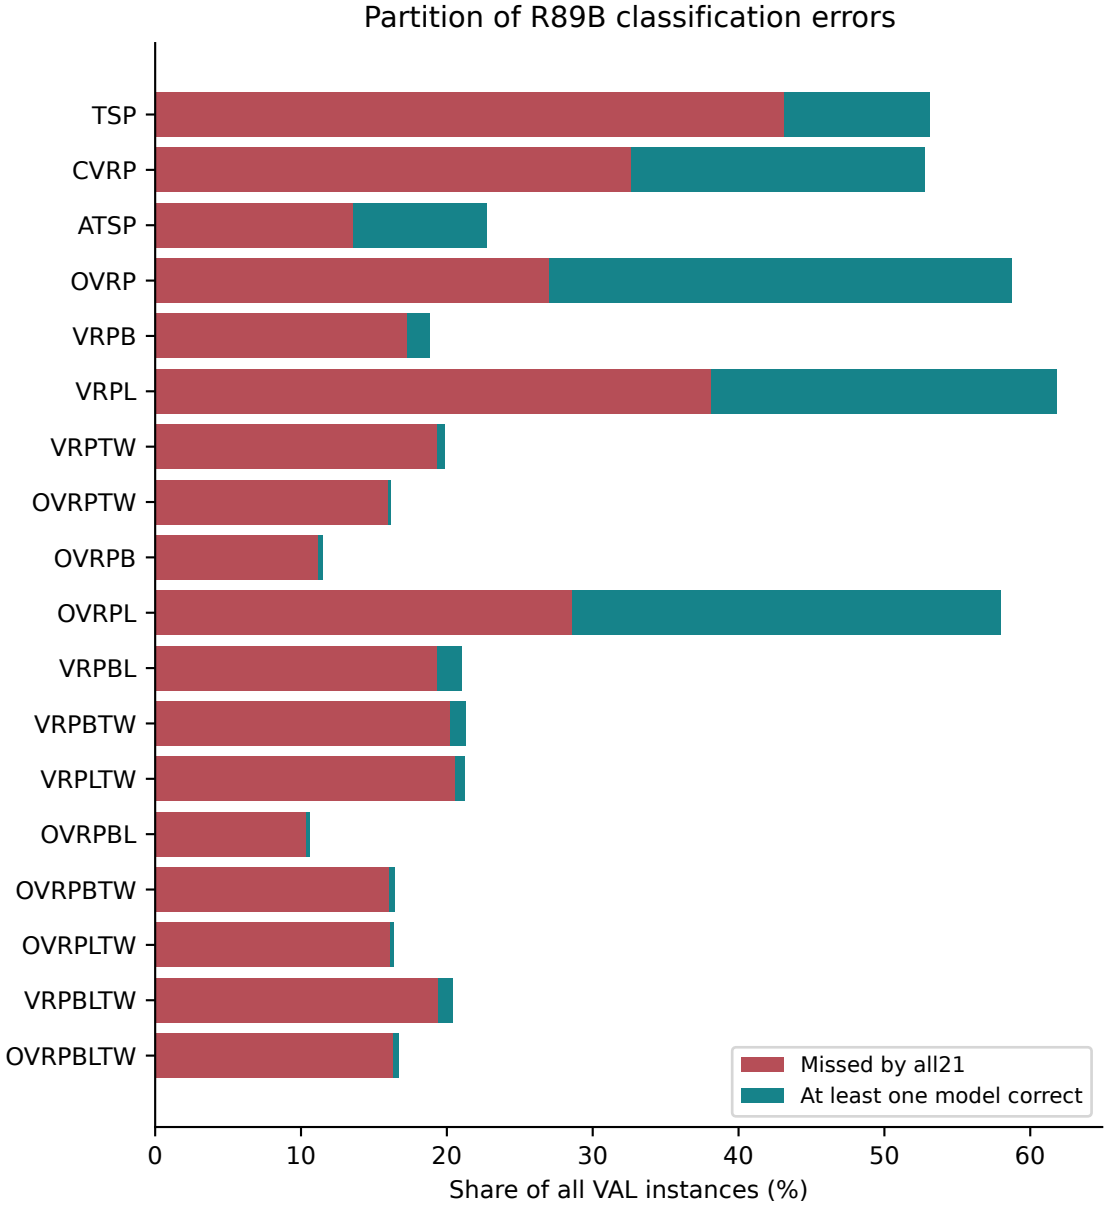

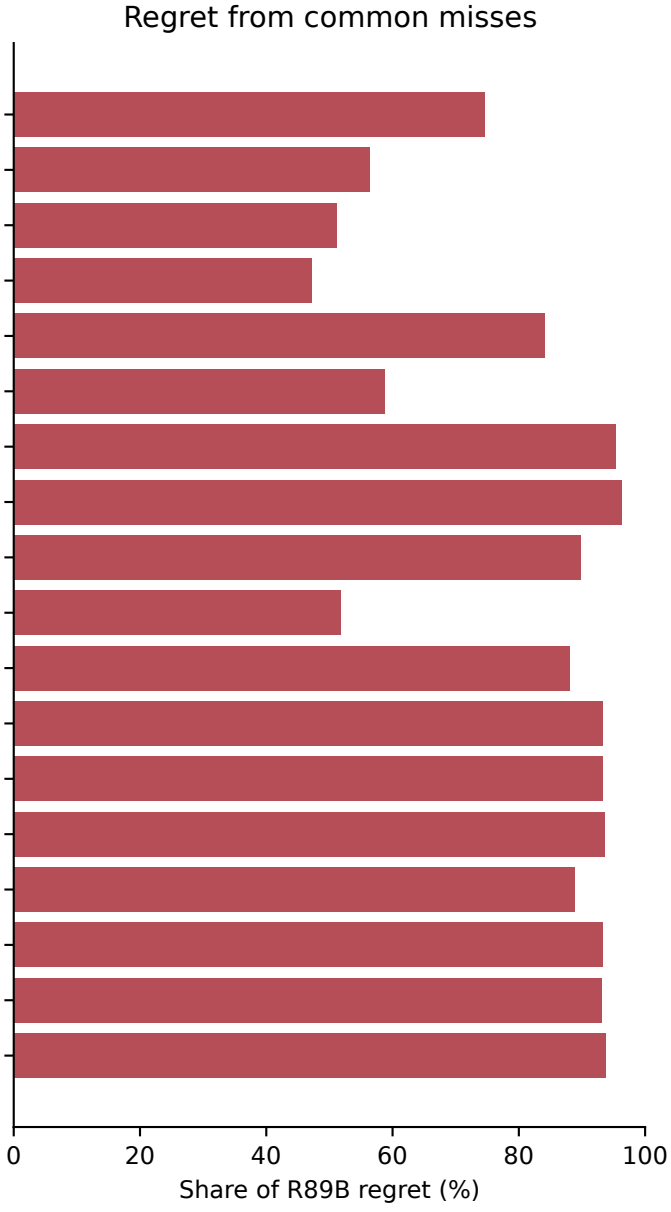

VAL only / native FP64 winners / frozen NEW90k / no new fit or inference rule.
Coverage uses known answers across saved models: NOT deployable performance or an accuracy ceiling.

Fixed36 missed-method groups: coverage is not a new selector

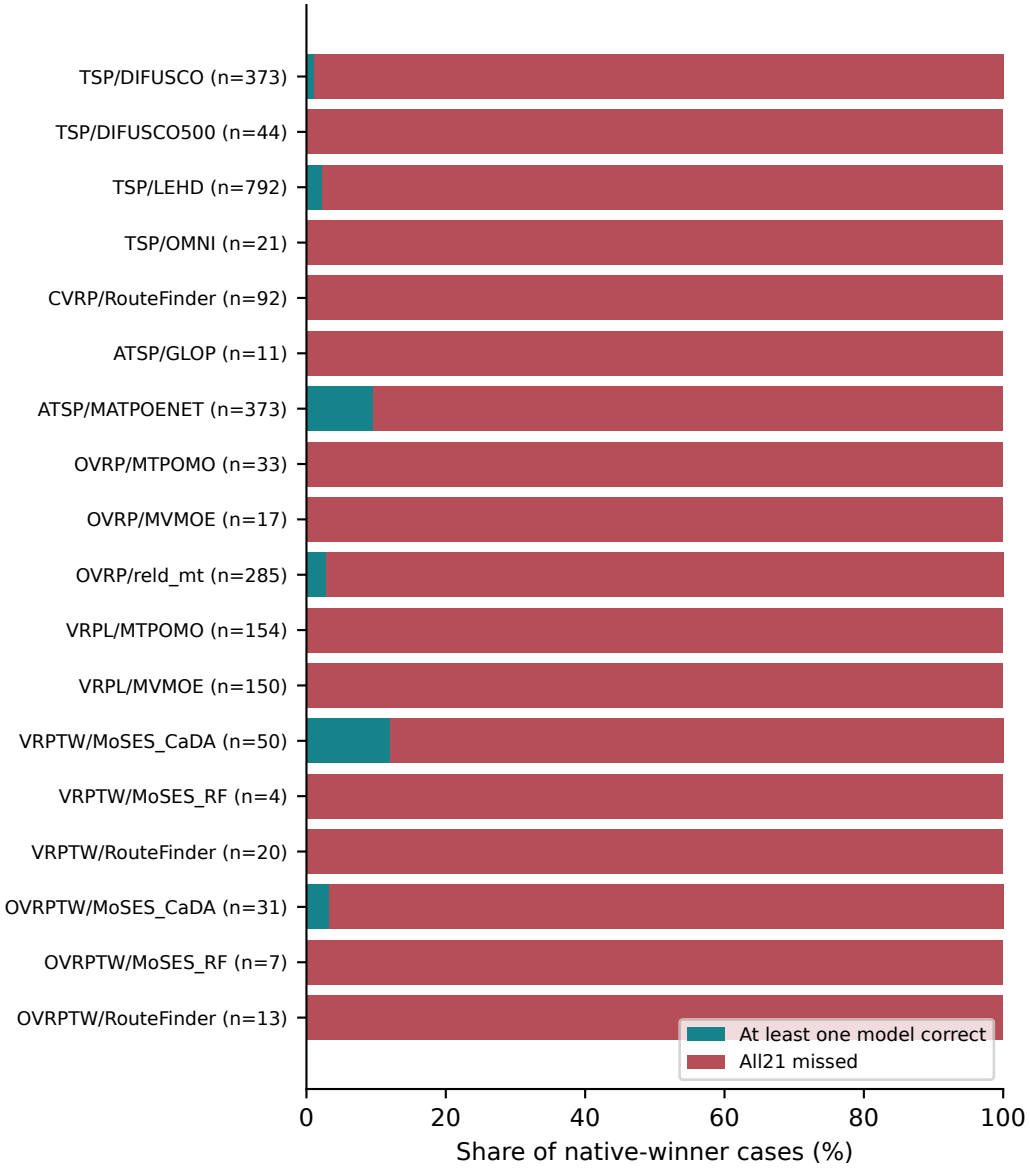

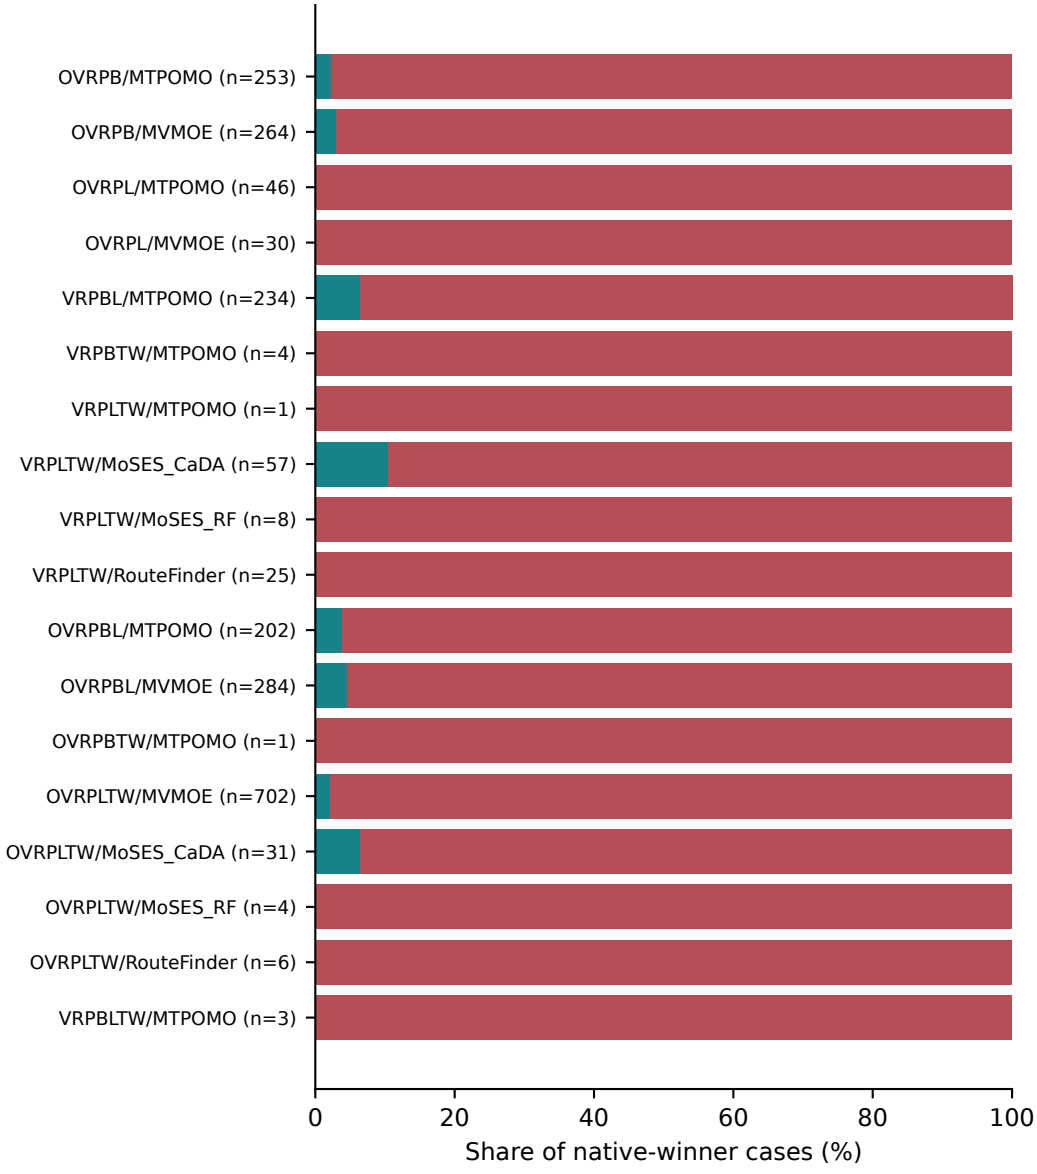

VAL only / native FP64 winners / frozen NEW90k / no new fit or inference rule.
Coverage uses known answers across saved models: NOT deployable performance or an accuracy ceiling.
Weak models can add coverage with many false picks; see per-model recall/false-pick CSV.