R100 | Pairwise Context And Evidence Boundaries

One predeclared full18 R100 pairwise-binary frozen TabICL bank. No SGD, fine-tuning, epochs, TRAIN accuracy or VAL aggregation search.

For every legal pair: ALL historical TRAIN rows in natural frequency/order, including third-method full-pool winners. Only exact native FP64 pair ties are excluded.

Binary labels: 0 first method wins, 1 second wins. Pair orientation follows ascending stable global IDs; no FP32 reconstruction or full-winner filtering.

Constant shortcut ONLY for exactly one strict class, or no strict rows (.5 neutral). Support/tie/class counts are identity-bound and reported.

P[i,j] predicts i beating j conditional on a strict TRAIN relation; P[j,i]=1-P[i,j]. These are NOT full-pool winner probabilities.

Decision: argmax mean opponent win probability, fixed expected-rank/Borda; ascending global-ID exact ties. No coupling, calibration, heuristic threshold or cost-sensitive claim.

Formal log-softmax(mean-win) is for descriptive CE/report interoperability only, not calibrated probability or the ranking rule.

Each nonconstant pair independently fits its strict TRAIN context and predicts the whole named dataset; only one pair classifier is resident. Runtime scales with K*(K-1)/2 contexts.

Original R97 public classifier weights/runtime: max_classes=10 is sufficient for binary2. n_estimators=1, norm=none, batch_size=1, seed2, AMP, kv_cache=False, GPU-only, Flash-only.

Immutable R97 input-only statistics and TRAIN-constant drop. Query labels/winners/costs never enter fit/forward; any query base ID in complete TRAIN is rejected BEFORE pair filtering.

SBS fixed on R84 full81000 TRAIN. Oracle is the per-instance legal-pool minimum. ALL is equal-problem macro; mixed raw costs descriptive only. vs_SBS/vs_Oracle are percent ratios of mean costs; actual regret is mean per-instance percent regret. Native/FP64 ind is never replaced by FP32 argmin.

Same current bank on VAL, full TEST and native LIB. Historical TEST members informed research, not wholly blind. Public pretraining compute is unmatched.

NSS reuses saved R79 historical70114 TRAIN weights, not a new81000 retrain. Missing stays missing; incompatible saved evidence fails.

R89B main, R84A secondary; fixed36 targets frozen from R84A VAL. Report rescued/lost true winners, false picks, raw false-pick cost and regret, not zero-count changes alone.

Per-pair fit/predict times are NESTED in enclosing calls and are not added twice. VAL/TEST use cached inputs: fresh raw-graph statistics/cache-build latency is NOT measured.

LIB native_input_to_statistics includes native loading/geometry/statistics without query-label loading. Recorded sums are not full end-to-end latency.

R100 | VAL | ALL And Full18

Task	n	Top1 %	Top2 %	Top3 %	mean_cost	SBS cost	Oracle	vs_SBS %	vs_Oracle %	Actual regret %
ALL	81000	71.070370	90.482716	96.033333	11.416165	11.427978	11.371692	-0.201078	0.386414	0.394268
TSP	4500	45.600000	70.888889	86.511111	8.023986	8.045354	7.990533	-0.265594	0.418656	0.409313
CVRP	4500	46.777778	67.155556	78.222222	18.147534	18.269377	18.064087	-0.666926	0.461949	0.595322
ATSP	4500	77.911111	95.355556	98.933333	1.637724	1.676905	1.632471	-2.336494	0.321801	0.319152
OVRP	4500	42.044444	73.266667	92.533333	7.459348	7.469579	7.415777	-0.136966	0.587554	0.597930
VRPB	4500	80.866667	97.244444	100.000000	9.354679	9.354876	9.318750	-0.002104	0.385563	0.394564
VRPL	4500	37.533333	65.311111	82.422222	12.429751	12.437959	12.383444	-0.065994	0.373944	0.380305
VRPTW	4500	79.955556	98.422222	99.288889	18.819473	18.819431	18.766139	0.000223	0.284205	0.301462
OVRPTW	4500	83.777778	99.222222	99.755556	10.143730	10.143730	10.116859	0.000000	0.265602	0.262806
OVRPB	4500	88.511111	96.844444	100.000000	6.357254	6.357254	6.341078	0.000000	0.255104	0.265013
OVRPL	4500	41.688889	73.422222	91.822222	7.436408	7.445709	7.390416	-0.124924	0.622322	0.639050
VRPBL	4500	78.555556	96.933333	100.000000	9.370291	9.371523	9.328366	-0.013138	0.449437	0.452938
VRPBTW	4500	78.422222	99.844444	100.000000	19.950456	19.950430	19.851084	0.000128	0.500588	0.477917
VRPLTW	4500	78.755556	98.422222	99.422222	18.765968	18.765871	18.707843	0.000518	0.310703	0.326395
OVRPBL	4500	89.222222	97.177778	100.000000	6.369942	6.370093	6.352824	-0.002370	0.269460	0.280748
OVRPBTW	4500	83.466667	100.000000	100.000000	10.660992	10.661098	10.623245	-0.001001	0.355322	0.329802
OVRPLTW	4500	83.444444	99.266667	99.688889	10.157471	10.157423	10.129818	0.000474	0.272986	0.276925
VRPBLTW	4500	79.444444	99.911111	100.000000	19.796009	19.797044	19.706420	-0.005230	0.454615	0.443949
OVRPBLTW	4500	83.288889	100.000000	100.000000	10.609955	10.609955	10.571302	0.000000	0.365644	0.343234

Native/FP64 ind; fixed expected-rank/Borda, no calibration/aggregation search.

SBS: R84 full81000 TRAIN. ALL equal-problem macro; mixed raw cost descriptive.

vs_* ratios of means differ from mean per-instance actual regret.

R100 | TEST | ALL And Full18

Task	n	Top1 %	Top2 %	Top3 %	mean_cost	SBS cost	Oracle	vs_SBS %	vs_Oracle %	Actual regret %
ALL	81000	71.316049	90.483951	95.969136	11.373466	11.385299	11.329932	-0.200894	0.382209	0.389043
TSP	4500	46.444444	70.088889	86.200000	7.970024	7.989621	7.937593	-0.245278	0.408573	0.402728
CVRP	4500	47.577778	68.644444	78.977778	18.000121	18.124983	17.918509	-0.688899	0.455461	0.590110
ATSP	4500	77.466667	95.733333	98.911111	1.635020	1.674207	1.629441	-2.340671	0.342334	0.340713
OVRP	4500	41.200000	72.266667	91.911111	7.445250	7.455750	7.400487	-0.140823	0.604865	0.614165
VRPB	4500	80.511111	97.133333	100.000000	9.350952	9.351810	9.315298	-0.009174	0.382745	0.386606
VRPL	4500	37.622222	64.666667	81.511111	12.347549	12.356138	12.300958	-0.069511	0.378757	0.386561
VRPTW	4500	80.866667	98.844444	99.377778	18.690730	18.690545	18.637408	0.000987	0.286100	0.297267
OVRPTW	4500	83.933333	99.288889	99.555556	10.093063	10.093063	10.068135	0.000000	0.247590	0.238522
OVRPB	4500	88.133333	96.933333	100.000000	6.370826	6.370826	6.352225	0.000000	0.292814	0.301388
OVRPL	4500	41.955556	74.000000	91.800000	7.463368	7.472542	7.418033	-0.122762	0.611145	0.621077
VRPBL	4500	79.533333	96.933333	100.000000	9.397090	9.396379	9.355967	0.007569	0.439541	0.443116
VRPBTW	4500	78.800000	99.800000	100.000000	19.665560	19.665725	19.577666	-0.000839	0.448952	0.446033
VRPLTW	4500	79.266667	98.688889	99.444444	18.803004	18.803136	18.745143	-0.000702	0.308673	0.322482
OVRPBL	4500	89.288889	96.600000	100.000000	6.375741	6.375741	6.359310	0.000000	0.258375	0.269531
OVRPBTW	4500	84.822222	100.000000	100.000000	10.659591	10.659591	10.624748	0.000000	0.327942	0.299912
OVRPLTW	4500	82.911111	99.222222	99.755556	10.085766	10.086112	10.059001	-0.003422	0.266088	0.266661
VRPBLTW	4500	79.755556	99.866667	100.000000	19.706624	19.707068	19.613284	-0.002254	0.475901	0.454619
OVRPBLTW	4500	83.600000	100.000000	100.000000	10.662116	10.662149	10.625574	-0.000306	0.343910	0.321275

Native/FP64 ind; fixed expected-rank/Borda, no calibration/aggregation search.

SBS: R84 full81000 TRAIN. ALL equal-problem macro; mixed raw cost descriptive.

vs_* ratios of means differ from mean per-instance actual regret.

R100 | Native LIB Metrics

Dataset	n	Top1 %	Top2 %	Top3 %	mean_cost	SBS cost	Oracle	vs_SBS %	vs_Oracle %	Actual regret %
TSPLIB	49	20.408163	38.775510	67.346939	8.224506	8.240435	8.142167	-0.193312	1.011260	1.037809
CVRPLIB	100	59.000000	81.000000	90.000000	67.423326	67.245364	66.945093	0.264646	0.714366	0.594214

Same unique context/weight bank as full18 VAL/TEST; original native LIB, no per-dataset tuning.

R100 | NSS Four-Dataset Historical Comparison

Dataset	Status	n	NSS Top1 %	R100 Top1 %	NSS cost	R100 cost
TSP TEST	complete	4500	45.888889	46.444444	7.970401	7.970024
CVRP TEST	complete	4500	45.977778	47.577778	18.004802	18.000121
TSPLIB	complete	49	20.408163	20.408163	8.217907	8.224506
CVRPLIB	complete	100	70.000000	59.000000	67.101112	67.423326

NSS: saved R79 historical70114 TRAIN weights, not a new81000 retrain. Current R100 bank fixed across all four.

Native rows/costs/winner/pool aligned; CSV records versions, pools, checkpoints and missing evidence.

R100 | Full18 TEST vs R89B

Status	Rescued	Harmed	Net correct	Top1 delta pp	Cost delta	Regret delta pp
complete	1693	1777	-84	-0.103704	-0.000004	-0.000603

R89B main, R84A secondary, historical and not compute matched. Positive cost/regret delta is worse.

Fixed36 status: complete; targets frozen from R84A VAL, not R100 TEST.

R100 | Fixed36 TEST Recovery vs R89B

Task / Method	Winners	Rescued	Lost	Net	New false	False cost delta	Regret delta pp
TSP/DIFUSCO	334	0	0	0	0	0.000000	-0.002806
TSP/DIFUSCO500	42	0	0	0	0	0.000000	0.001447
TSP/LEHD	762	0	4	-4	0	0.000000	0.003498
TSP/OMNI	28	0	0	0	0	0.000000	-0.000217
CVRP/RouteFinder	67	0	0	0	1	0.000012	-0.003156
ATSP/GLOP	11	0	0	0	0	0.000000	-0.002683
ATSP/MATPOENET	370	26	0	26	30	0.000201	-0.021371
OVRP/MTPOMO	44	0	0	0	0	0.000000	-0.000870
OVRP/MVMOE	24	0	0	0	0	0.000000	-0.000401
OVRP/reld_mt	310	0	6	-6	0	0.000000	-0.002359
VRPL/MTPOMO	149	0	0	0	0	0.000000	0.000945
VRPL/MVMOE	155	0	0	0	0	0.000000	-0.000633
VRPTW/MoSES_CaDA	42	1	1	0	0	0.000000	0.000162
VRPTW/MoSES_RF	5	0	0	0	0	0.000000	0.002330
VRPTW/RouteFinder	13	0	0	0	0	0.000000	0.000000
OVRPTW/MoSES_CaDA	26	0	1	-1	0	0.000000	0.002423
OVRPTW/MoSES_RF	6	0	0	0	0	0.000000	0.000000
OVRPTW/RouteFinder	10	0	0	0	0	0.000000	0.000000

New false-pick cost delta is relative to reference selections, normalized by whole task n.

CSV also includes false raw cost sums/contributions, precision, recall and new-false regret.

R100 | Fixed36 TEST Recovery vs R89B

Task / Method	Winners	Rescued	Lost	Net	New false	False cost delta	Regret delta pp
OVRPB/MTPOMO	256	0	6	-6	0	0.000000	0.010573
OVRPB/MVMOE	278	0	6	-6	0	0.000000	0.017817
OVRPL/MTPOMO	58	0	0	0	0	0.000000	-0.002261
OVRPL/MVMOE	34	0	0	0	0	0.000000	-0.000154
VRPBL/MTPOMO	229	0	2	-2	0	0.000000	0.019181
VRPBTW/MTPOMO	8	0	0	0	0	0.000000	0.000000
VRPLTW/MTPOMO	0	0	0	0	0	0.000000	0.000000
VRPLTW/MoSES_CaDA	53	0	1	-1	2	0.000050	0.001001
VRPLTW/MoSES_RF	9	0	0	0	0	0.000000	0.000000
VRPLTW/RouteFinder	19	0	0	0	0	0.000000	-0.000058
OVRPBL/MTPOMO	219	0	2	-2	0	0.000000	0.004394
OVRPBL/MVMOE	263	0	5	-5	0	0.000000	0.004660
OVRPBTW/MTPOMO	0	0	0	0	0	0.000000	0.000000
OVRPLTW/MVMOE	725	0	7	-7	0	0.000000	0.007261
OVRPLTW/MoSES_CaDA	30	0	1	-1	0	0.000000	0.000388
OVRPLTW/MoSES_RF	5	0	0	0	0	0.000000	0.000000
OVRPLTW/RouteFinder	11	0	0	0	0	0.000000	0.000000
VRPBLTW/MTPOMO	3	0	0	0	0	0.000000	0.000000

New false-pick cost delta is relative to reference selections, normalized by whole task n.

CSV also includes false raw cost sums/contributions, precision, recall and new-false regret.

R100 | Full18 TEST vs R84A

Status	Rescued	Harmed	Net correct	Top1 delta pp	Cost delta	Regret delta pp
complete	1636	1571	65	0.080247	-0.000377	-0.005804

R89B main, R84A secondary, historical and not compute matched. Positive cost/regret delta is worse.

Fixed36 status: complete; targets frozen from R84A VAL, not R100 TEST.

R100 | Fixed36 TEST Recovery vs R84A

Task / Method	Winners	Rescued	Lost	Net	New false	False cost delta	Regret delta pp
TSP/DIFUSCO	334	0	0	0	0	0.000000	-0.001260
TSP/DIFUSCO500	42	0	0	0	0	0.000000	0.001752
TSP/LEHD	762	0	0	0	0	0.000000	0.001664
TSP/OMNI	28	0	0	0	0	0.000000	-0.000138
CVRP/RouteFinder	67	0	0	0	1	0.000003	-0.003219
ATSP/GLOP	11	0	0	0	0	0.000000	-0.002683
ATSP/MATPOENET	370	26	0	26	30	0.000201	-0.024398
OVRP/MTPOMO	44	0	0	0	0	0.000000	-0.001110
OVRP/MVMOE	24	0	0	0	0	0.000000	-0.000313
OVRP/reld_mt	310	0	1	-1	0	0.000000	-0.001431
VRPL/MTPOMO	149	0	0	0	0	0.000000	0.000884
VRPL/MVMOE	155	0	0	0	0	0.000000	-0.000710
VRPTW/MoSES_CaDA	42	1	1	0	0	0.000000	0.000162
VRPTW/MoSES_RF	5	0	0	0	0	0.000000	0.000000
VRPTW/RouteFinder	13	0	0	0	0	0.000000	0.000000
OVRPTW/MoSES_CaDA	26	0	0	0	0	0.000000	0.000000
OVRPTW/MoSES_RF	6	0	0	0	0	0.000000	0.000000
OVRPTW/RouteFinder	10	0	0	0	0	0.000000	0.000000

New false-pick cost delta is relative to reference selections, normalized by whole task n.

CSV also includes false raw cost sums/contributions, precision, recall and new-false regret.

R100 | Fixed36 TEST Recovery vs R84A

Task / Method	Winners	Rescued	Lost	Net	New false	False cost delta	Regret delta pp
OVRPB/MTPOMO	256	0	1	-1	0	0.000000	0.000053
OVRPB/MVMOE	278	0	1	-1	0	0.000000	0.003196
OVRPL/MTPOMO	58	0	0	0	0	0.000000	-0.001641
OVRPL/MVMOE	34	0	0	0	0	0.000000	-0.000442
VRPBL/MTPOMO	229	0	0	0	0	0.000000	0.004147
VRPBTW/MTPOMO	8	0	0	0	0	0.000000	0.000000
VRPLTW/MTPOMO	0	0	0	0	0	0.000000	0.000000
VRPLTW/MoSES_CaDA	53	1	0	1	2	0.000050	-0.000981
VRPLTW/MoSES_RF	9	0	0	0	0	0.000000	0.000000
VRPLTW/RouteFinder	19	0	0	0	0	0.000000	-0.000058
OVRPBL/MTPOMO	219	0	2	-2	0	0.000000	0.005695
OVRPBL/MVMOE	263	0	2	-2	0	0.000000	0.002967
OVRPBTW/MTPOMO	0	0	0	0	0	0.000000	0.000000
OVRPLTW/MVMOE	725	0	2	-2	0	0.000000	0.001901
OVRPLTW/MoSES_CaDA	30	0	1	-1	0	0.000000	0.000388
OVRPLTW/MoSES_RF	5	0	0	0	0	0.000000	0.000000
OVRPLTW/RouteFinder	11	0	0	0	0	0.000000	0.000000
VRPBLTW/MTPOMO	3	0	0	0	0	0.000000	0.000000

New false-pick cost delta is relative to reference selections, normalized by whole task n.

CSV also includes false raw cost sums/contributions, precision, recall and new-false regret.

R100 | Independent Pair Context Runtime

Task	Methods	Pairs	Fits	Predicts	Load s	Fit s	Predict s	Outer s	Recorded sum s
TSP	8	28	84	84	0.431344	128.354628	290.136677	433.113939	434.106473
CVRP	10	45	135	135	0.575577	220.638755	488.667261	729.455742	730.732926
ATSP	5	10	20	20	0.149122	18.829868	51.485460	73.275023	73.600679
OVRP	6	15	30	30	0.245355	49.261884	110.787235	164.832919	165.255991
VRPB	3	3	6	6	0.197475	10.003180	21.982904	32.891624	33.267400
VRPL	6	15	30	30	0.242053	49.657712	110.738082	165.145273	165.564697
VRPTW	6	15	30	30	0.282720	57.037068	126.547630	188.237058	188.697814
OVRPTW	6	15	30	30	0.280318	56.657979	125.727377	186.820606	187.278493
OVRPB	3	3	6	6	0.197347	9.893837	21.966858	32.751566	33.124315
OVRPL	6	15	30	30	0.242273	49.122750	112.025042	165.543527	165.962703
VRPBL	3	3	6	6	0.198461	9.803840	22.132810	32.828538	33.203838
VRPBTW	3	3	6	6	0.254746	11.442507	25.101547	37.444366	37.880127
VRPLTW	6	15	30	30	0.280858	57.652724	126.359093	188.858322	189.317736
OVRPBL	3	3	6	6	0.198559	9.904051	22.114159	32.912264	33.287874
OVRPBTW	3	3	6	6	0.241445	11.463061	25.529867	37.892894	38.315765
OVRPLTW	6	15	30	30	0.281523	57.448699	126.903854	189.043747	189.504086
VRPBLTW	3	3	6	6	0.237764	12.025260	25.808992	38.814882	39.230174
OVRPBLTW	3	3	6	6	0.233705	11.506159	25.275239	37.692606	38.102964

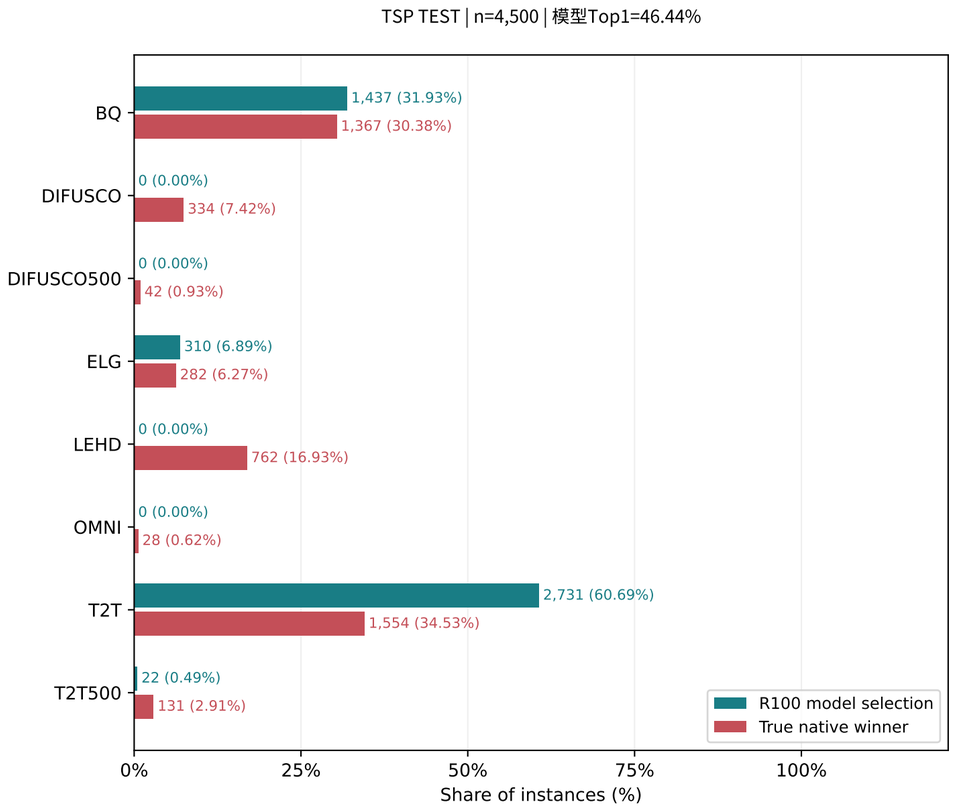
K(K-1)/2 contexts per named dataset; only exact one-class/no-strict-row shortcuts. One classifier resident.

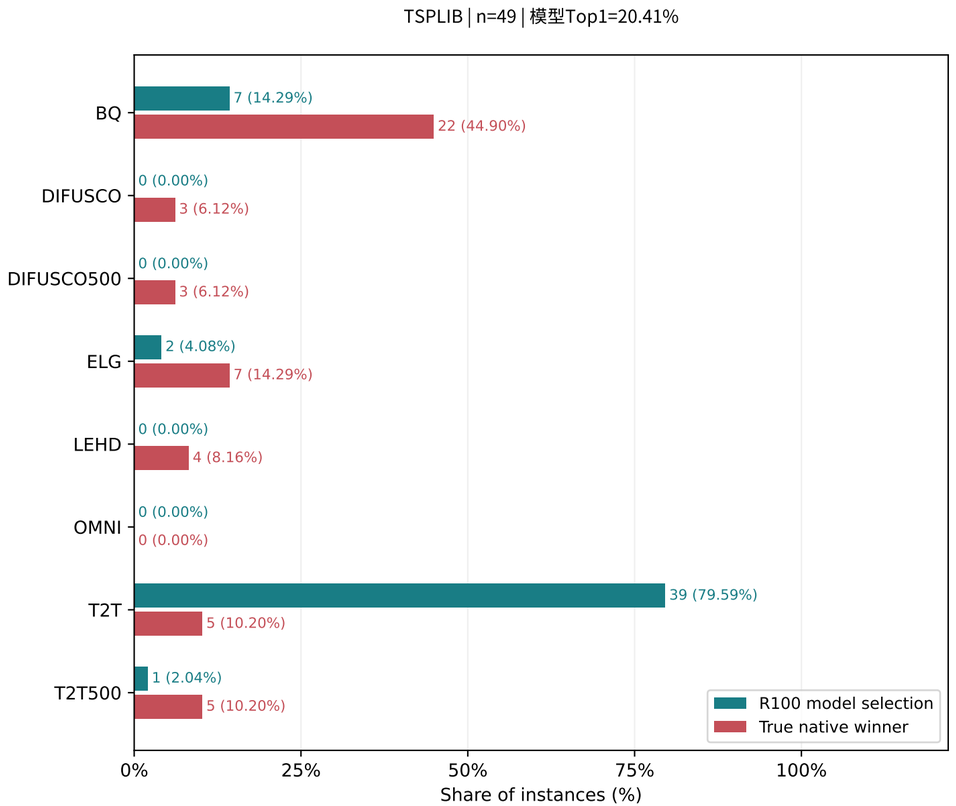
Pair breakdowns are nested in outer calls, not added twice. Cached VAL/TEST exclude fresh graph statistics.

LIB includes native input/geometry/statistics; recorded sum is not full end-to-end runtime.

R100 | TSP | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea



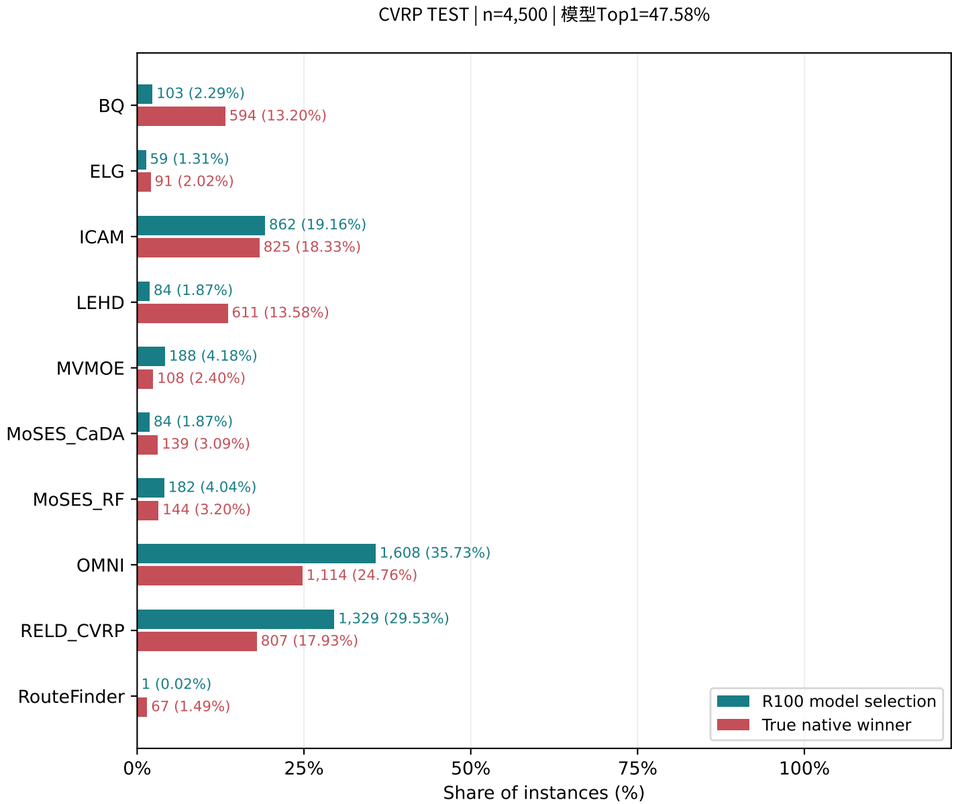


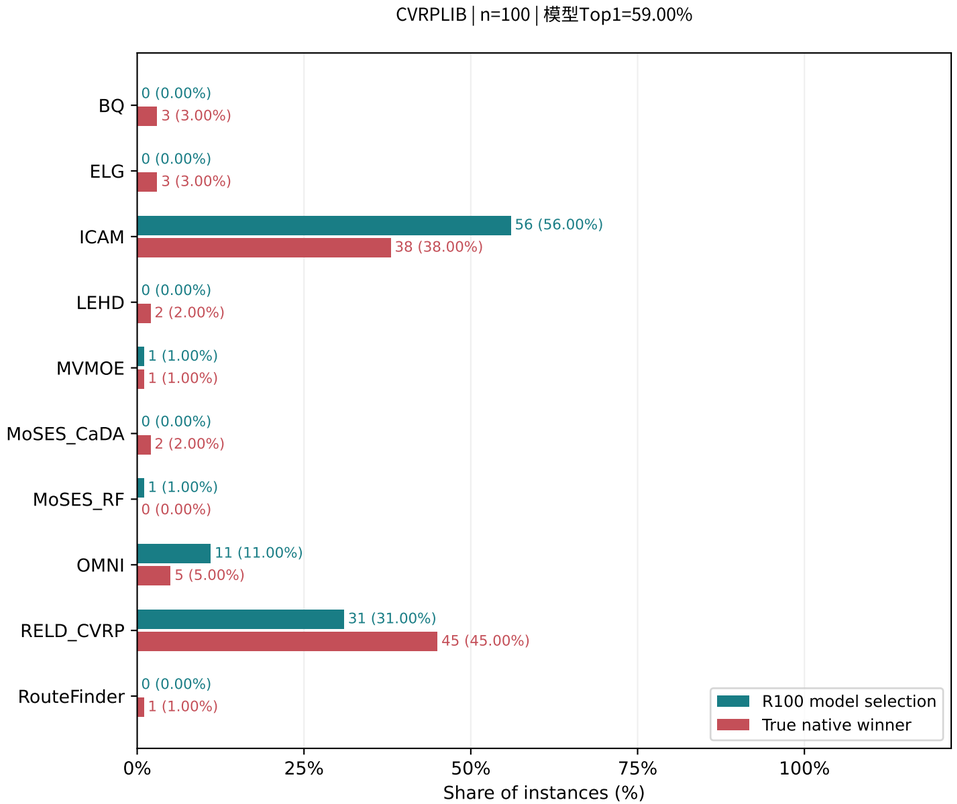
Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | CVRP | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea



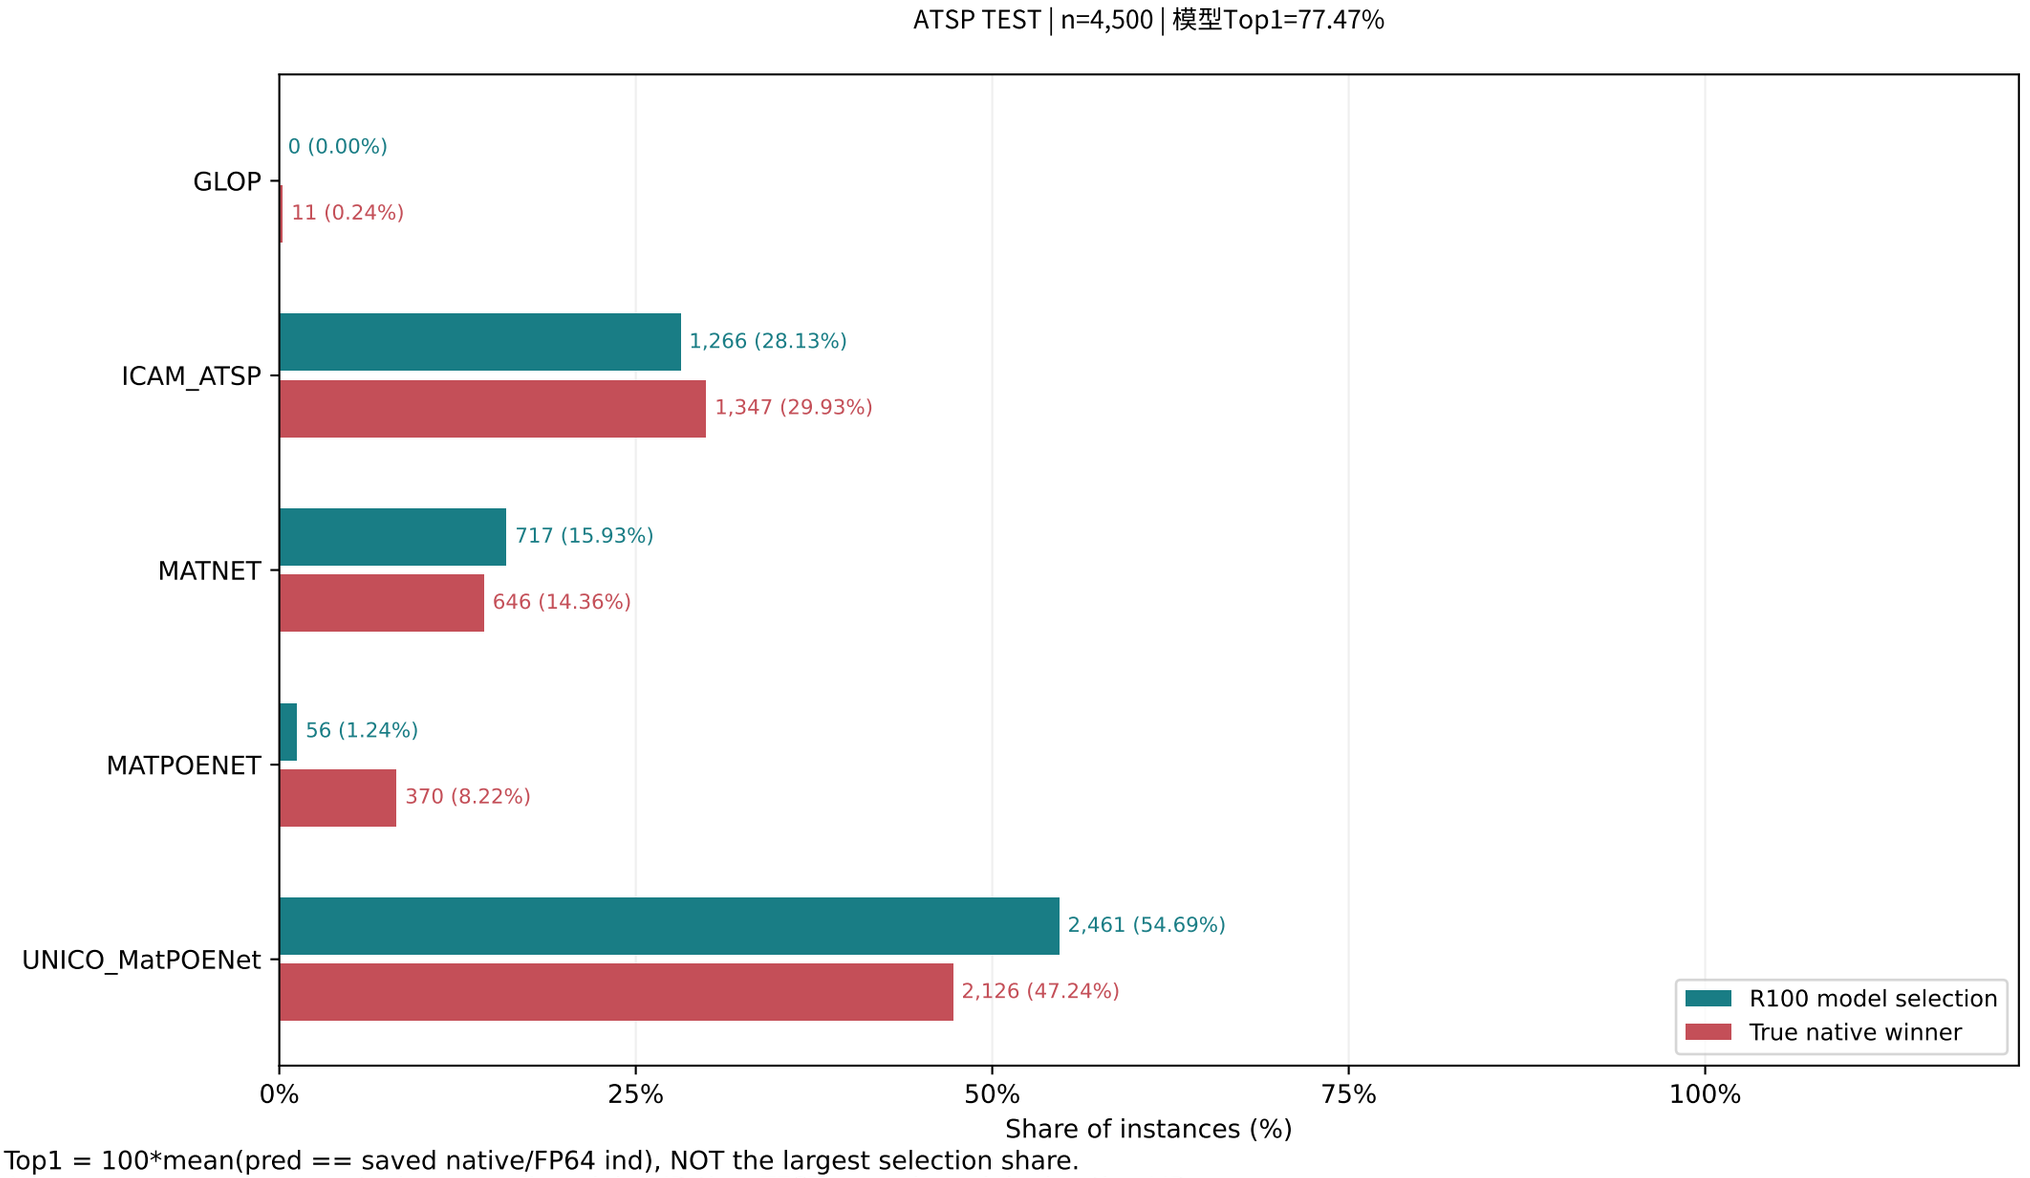


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | ATSP | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

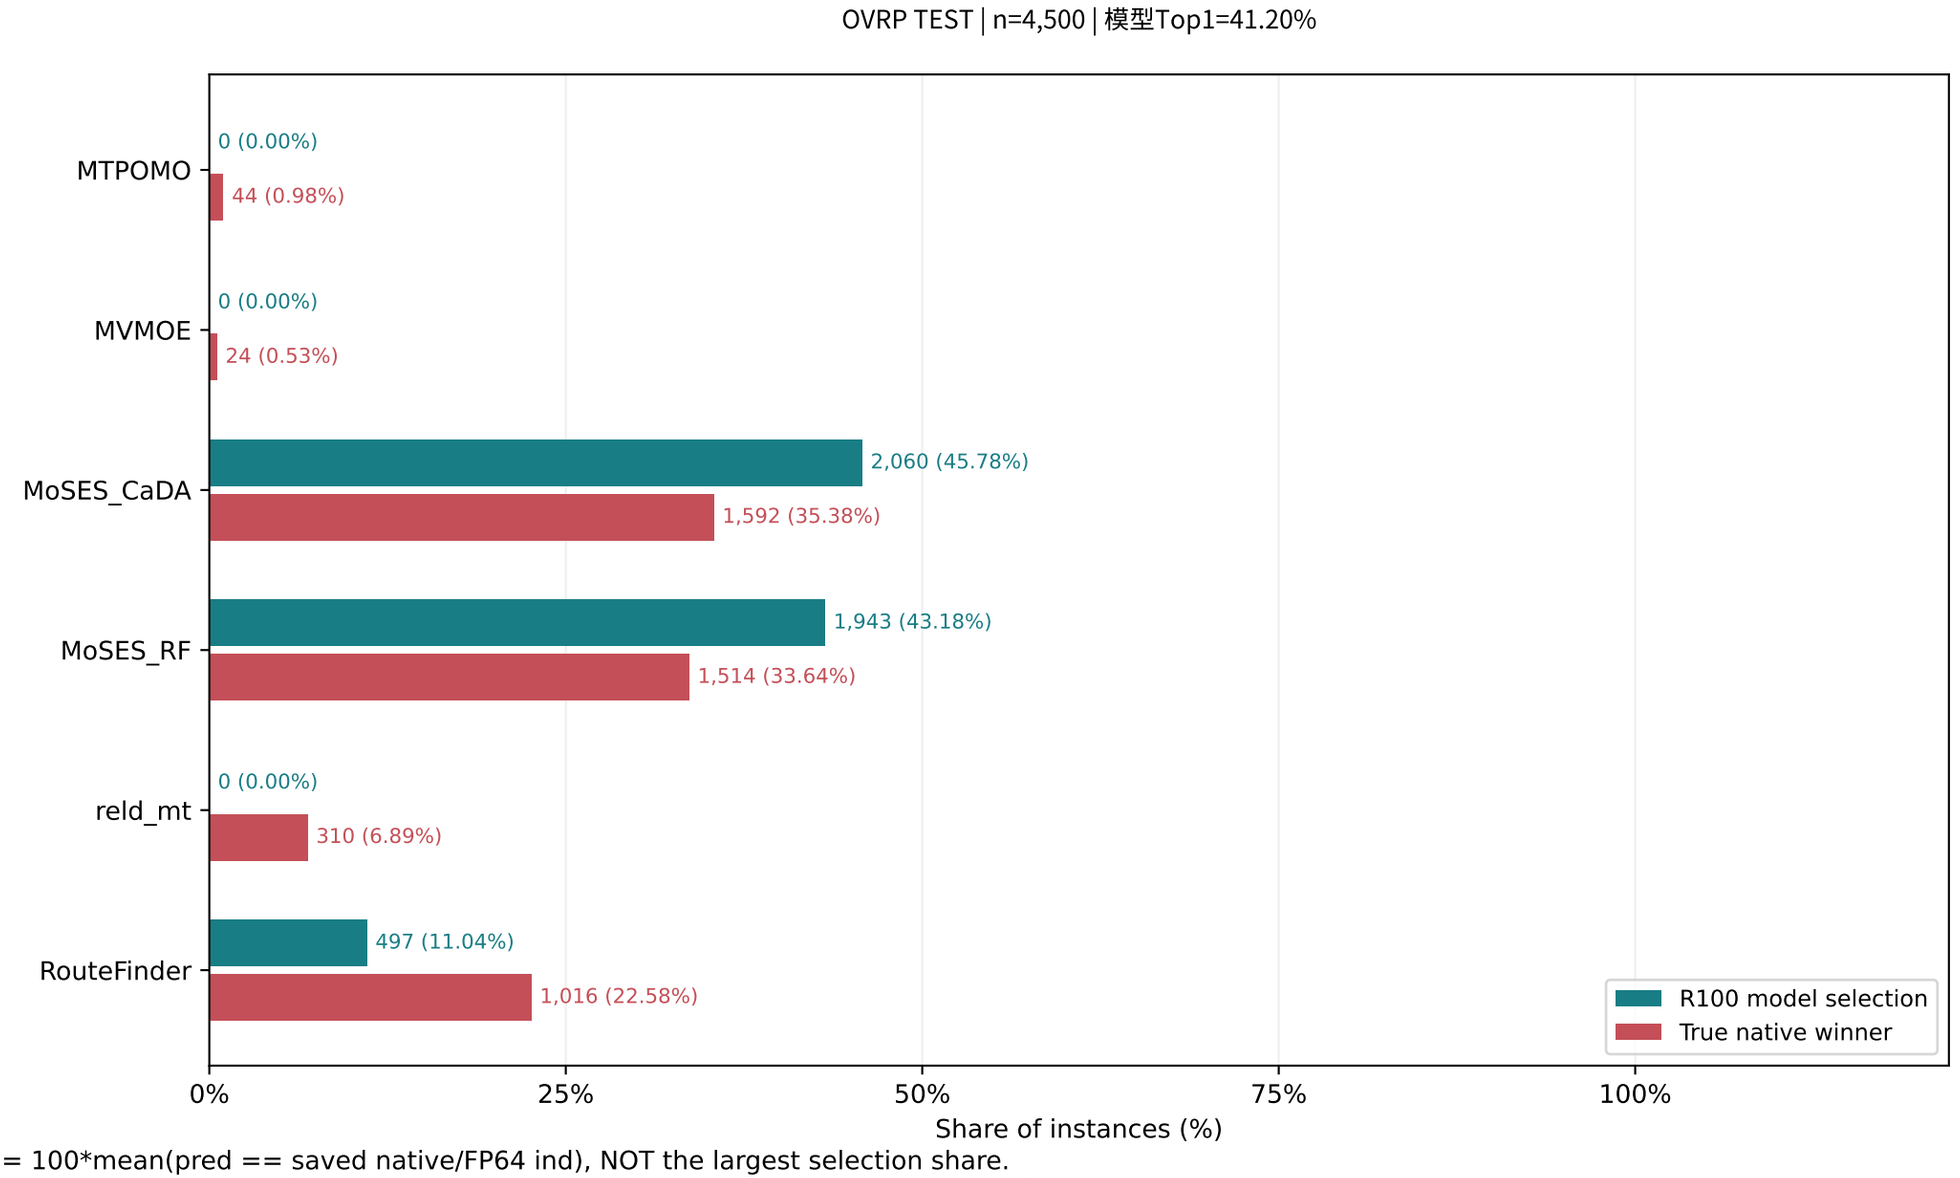


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | OVRP | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

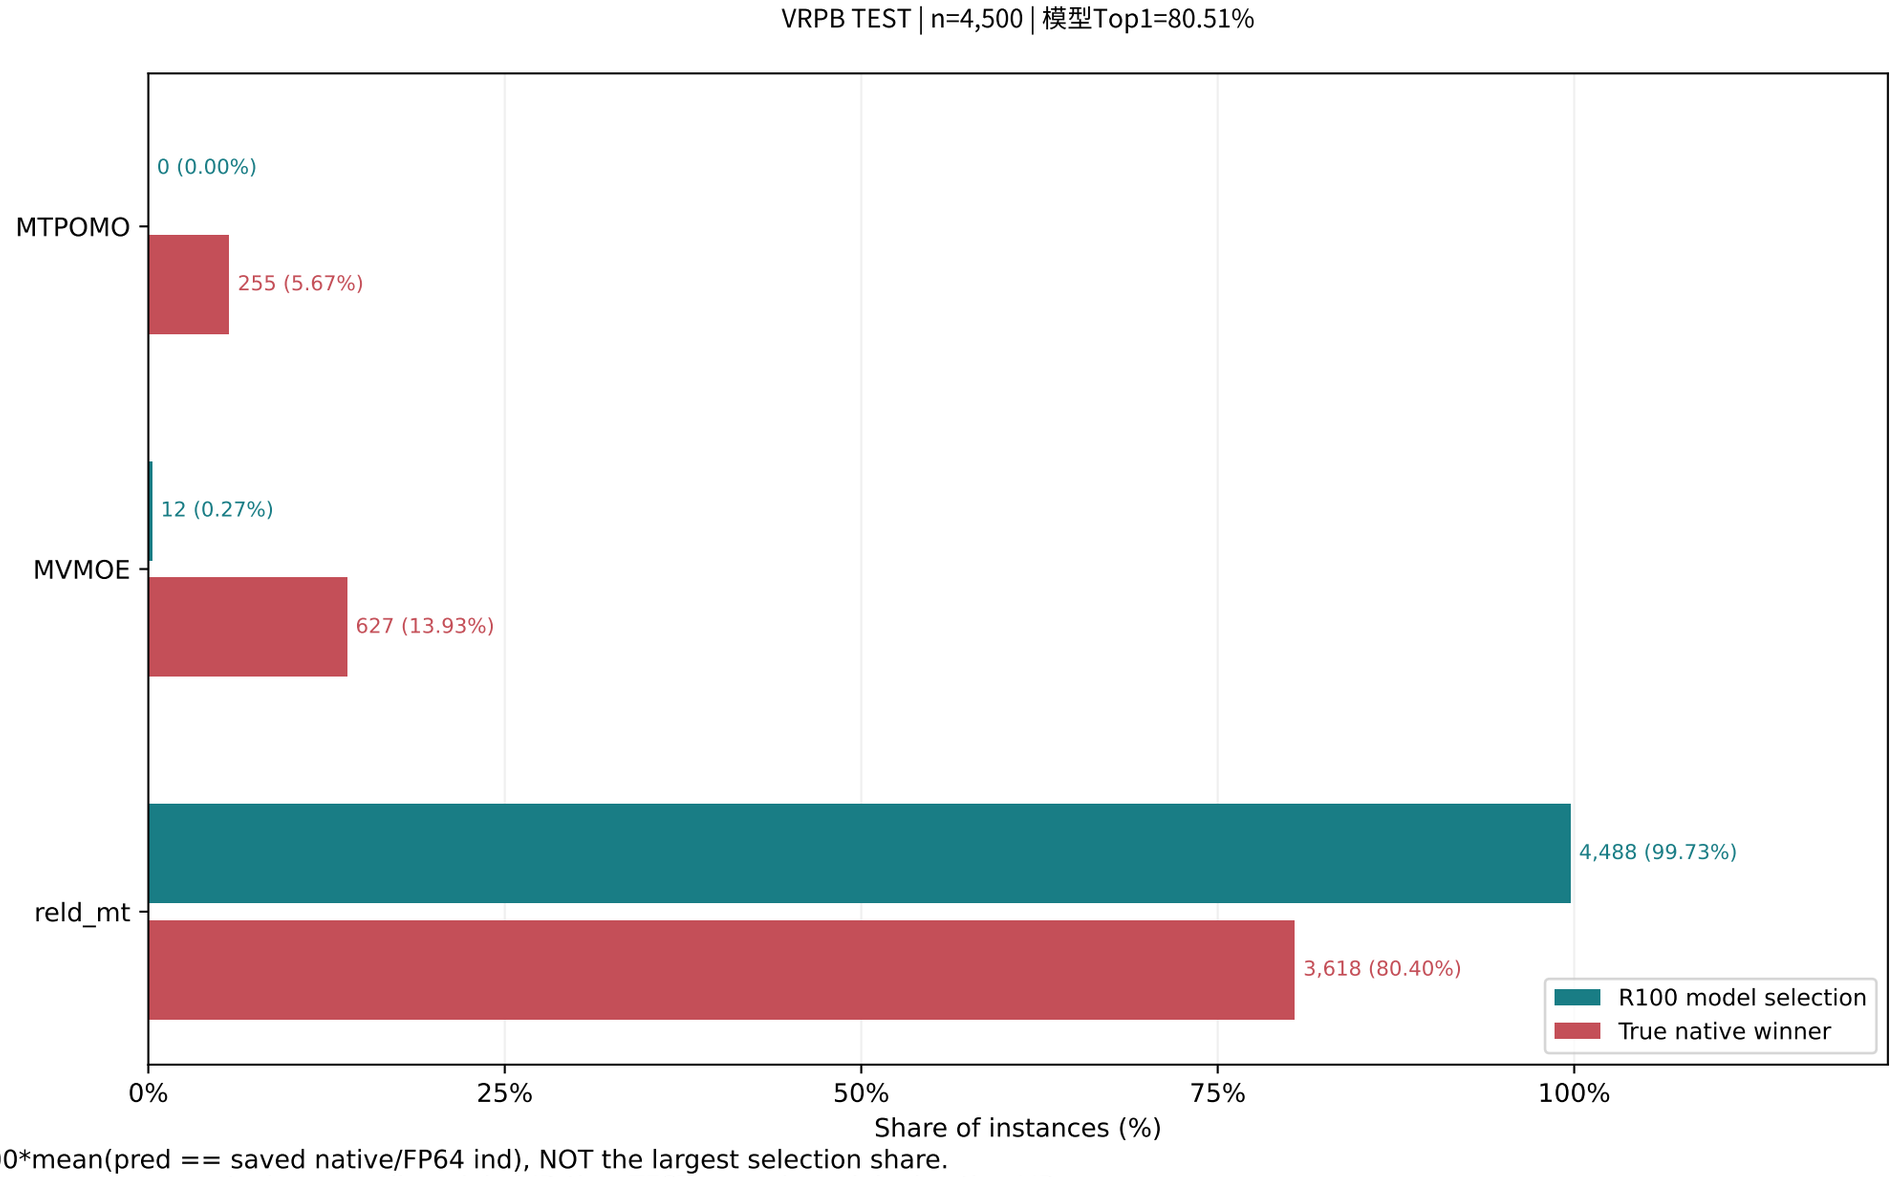


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | VRPB | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

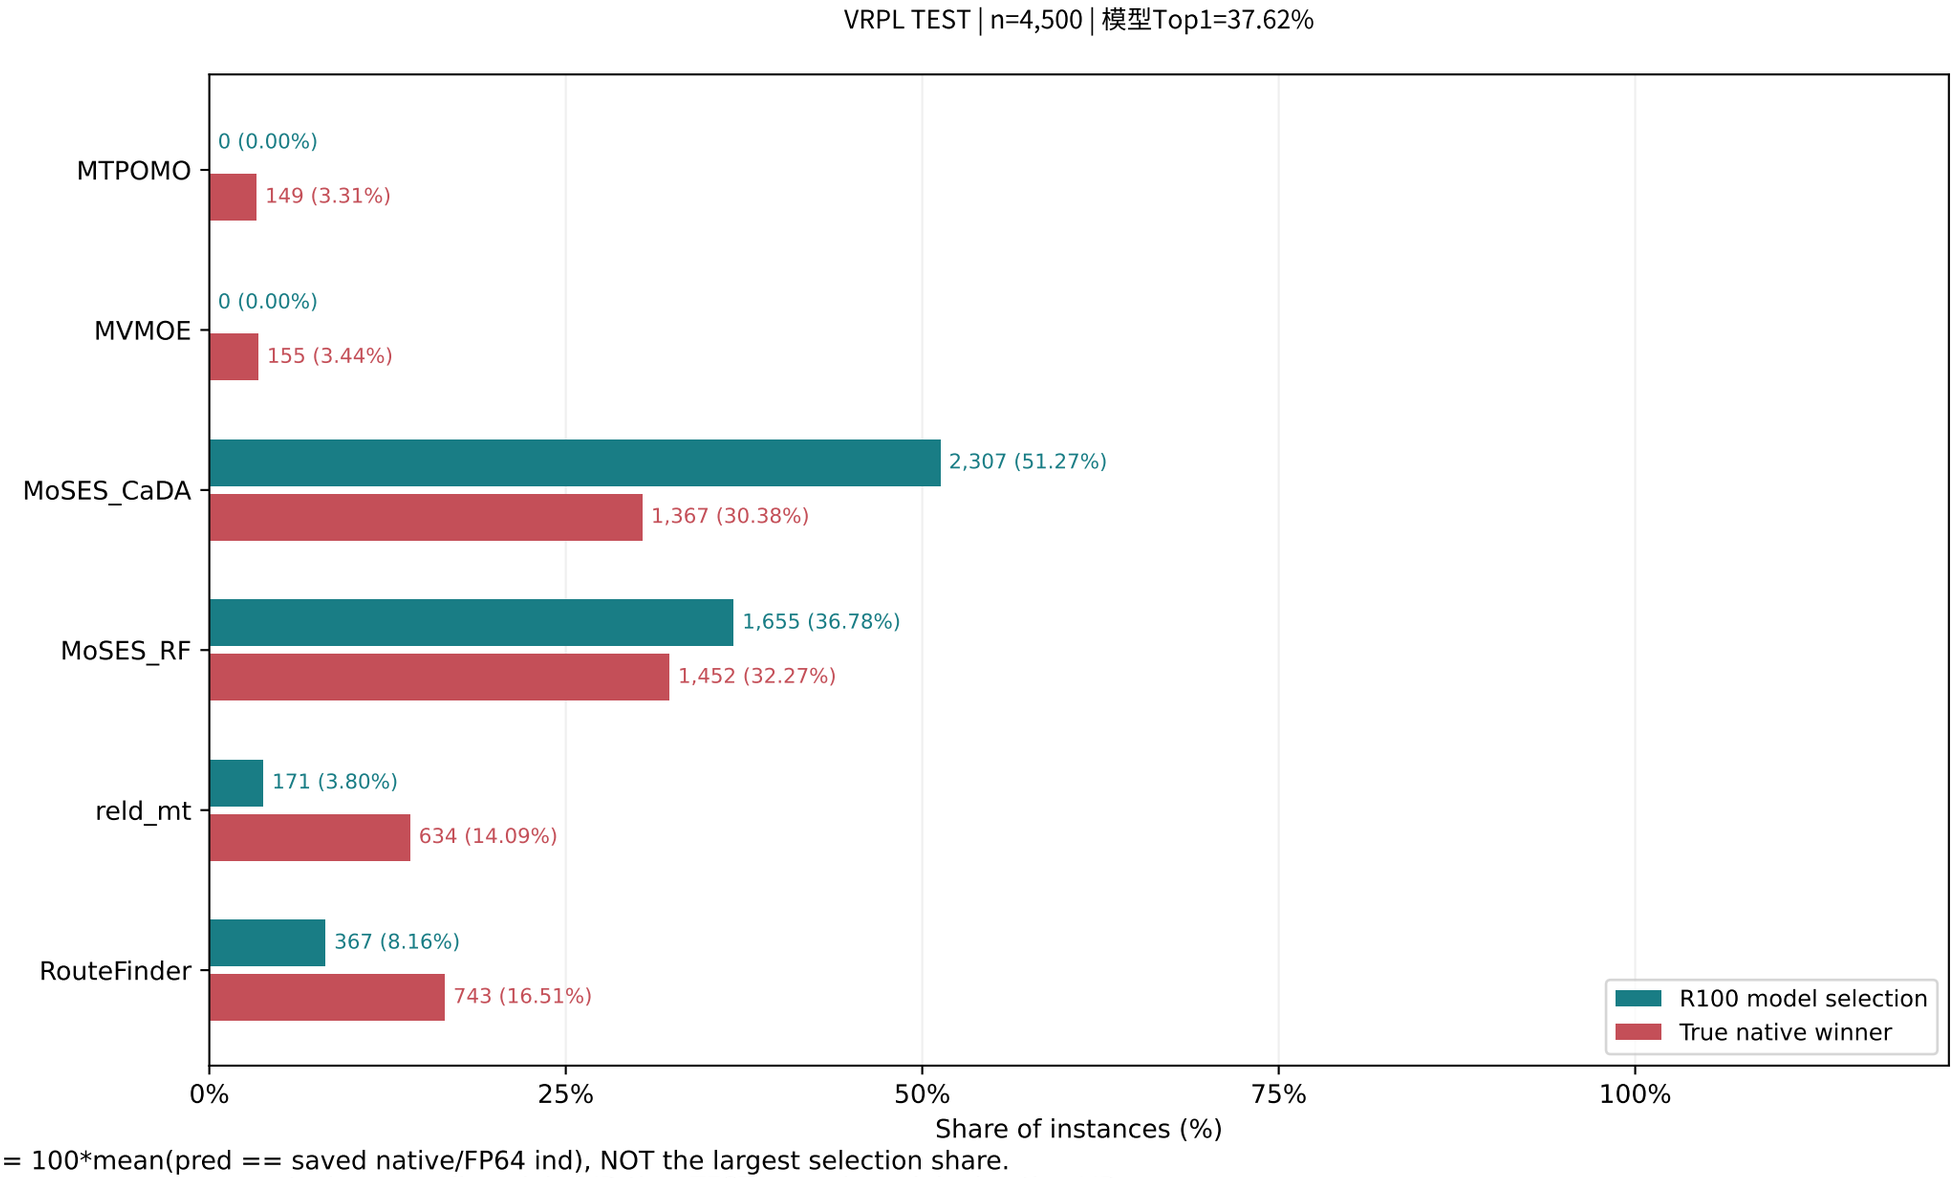


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | VRPL | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

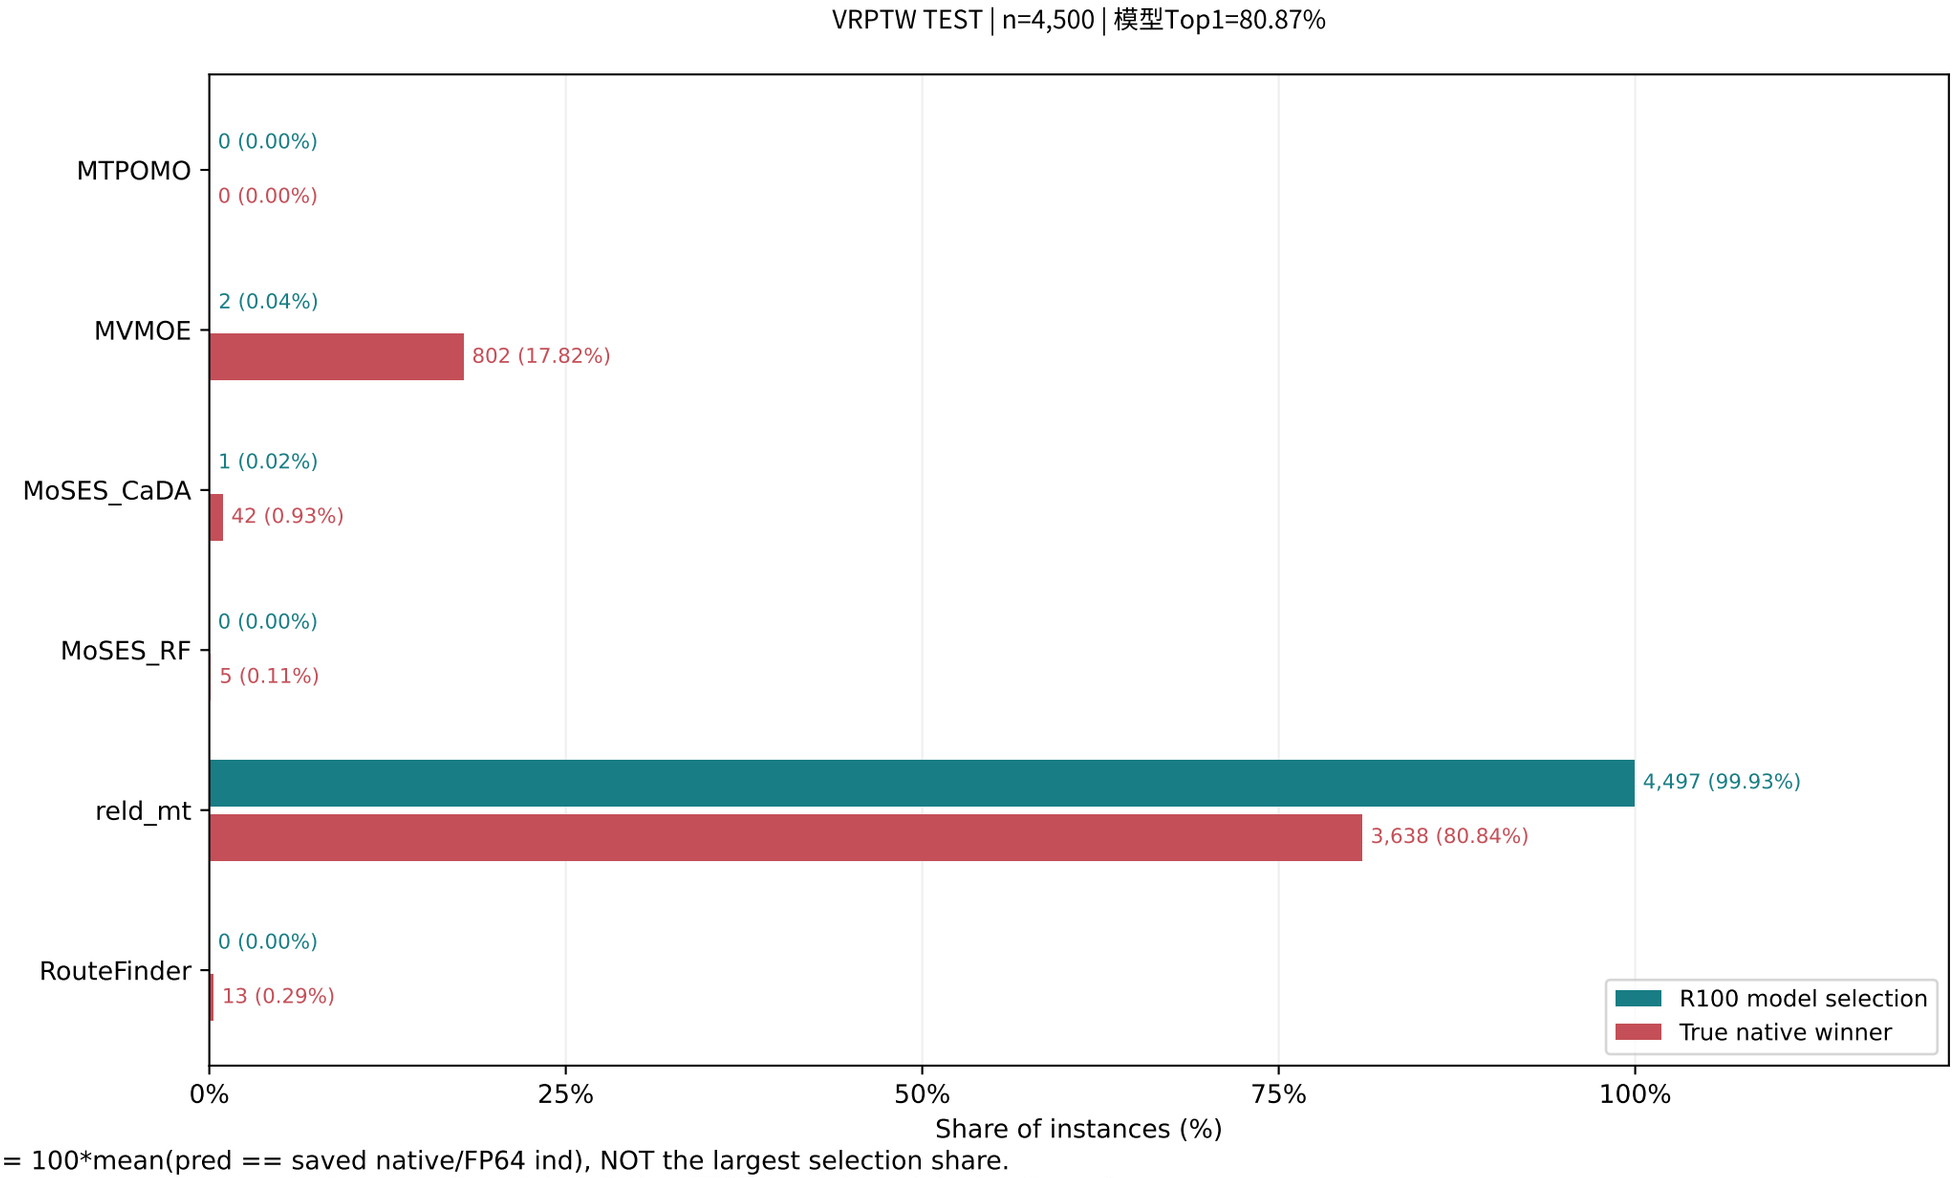


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | VRPTW | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

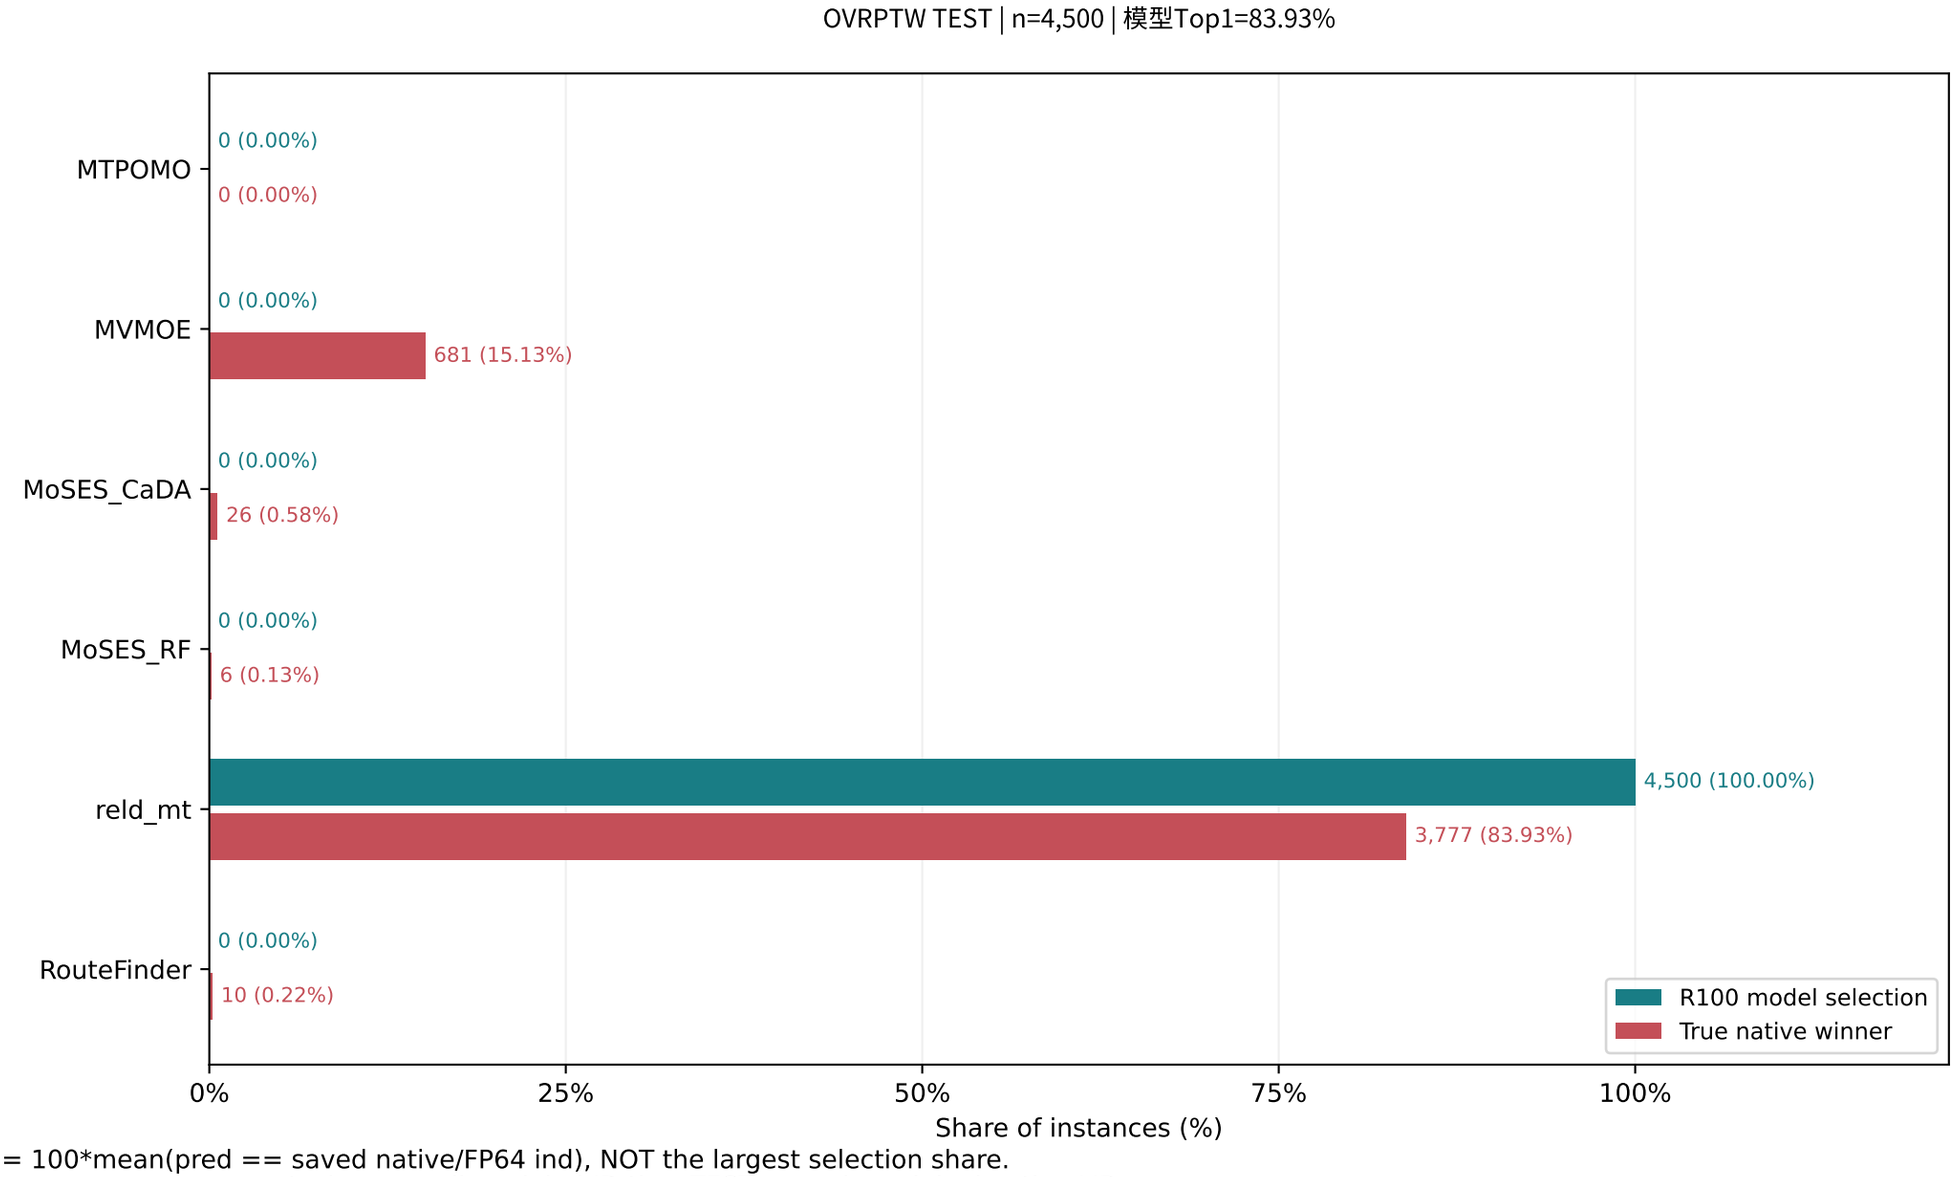


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | OVRPTW | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

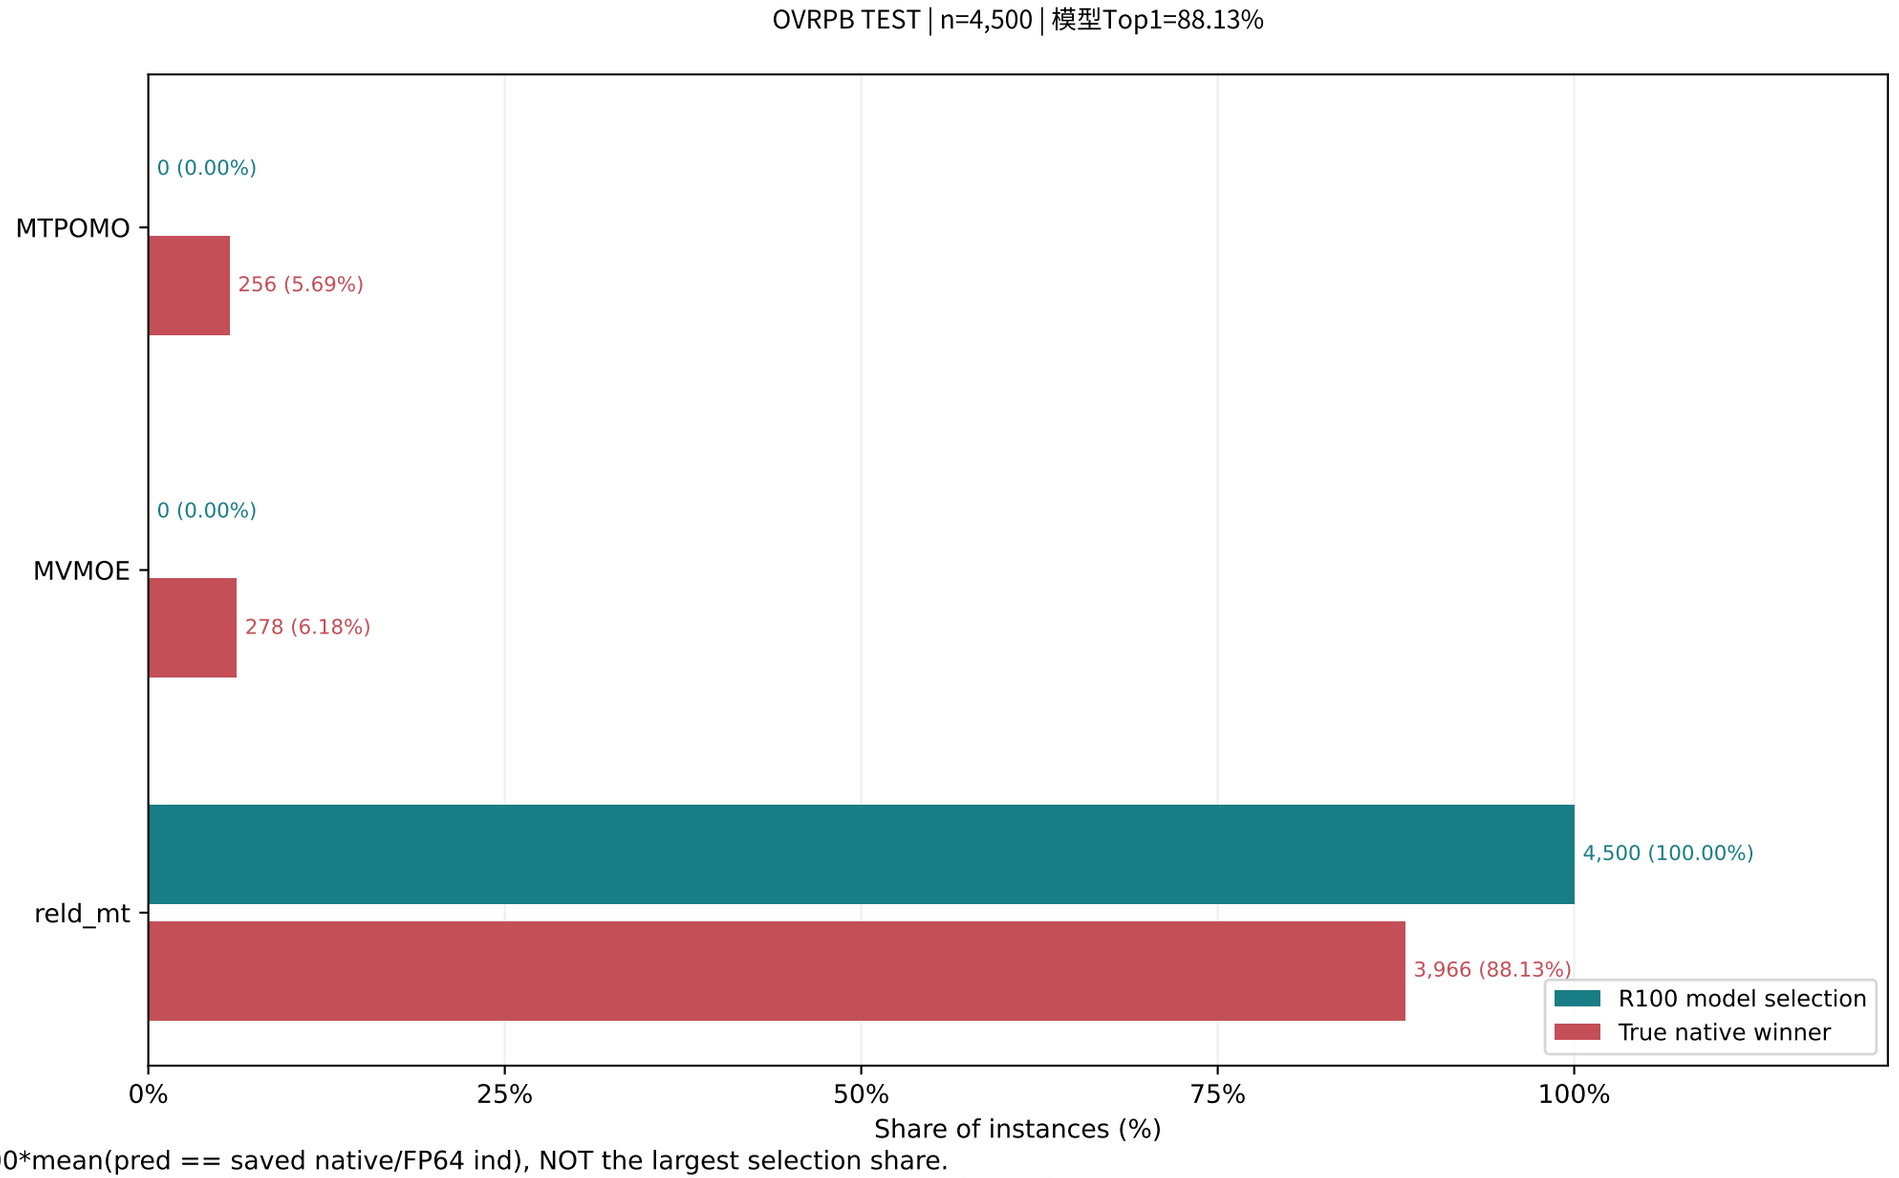


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | OVRPB | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

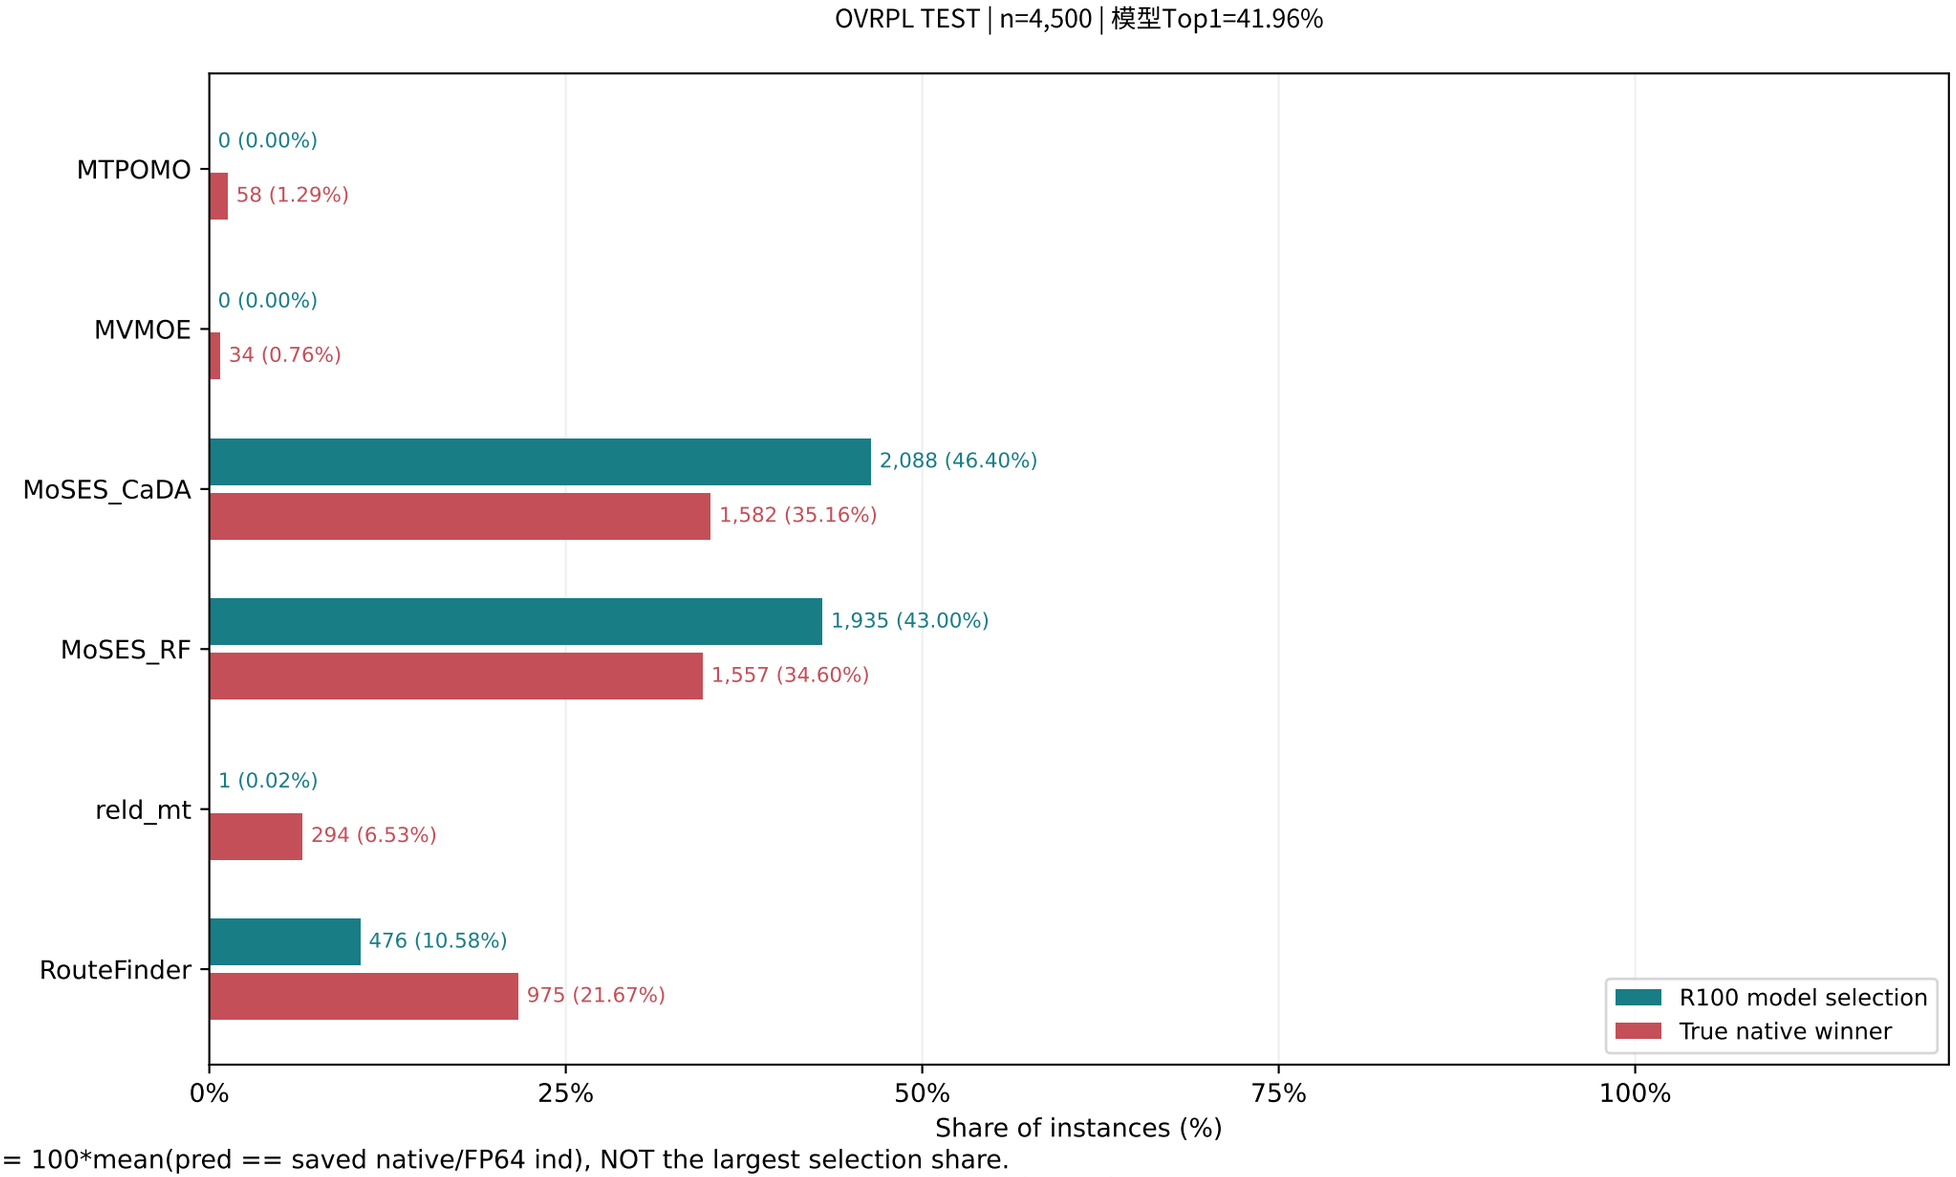


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | OVRPL | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

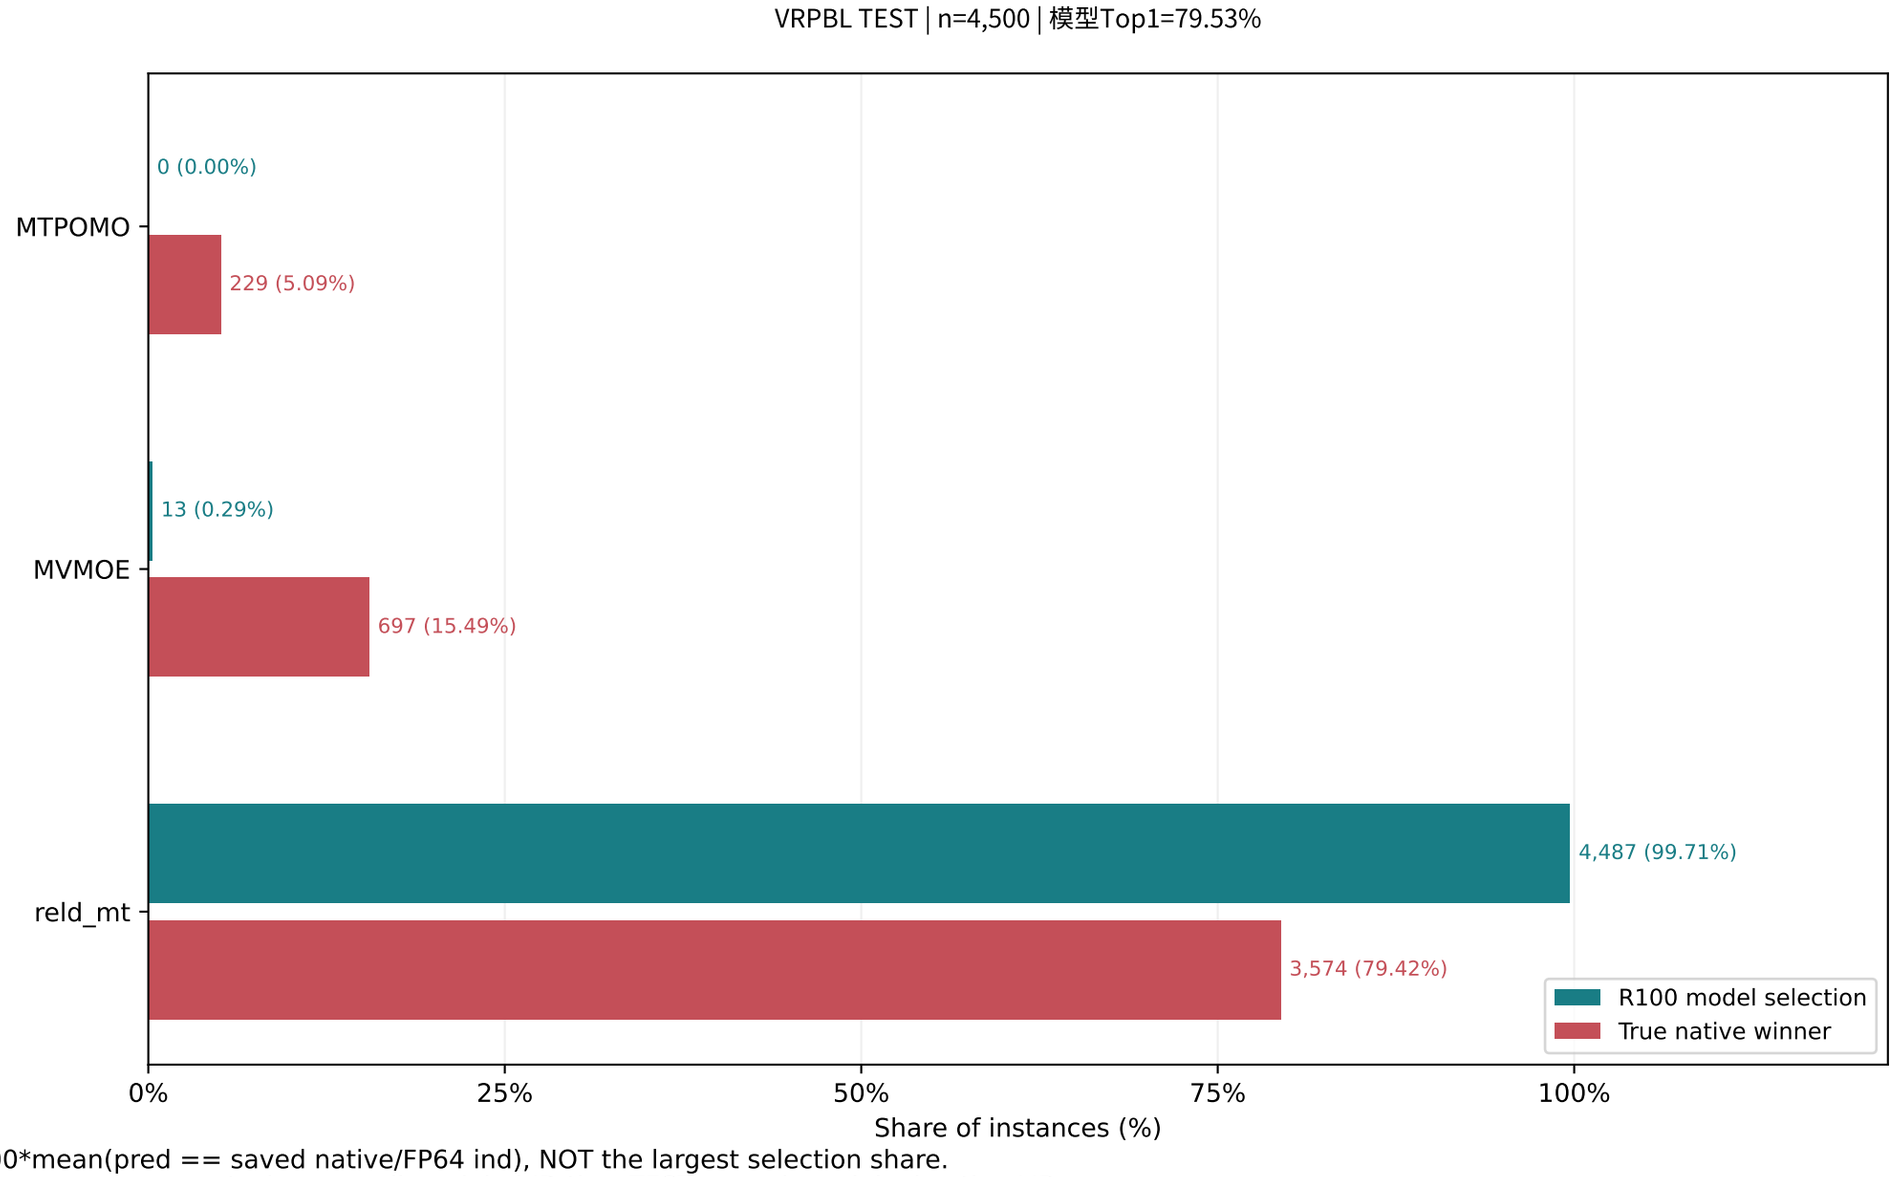


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | VRPBL | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

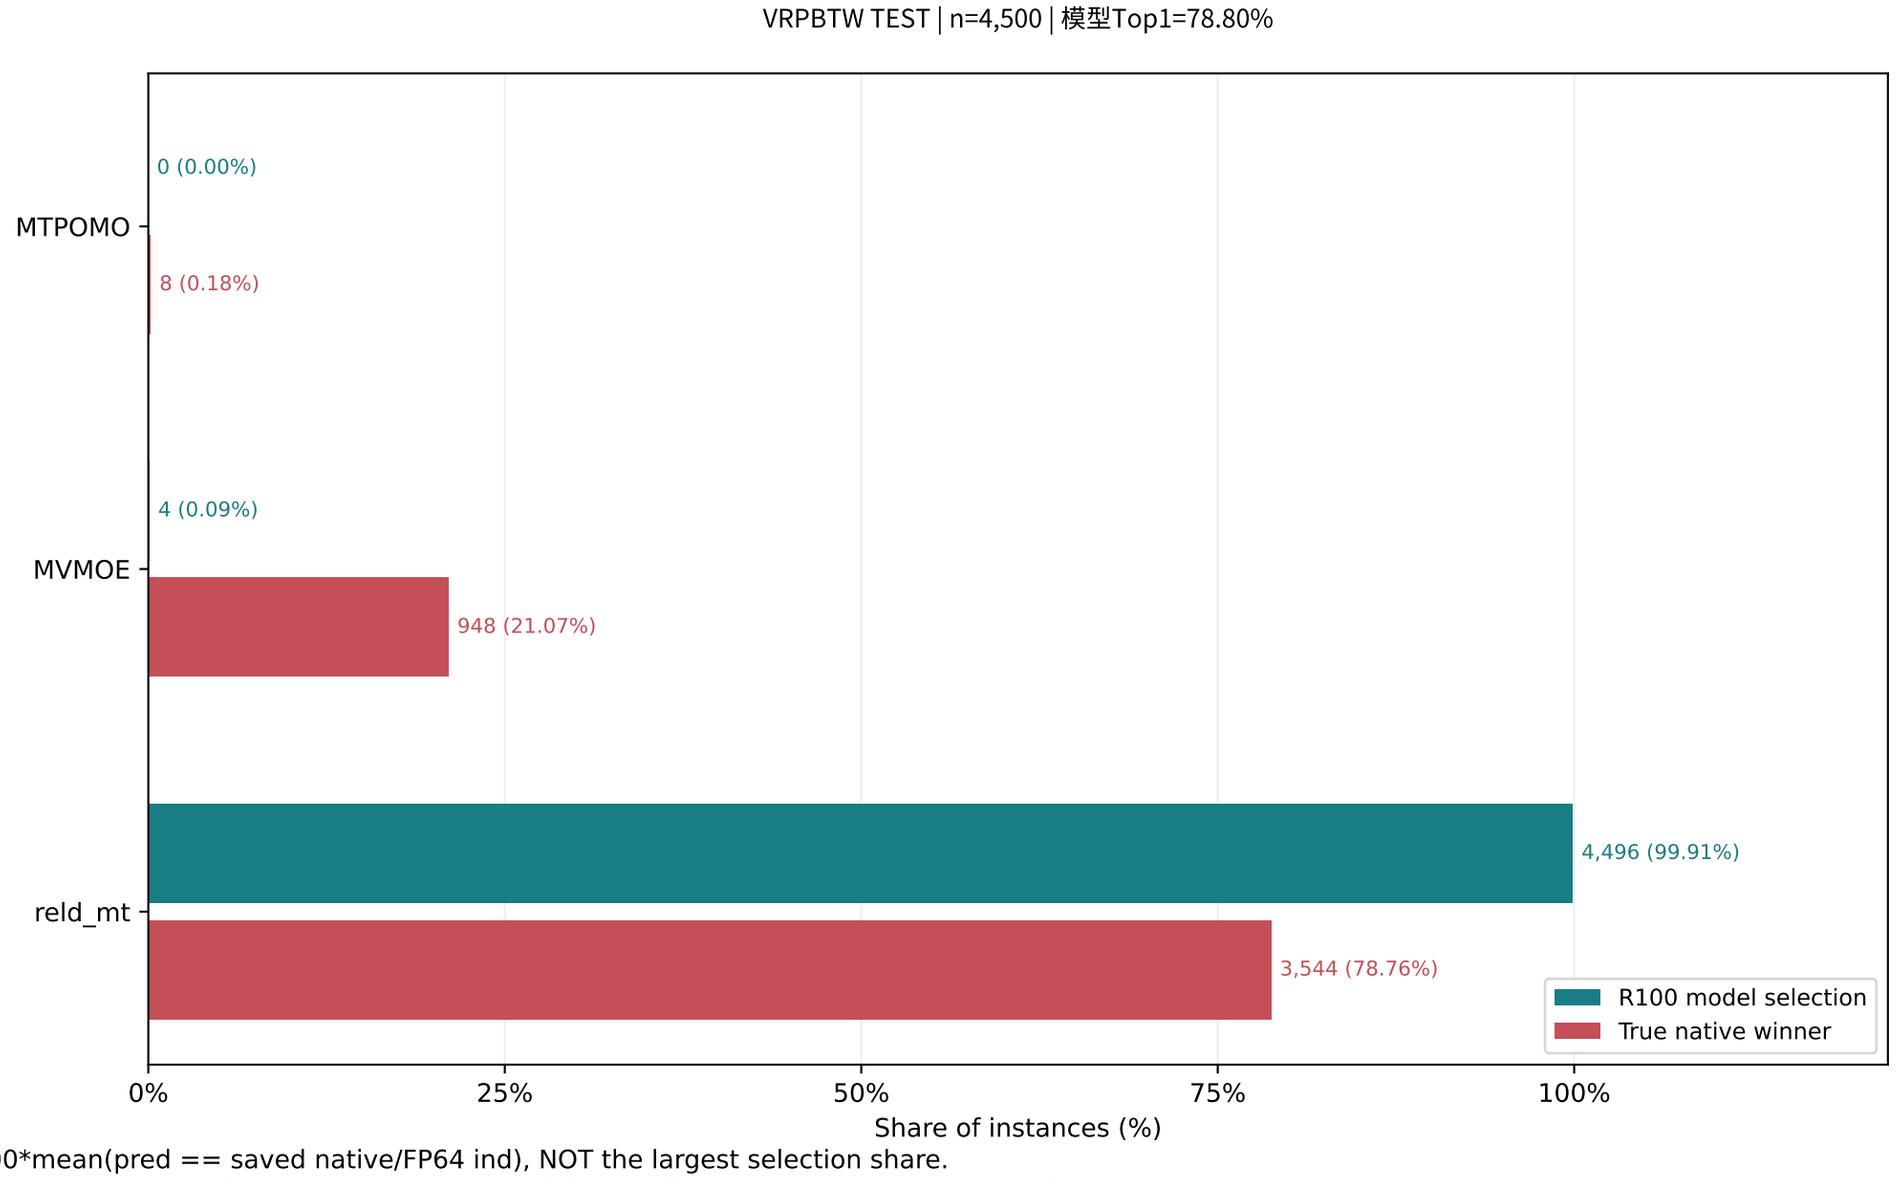


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | VRPBTW | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

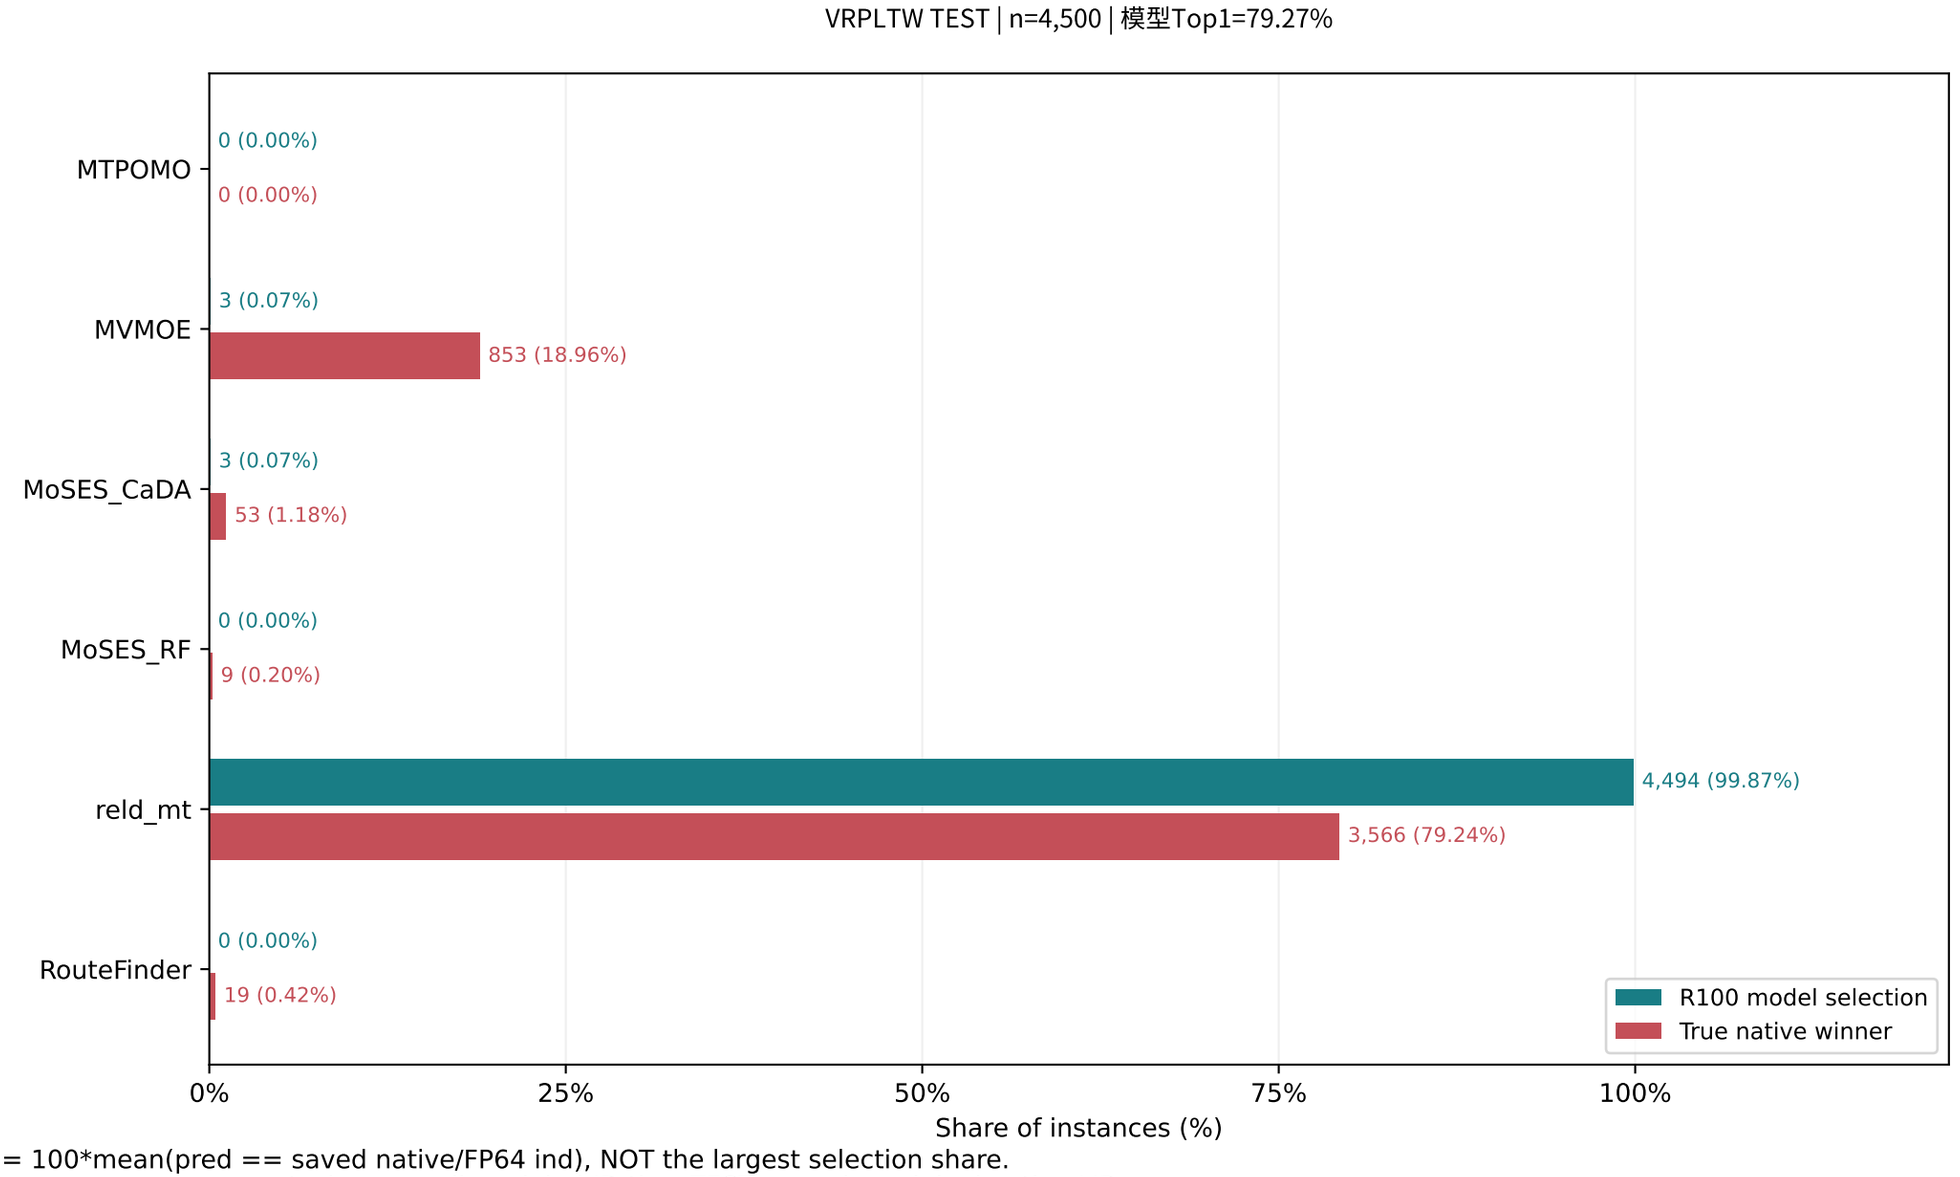


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | VRPLTW | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

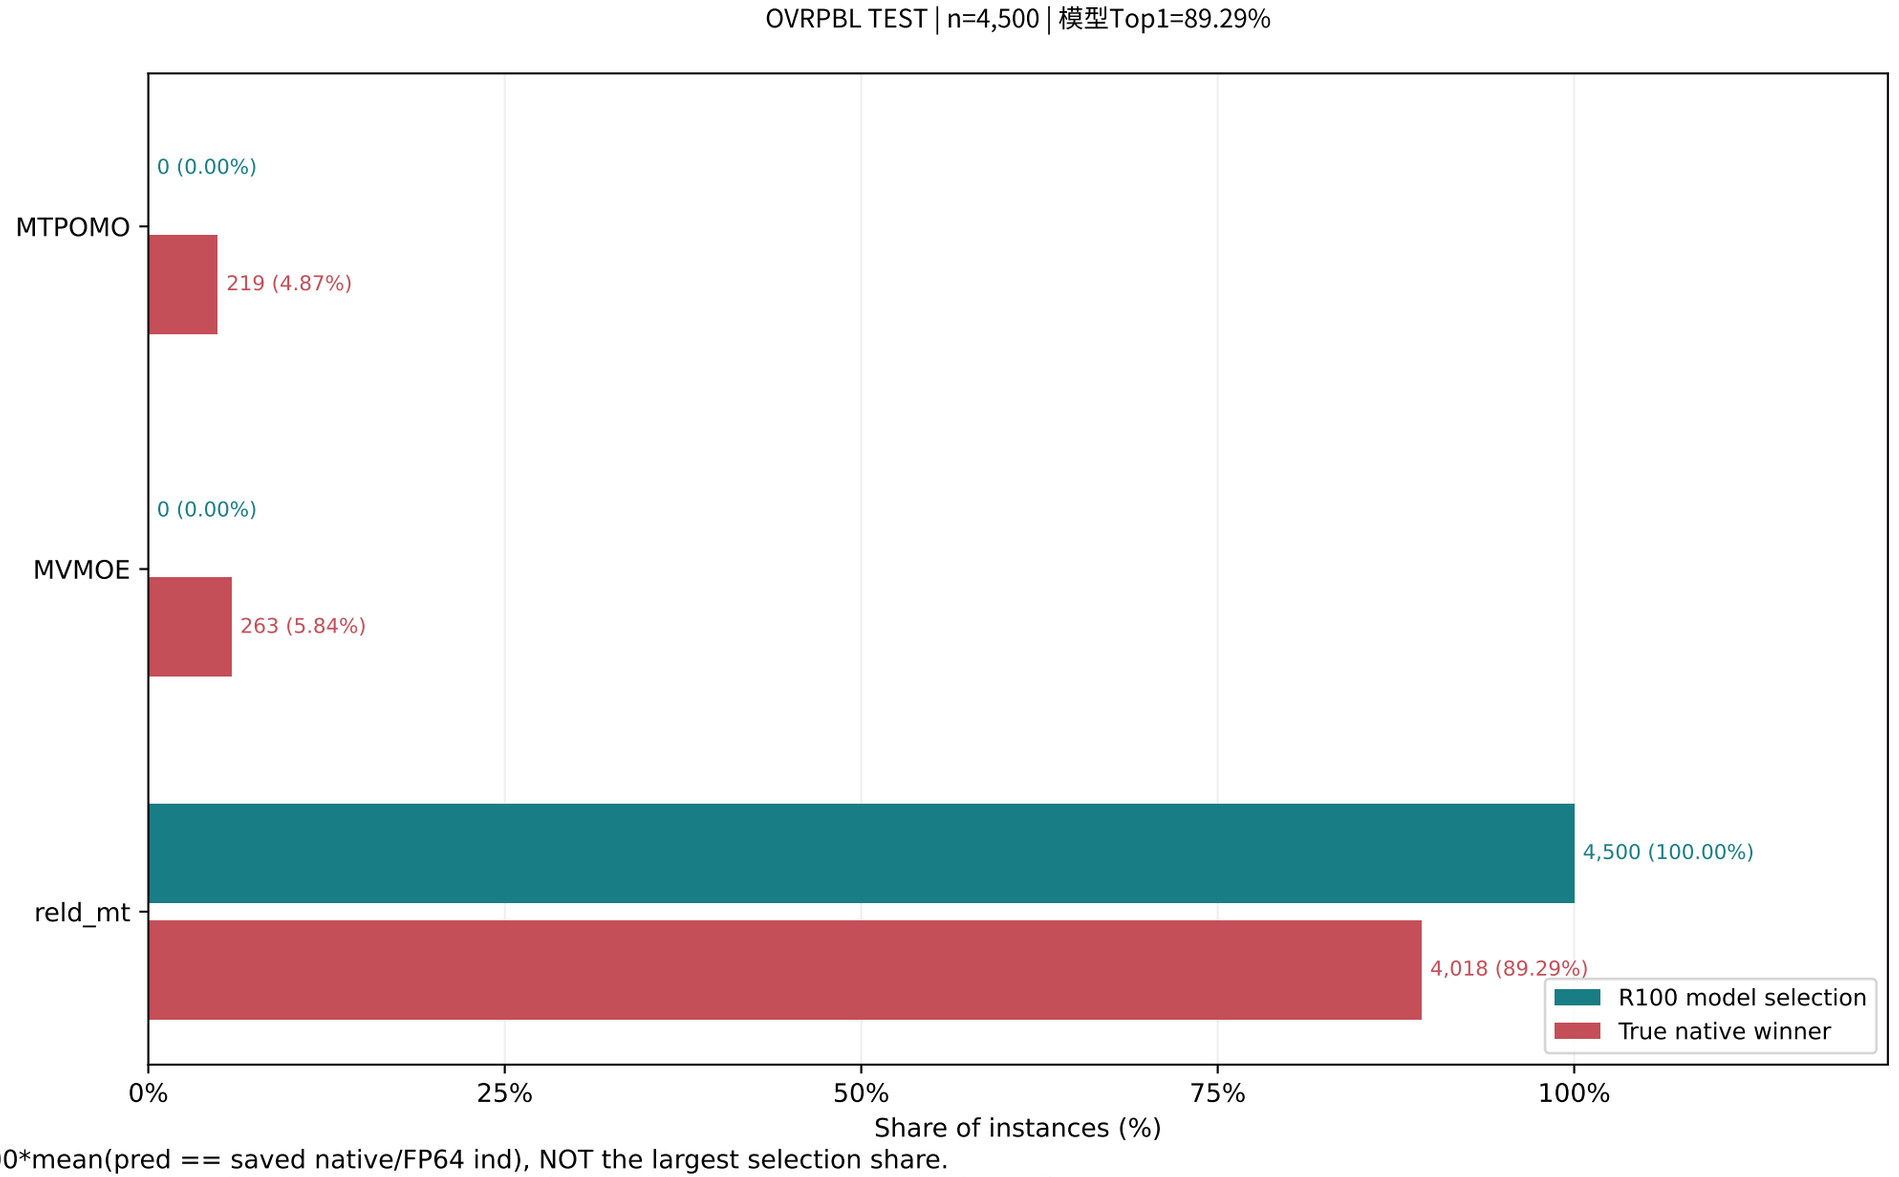


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | OVRPBL | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

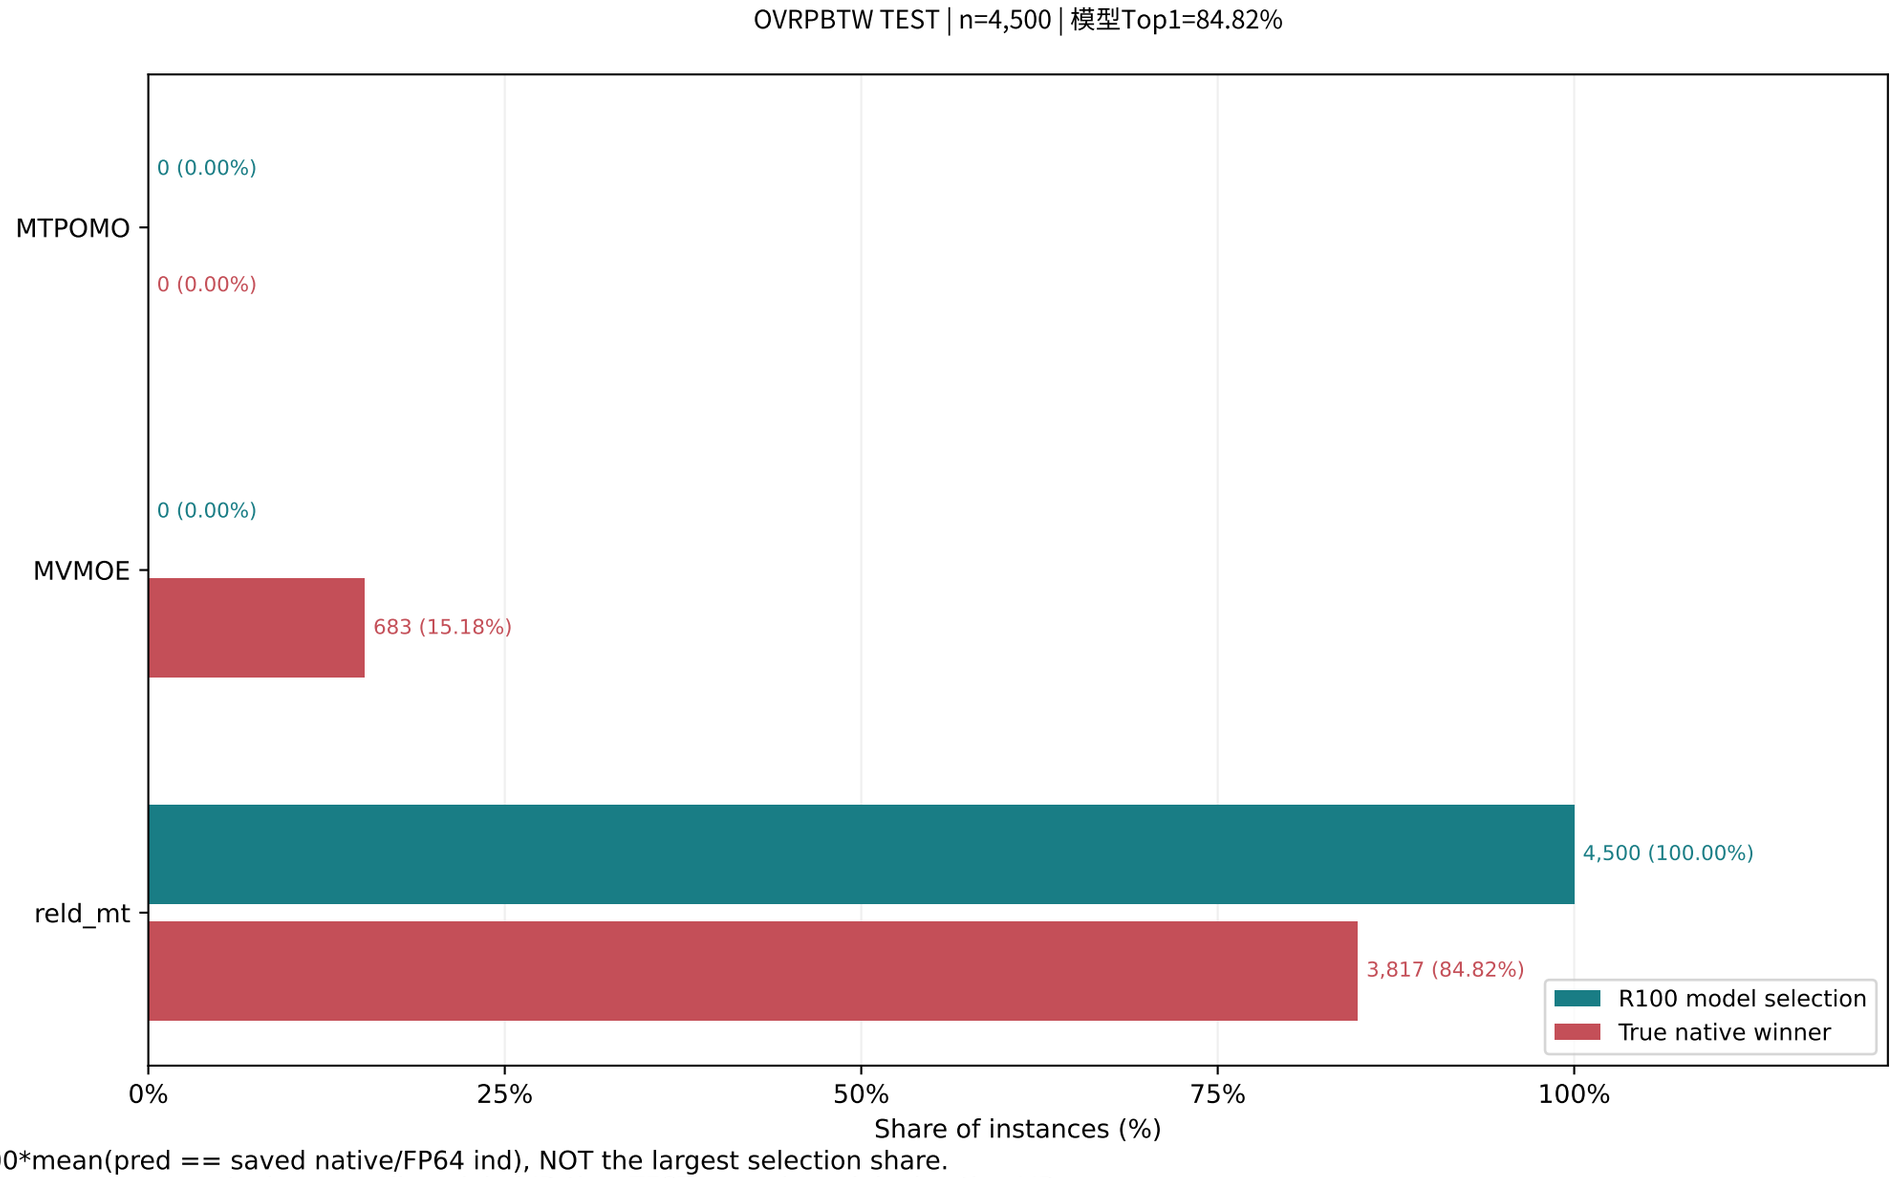


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | OVRPBTW | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

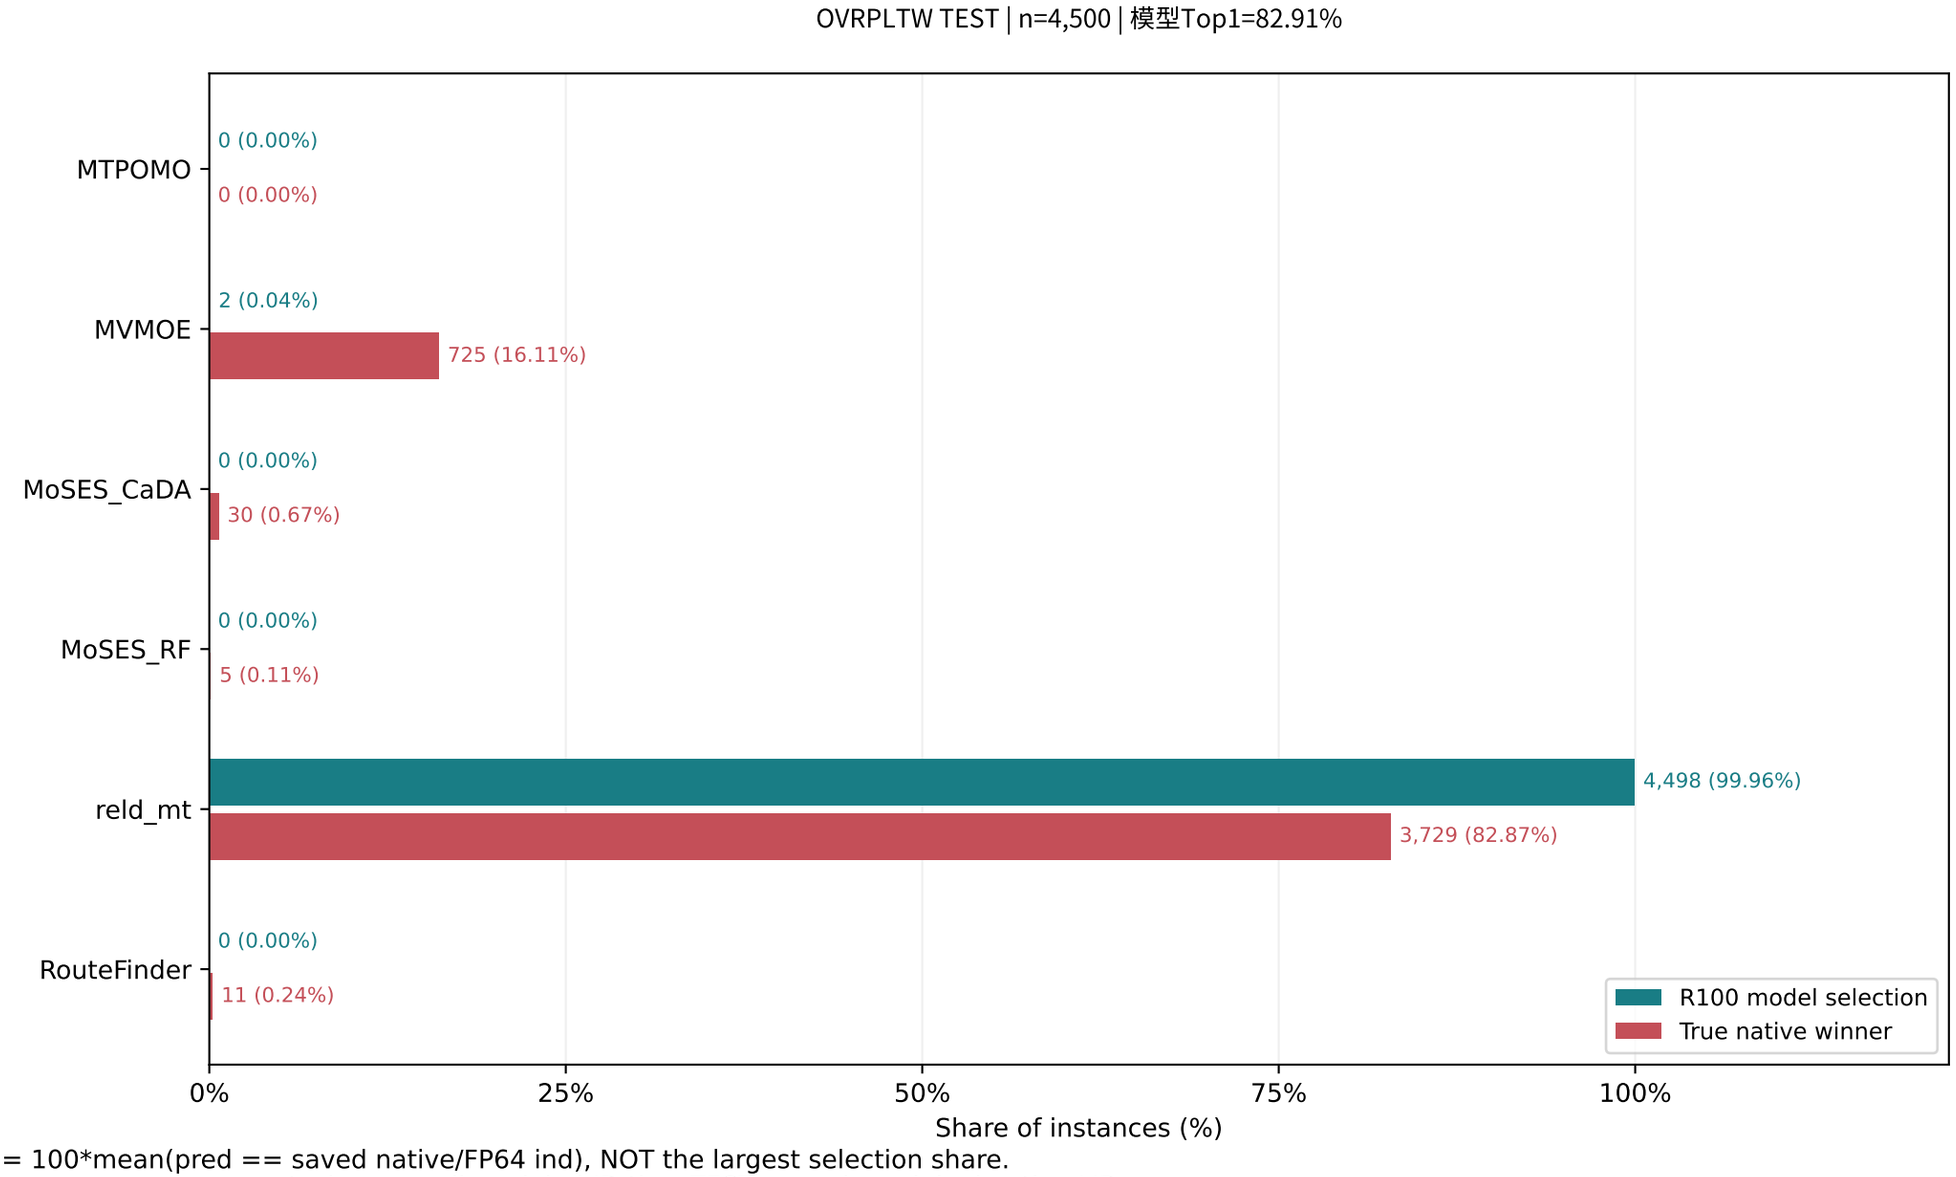


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | OVRPLTW | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

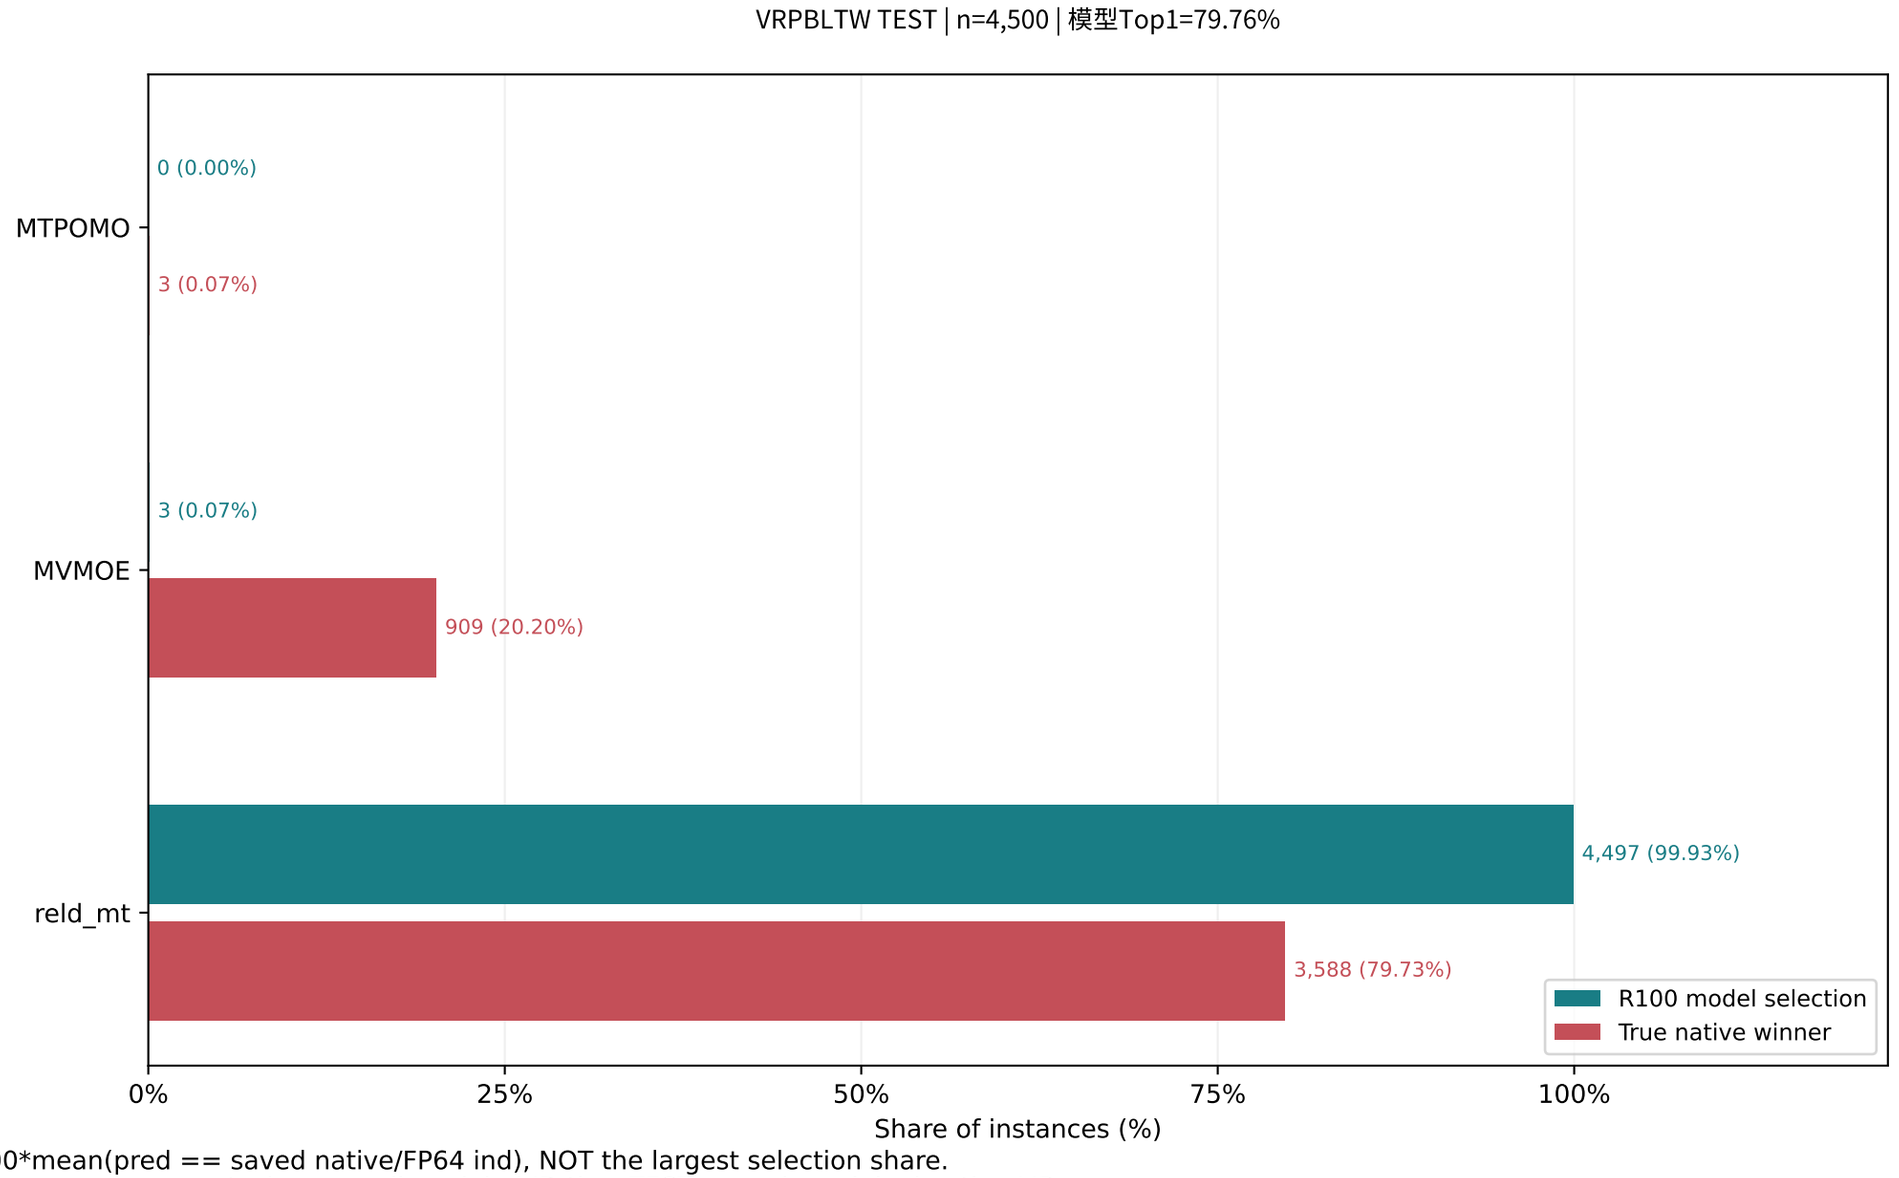


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | VRPBLTW | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea

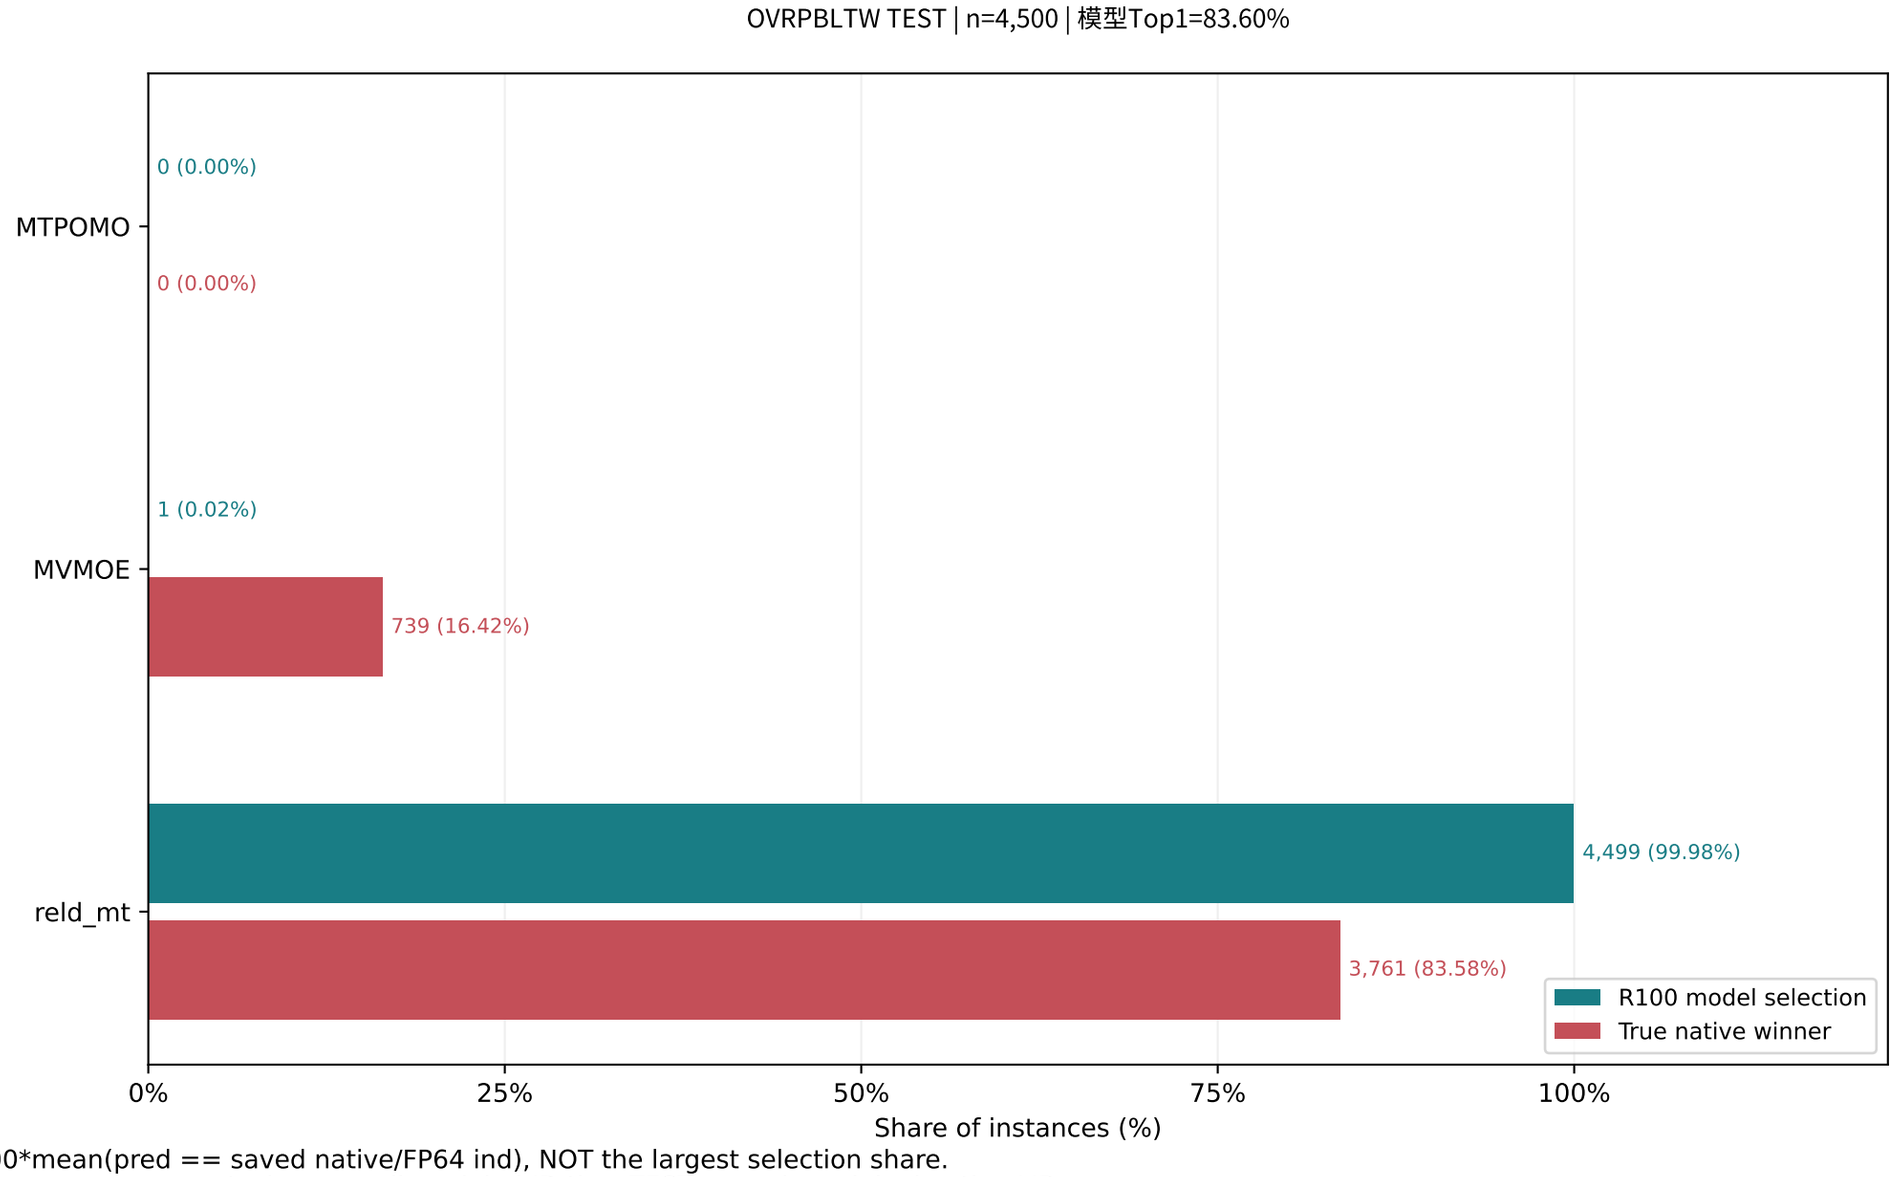


Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.

R100 | OVRPBLTW | Frozen Pairwise Binary Context

Data: selection_data_new90k_B_filtered_single_reld_mt_seed2_20261018
Bank SHA256: d46565818b3e4173afa37a84183be9f3a2d0970b0db14b85f5d912a529c7f9ea



Top1 = 100*mean(pred == saved native/FP64 ind), NOT the largest selection share.
Argmax raw mean opponent win; ascending global-ID ties. TEST precedes original native LIB.

Pair win probabilities and formal exp(mean-win) views are NOT full-pool winner probabilities.