R101/B | Two-stage provenance

Fixed TRAIN hint endpoint, then independently reset original full18 learning.

Teacher: pretrained. Stage I:12800 updates, logical batch1024, seed2; CVRP+15MVRP TRAIN81000/task only.

Separate non-affine per-node LayerNorm MSE; one dropout forward. No outcome labels, learned adapter or teacher VAL/TEST. Fixed endpoint, not validation selected.

Stage II: fresh original R84/R88B student; ONLY instance_encoder transferred. All non-encoder parameters remain identical fresh R84. Optimizer, scaler, RNG and task/sample streams reset.

Original1842953 parameters; no added heads or features. R77 CE .35 / comparison .35 / risk .02 / R-Drop .35. No hint loss, teacher arrays or online donor in Stage II or deployment.

Both A/B have12800+57600 updates; R89B/R84A are NOT total-budget-matched. B teacher has additional external solver pretraining history. The paper memo provides conditional motivation only.

Native/FP64 ind winners; legal-pool Oracle and TRAIN-fitted SBS; ALL is task-equal. Development-exposed TEST; no TEST/LIB checkpoint selection.

Stage-I encoder SHA256: 6461d4d78289f6cbb92f534034602eea4eba389fc606f1fc0b386ed2662c657a

Unchanged non-encoder SHA256: 970a8aa2a548f9034f34b73d060662f71808e46359335185ce07bfc5b1bb2fa6

R101/B | TRAIN-only pretrained node hints; fixed12800 endpoint

Logged LN-MSE (last update, not interval mean)

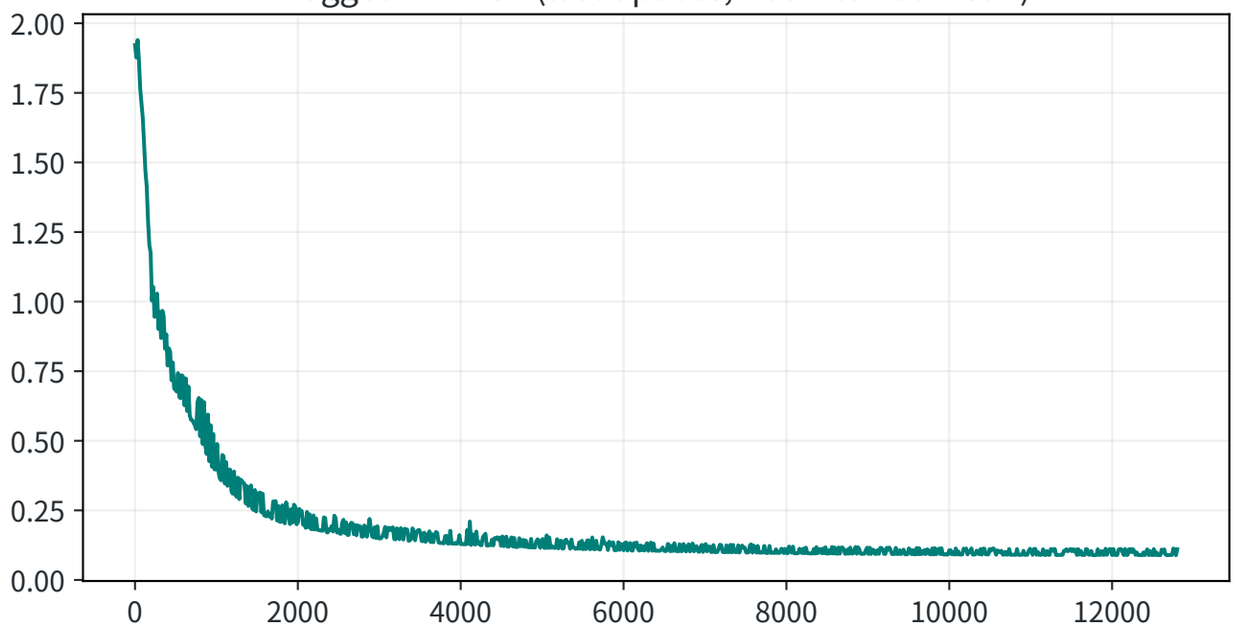

AdamW learning rate

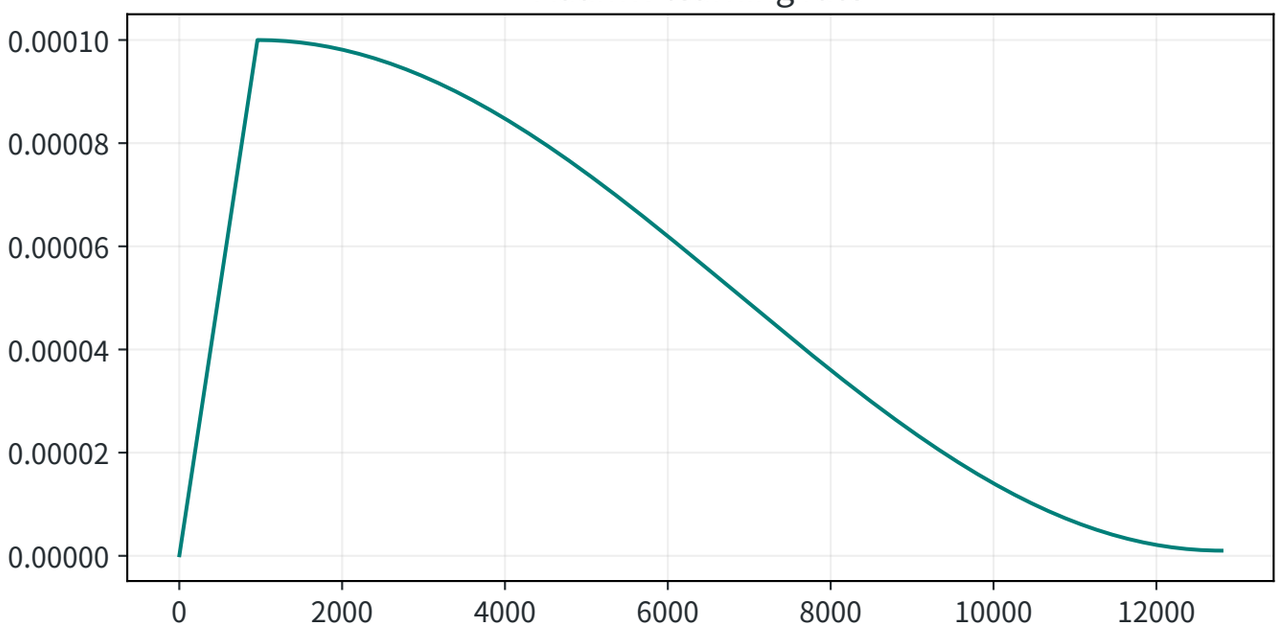

Gradient norm before clipping

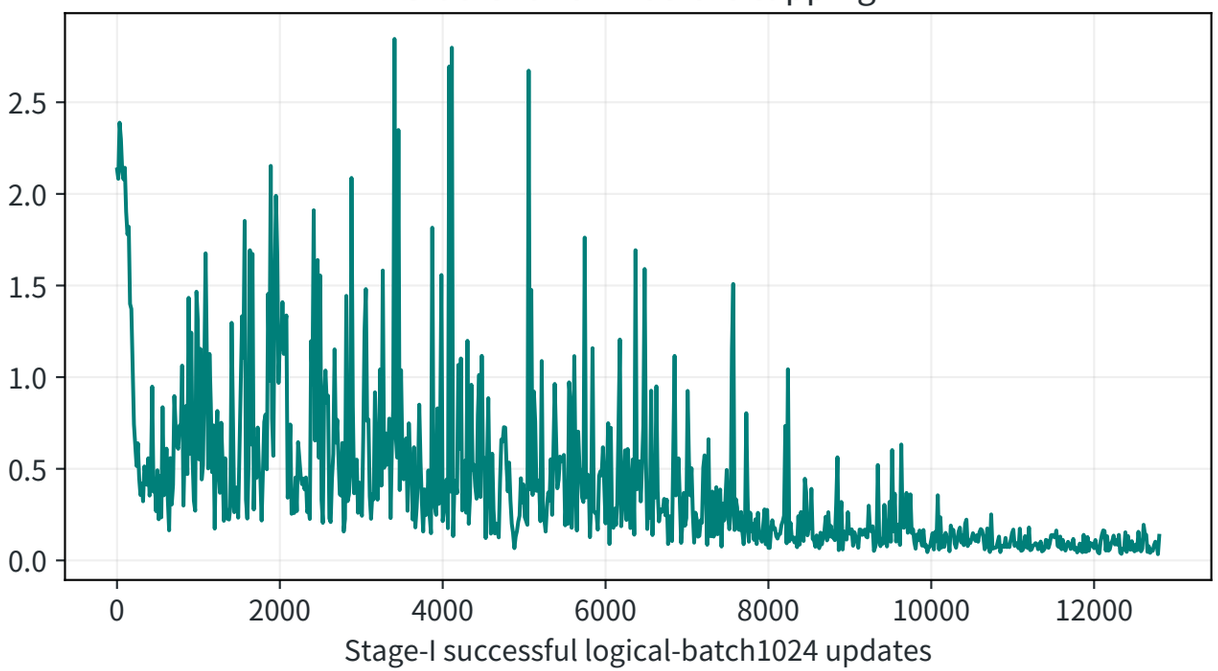
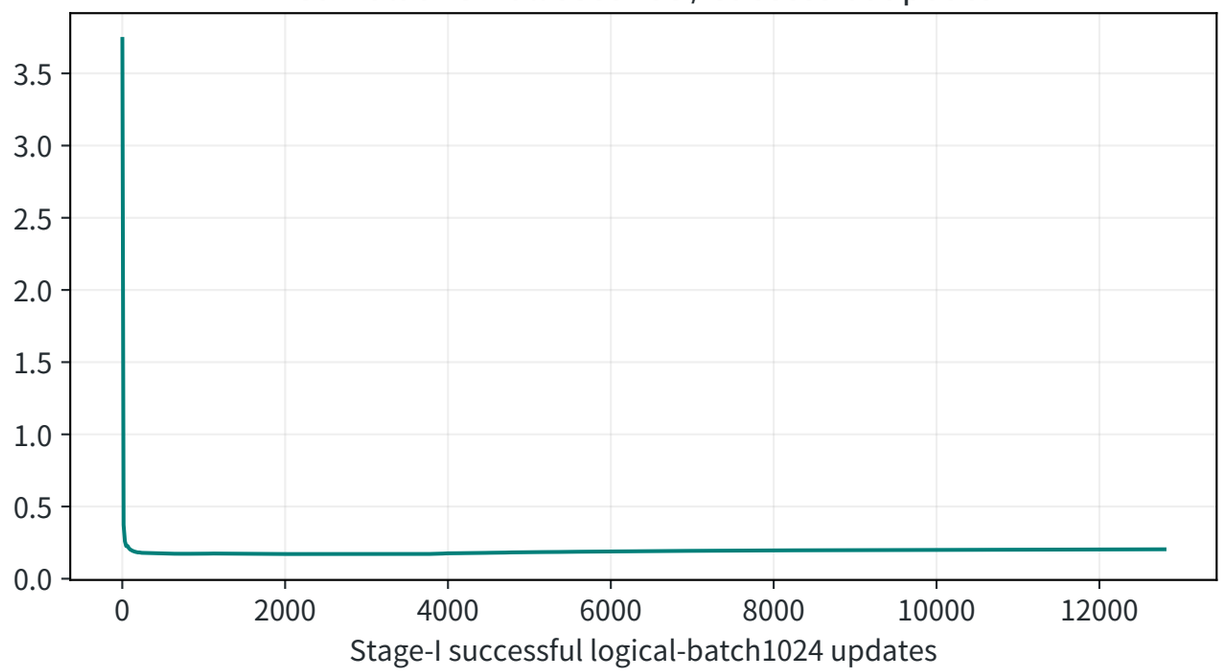
Cumulative wall seconds / successful update