R101 pretrained B vs random A | Fixed target36 vs R101/A

Native/FP64 winners; true rescues/losses and false-pick cost, not selection-frequency accuracy.

Split	True winners	Correct	Rescued	Lost	New false	Recall (%)
TRAIN	83279	919	639	127	535	1.1035
VAL	4625	52	42	5	29	1.1243
TEST	4598	50	38	7	38	1.0874

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	0.116324	24168	22472	-0.000724919	-0.006093
VAL	0.049383	1308	1268	-0.000272309	-0.002608
TEST	0.112346	1326	1235	-0.000388076	-0.002009

Split	New-false cost contribution	New-false cost delta	New-false regret pp
TRAIN	0.003040981	0.000079047	0.001686
VAL	0.003171510	0.000128414	0.002138
TEST	0.003721341	0.000180927	0.002915

Cost contributions: sum over newly false-picked instances / task n; ALL averages18 tasks.

R89B/R84A are historical, not12800+57600 budget-matched; zero-count movement is not success.

R101 B vs A | Original student latency

First128 GPU-resident original TEST inputs/task; three synchronized repetitions.

Problem	A student ms/row	B student ms/row
TSP	0.243893	0.242471
CVRP	0.291117	0.291154
ATSP	0.065964	0.061758
OVRP	0.064915	0.061512
VRPB	0.065113	0.061487
VRPL	0.065013	0.061102
VRPTW	0.064890	0.061006
OVRPTW	0.065490	0.060956
OVRPB	0.064941	0.061135
OVRPL	0.064911	0.060924
VRPBL	0.065333	0.060681
VRPBTW	0.065099	0.061029
VRPLTW	0.064982	0.060763
OVRPBL	0.065081	0.061138
OVRPBTW	0.065137	0.061167
OVRPLTW	0.064811	0.061032
VRPBLTW	0.065158	0.060978
OVRPBLTW	0.065094	0.060994

Original1842953-parameter student only; no online donor, teacher feature I/O or solver decoder.

Excludes dataset I/O, host transfer and Stage-I pretraining; not single-instance/E2E latency.

No latency or speedup claim is inferred from R94 cached-selector numbers.

Hardware is disclosed, not bitwise matched; no model reselection based on timing.