R104/A_width256 | Capacity and actual latency

Complete standalone model; synchronized GPU-resident original inputs; no donor.

Problem	Instances	Milliseconds / instance
TSP	128	0.518652
CVRP	128	0.611607
ATSP	128	0.099529
OVRP	128	0.097071
VRPB	128	0.096139
VRPL	128	0.096957
VRPTW	128	0.096874
OVRPTW	128	0.096459
OVRPB	128	0.096244
OVRPL	128	0.097163
VRPBL	128	0.096288
VRPBTW	128	0.095911
VRPLTW	128	0.096844
OVRPBL	128	0.096046
OVRPBTW	128	0.096209
OVRPLTW	128	0.096729
VRPBLTW	128	0.096233
OVRPBLTW	128	0.096489

Parameters: 7,288,969; copied fresh: 1,417; independent: 7,287,552

Width256 instance4/joint2 and width128 instance8/joint4 are two capacity axes, not equal-parameter, equal-FLOPs or identical-full-initialization A/B controls. Same-name/shape fresh R84 seed2 parameters are copied; resized/new parameters use isolated seed104002. Original R88B physical TSP node0, raw inputs/geometry/distance, methodID+34, original full score MLP, utility differences/maximin and R77 CE/comparison/risk/two-view R-Drop remain. No teacher, context labels, features, routes or new targets. R89A AdamW, R89B Muon and R84A are historical same-data references, not freshly matched controls; R89B differs in optimizer. Hardware and microbatch differences preclude a clean capacity-only causal claim. Fixed R84 target36, true rescues/losses, false picks and net cost are separate from selection frequency. Native/FP64 ind and legal pools; TRAIN-fitted SBS; ALL task-equal. TEST is development-exposed; one whole-model minimum full18 VAL lock, no TEST/LIB reselection. Timing includes the complete current standalone model on original GPU-resident inputs, not donor/cached-feature latency; excludes dataset I/O, transfer and solver execution.

R104/A_width256 | Fixed target36 vs R89A

True rescues, losses, false picks and aligned full-population net cost; no reselection.

Split	Winners	Rescued	Lost	New false	Net cost delta
TRAIN	83279	1188	51	964	-0.000387700
VAL	4625	51	2	59	0.000185600
TEST	4598	50	6	73	0.000325119

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	0.223731	32990	29728	-0.000387700	-0.003027
VAL	-0.098765	1697	1777	0.000185600	0.003219
TEST	-0.106173	1665	1751	0.000325119	0.006257

Historical only, not retrained/matched; R89A AdamW, R89B Muon, R84A original AdamW.

ALL net cost is the equal-task mean of aligned instance cost differences, not target frequency.

R104/A_width256 | Fixed target36 vs R89B

True rescues, losses, false picks and aligned full-population net cost; no reselection.

Split	Winners	Rescued	Lost	New false	Net cost delta
TRAIN	83279	812	368	795	0.000648626
VAL	4625	30	18	47	0.000723579
TEST	4598	39	24	59	0.000640153

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	-0.222908	25774	29024	0.000648626	0.008039
VAL	-0.296296	1379	1619	0.000723579	0.008224
TEST	-0.293827	1370	1608	0.000640153	0.010847

Historical only, not retrained/matched; R89A AdamW, R89B Muon, R84A original AdamW.

ALL net cost is the equal-task mean of aligned instance cost differences, not target frequency.

R104/A_width256 | Fixed target36 vs R84A

True rescues, losses, false picks and aligned full-population net cost; no reselection.

Split	Winners	Rescued	Lost	New false	Net cost delta
TRAIN	83279	1224	53	987	-0.000435357
VAL	4625	55	2	62	0.000406386
TEST	4598	53	6	76	0.000267561

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	0.224348	32718	29447	-0.000435357	-0.003525
VAL	-0.133333	1702	1810	0.000406386	0.005406
TEST	-0.109877	1620	1709	0.000267561	0.005646

Historical only, not retrained/matched; R89A AdamW, R89B Muon, R84A original AdamW.

ALL net cost is the equal-task mean of aligned instance cost differences, not target frequency.