R104/B_depth8 | Capacity and actual latency

Complete standalone model; synchronized GPU-resident original inputs; no donor.

Problem	Instances	Milliseconds / instance
TSP	128	0.646720
CVRP	128	0.787508
ATSP	128	0.110263
OVRP	128	0.105366
VRPB	128	0.105402
VRPL	128	0.105198
VRPTW	128	0.105517
OVRPTW	128	0.105573
OVRPB	128	0.105670
OVRPL	128	0.105495
VRPBL	128	0.105677
VRPBTW	128	0.105568
VRPLTW	128	0.105783
OVRPBL	128	0.105781
OVRPBTW	128	0.105529
OVRPLTW	128	0.105516
VRPBLTW	128	0.105383
OVRPBLTW	128	0.105505

Parameters: 3,426,057; copied fresh: 1,842,953; independent: 1,583,104

Width256 instance4/joint2 and width128 instance8/joint4 are two capacity axes, not equal-parameter, equal-FLOPs or identical-full-initialization A/B controls. Same-name/shape fresh R84 seed2 parameters are copied; resized/new parameters use isolated seed104002. Original R88B physical TSP node0, raw inputs/geometry/distance, methodID+34, original full score MLP, utility differences/maximin and R77 CE/comparison/risk/two-view R-Drop remain. No teacher, context labels, features, routes or new targets. R89A AdamW, R89B Muon and R84A are historical same-data references, not freshly matched controls; R89B differs in optimizer. Hardware and microbatch differences preclude a clean capacity-only causal claim. Fixed R84 target36, true rescues/losses, false picks and net cost are separate from selection frequency. Native/FP64 ind and legal pools; TRAIN-fitted SBS; ALL task-equal. TEST is development-exposed; one whole-model minimum full18 VAL lock, no TEST/LIB reselection. Timing includes the complete current standalone model on original GPU-resident inputs, not donor/cached-feature latency; excludes dataset I/O, transfer and solver execution.

R104/B_depth8 | Fixed target36 vs R89A

True rescues, losses, false picks and aligned full-population net cost; no reselection.

Split	Winners	Rescued	Lost	New false	Net cost delta
TRAIN	83279	446	87	373	-0.000554436
VAL	4625	28	3	18	-0.000392184
TEST	4598	23	2	30	0.000024878

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	0.172840	23132	20612	-0.000554436	-0.004199
VAL	0.040741	1200	1167	-0.000392184	-0.003686
TEST	0.134568	1275	1166	0.000024878	0.000890

Historical only, not retrained/matched; R89A AdamW, R89B Muon, R84A original AdamW.

ALL net cost is the equal-task mean of aligned instance cost differences, not target frequency.

R104/B_depth8 | Fixed target36 vs R89B

True rescues, losses, false picks and aligned full-population net cost; no reselection.

Split	Winners	Rescued	Lost	New false	Net cost delta
TRAIN	83279	182	516	254	0.000481890
VAL	4625	10	22	12	0.000145796
TEST	4598	12	20	20	0.000339912

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	-0.273800	19900	23892	0.000481890	0.006868
VAL	-0.156790	1109	1236	0.000145796	0.001318
TEST	-0.053086	1146	1189	0.000339912	0.005480

Historical only, not retrained/matched; R89A AdamW, R89B Muon, R84A original AdamW.

ALL net cost is the equal-task mean of aligned instance cost differences, not target frequency.

R104/B_depth8 | Fixed target36 vs R84A

True rescues, losses, false picks and aligned full-population net cost; no reselection.

Split	Winners	Rescued	Lost	New false	Net cost delta
TRAIN	83279	458	65	380	-0.000602093
VAL	4625	30	1	18	-0.000171397
TEST	4598	26	2	32	-0.000032680

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	0.173457	22443	19914	-0.000602093	-0.004697
VAL	0.006173	1179	1174	-0.000171397	-0.001500
TEST	0.130864	1215	1109	-0.000032680	0.000279

Historical only, not retrained/matched; R89A AdamW, R89B Muon, R84A original AdamW.

ALL net cost is the equal-task mean of aligned instance cost differences, not target frequency.