R106 | Native adaptation: whole A/B

Identical initial complete state, different trainable freedom/compute; no TEST/LIB model selection.

Model	Split / Problem	Top1 (%)	mean_cost	vs_SBS (%)	regret (%)
R106/A_frozen	VAL/ALL	71.33580	11.41603475	-0.19949	0.39749
R106/A_frozen	TEST/ALL	71.52346	11.37224262	-0.20925	0.38187
R106/A_frozen	TSPLIB/TSP	20.40816	8.21911394	-0.25874	0.96807
R106/A_frozen	CVRPLIB/CVRP	64.00000	67.14442315	-0.15011	0.54603
R106/B_finetune	VAL/ALL	70.94938	11.41643226	-0.19568	0.40154
R106/B_finetune	TEST/ALL	71.16049	11.37339918	-0.19961	0.39079
R106/B_finetune	TSPLIB/TSP	22.44898	8.21436358	-0.31639	0.89645
R106/B_finetune	CVRPLIB/CVRP	59.00000	67.16485442	-0.11973	0.55173

Native LR B = common x0.1 from first update; A frozen; both native eval and all splits online.

Initial state SHA256: 7381a063ecba3c898ed2ebabcb9d06a4bf30a70db6767ecf09cf230c165b0e41

R106 compares frozen A with trainable B native encoders under the same final original selector supervision. Both start from identical complete fresh R84 seed2 selector tensors, three projections and original R102B pretrained ELG/MatNet/ReLD weights. Native pretraining is not selector continuation; its data/overlap remains unknown. Total tensors are identical at initialization, but trainable freedom and training compute are NOT matched. Native stays eval in both arms; B receives gradients from update1 with effective LR0.1 times the common warmup/cosine LR. A does not update native parameters. One input-only native graph per microbatch is shared by the original two dropout selector forwards, with the unchanged CE/comparison/risk/R-Drop and real-row normalization. No auxiliary head or target. All TRAIN/VAL/TEST/LIB calls use online native encoders, never native-feature cache substitution. Native FP32 output precedes normal fusion/selector precision. Raw node4/joint2, physical TSP node0 role, methodID+hand34 and score MLP/maximin remain. Padding-adapted native execution is shared across arms; numerical fixture agreement is not bitwise equivalence or efficacy. Native/FP64 ind, legal-pool Oracle, full-TRAIN SBS; ALL is equal-task macro, mixed raw cost descriptive only. The fixed R84 target36 reports true rescues/losses, false picks and aligned net cost, not choice frequency as recovery. R94B and completed R102B recovery report view are historical, NOT newly matched controls; external pretraining and historical compute are not controlled. Development-exposed TEST; one whole18 minimum-VAL/first-tie checkpoint per arm, no TEST/LIB reselection or per-problem assembly. Actual all18 online native+selector timing is GPU-resident batch128 amortized, NOT singleton or I/O-inclusive end-to-end. Dataset I/O, CPU transfer, training and solver decoding/execution are excluded. Gradient/state witnesses are implementation checks, not selection-effectiveness evidence.

R106 | Fixed target36 B versus A

True recovery, false picks and aligned full-population net cost; not method-frequency changes.

Split	Winners	Rescued	Lost	New false	Full18 net cost delta
TRAIN	83279	336	1480	542	0.001413960
VAL	4625	18	60	41	0.000397503
TEST	4598	23	60	31	0.001156564

Split	Top1 delta pp	Corrected	Harmed	Cost delta	Regret delta pp
TRAIN	-0.588409	26756	35335	0.001413960	0.014404
VAL	-0.386420	1587	1900	0.000397503	0.004044
TEST	-0.362963	1610	1904	0.001156564	0.008920

ALL equal-task; raw mean cost is descriptive only. No per-problem checkpoint assembly.